Rafael Martí
Gerhard Reinelt

The Linear Ordering Problem

Exact and Heuristic Methods
in Combinatorial Optimization

Springer

Applied Mathematical Sciences
Volume 175

Rafael Martí   ·   Gerhard Reinelt

The Linear Ordering Problem

Exact and Heuristic Methods
in Combinatorial Optimization

Springer

Prof. Dr. Rafael Martí
Dpto Estadística e Inv. Operativa
Universitat de València
Dr. Moliner 50
46100 Valencia
Spain
rafael.marti@uv.es

Prof. Dr. Gerhard Reinelt
Universität Heidelberg
Inst. Informatik
Im Neuenheimer Feld 326
69120 Heidelberg
Germany
gerhard.reinelt@informatik.uni-heidelberg.de

Editors:

S.S. Antman
Department of Mathematics
and
Institute for Physical
Science and Technology
University of Maryland
College Park, MD 20742-4015
USA
ssa@math.umd.edu

J.E. Marsden
Control and Dynamical Systems, 107-81
California Institute of Technology
Pasadena, CA 91125
USA
marsden@cds.caltech.edu

L. Sirovich
Laboratory of Applied Mathematics
Department of Biomathematical Sciences
Mount Sinai School of Medicine
New York, NY 10029-6574
isirovich@rockefeller.edu

Cover design: deblik, Berlin

Printed on acid-free paper

*To Amparo Rico and Ximo Seró, el padrino,
for showing me the way.*

Rafa Martí

Preface

The idea for writing this book came up when the authors met at the University of Valencia in 2005. While comparing our experiences with regard to various aspects of the linear ordering problem (LOP), we realized that most of the optimization technologies had been successfully applied to solve this problem. We also found that there were only a small number of books covering all state-of-the-art optimization methods for hard optimization problems (especially considering both exact methods and heuristics together). We thought that the LOP would make an ideal example to survey these methods applied to one problem and felt the time was ripe to embark on the project of writing this monograph.

Faced with the challenge of solving hard optimization problems that abound in the real world, classical methods often encounter serious difficulties. Important applications in business, engineering or economics cannot be tackled by the solution methods that have been the predominant focus of academic research throughout the past three decades. Exact and heuristic approaches are dramatically changing our ability to solve problems of practical significance and are extending the frontier of problems that can be handled effectively. In this text we describe state-of-the-art optimization methods, both exact and heuristic, for the LOP. We actually employ the LOP to illustrate current optimization technologies and the design of successful implementations of exact and heuristic procedures. Therefore, we do not limit the scope of this book to the LOP but, on the contrary, we provide the reader with the background and strategies in optimization to tackle different combinatorial problems.

This monograph is devoted to the LOP, its origins, applications, instances and especially to methods for its effective approximate or exact solution. Our intention is to provide basic principles and fundamental ideas and reflect the state-of-the-art of heuristic and exact methods, thus allowing the reader to create his or her personal successful applications of the solution methods. The book is meant to be of interest for researchers and practitioners in computer science, mathematics, operations research, management science, industrial engineering, and economics. It can be used as a textbook on issues of practical optimization in a master's course or as a reference resource for engineering optimization algorithms.

To make the book accessible to a wider audience, it is to a large extent self-contained, providing the reader with the basic definitions and concepts in optimization. However, in order to limit the size of this monograph we have not included extensive introductions. Readers interested in further details are referred to appropriate textbooks such as [4, 84, 102, 117, 118, 124].

The structure of this book is as follows. Chapter 1 provides an introduction to the problem and its applications and describes the set of benchmark instances which we are using for our computational experiments and which have been made publically available. Chapter 2 describes such basic heuristic methods such as construction and local searches. Chapter 3 expands on Chapter 2 and covers meta-heuristics in which the simple methods are now embedded in complex solution algorithms based on different paradigms, such as evolution or learning strategies. Chapter 4 discusses branch-and-bound, the principal approach for solving difficult problems to optimality. A special version based on polyhedral combinatorics, branch-and-cut, is presented in Chapter 5. Chapter 6 deals in more detail with the linear ordering polytope which is at the core of branch-and-cut algorithms. The book concludes with Chapter 7, where a number of further aspects of the LOP and potential issues for further research are described.

Valencia, Heidelberg   *Rafael Martí*
October 2010   *Gerhard Reinelt*

Contents

Chapter 1
Introduction

Abstract The linear ordering problem (LOP) is one of the classical combinatorial optimization problems which was already classified as NP-hard in 1979 by Garey and Johnson [50]. It has received considerable attention in various application areas ranging from archeology and scheduling to economics and even mathematical psychology. Solution methods for the LOP have been proposed since 1958, when Chenery and Watanabe outlined some ideas on how to obtain solutions for this problem. The interest in this problem has continued over the years, resulting in the book [111] and many recent papers in scientific journals. This chapter surveys the main LOP applications and instances. We have compiled a comprehensive set of benchmark problems including all problem instances which have so far been used for conducting computational experiments. Furthermore we have included new instances. All of them form the new benchmark library LOLIB. We will use them in the next chapters to report our experiments with heuristics, meta-heuristics and exact approaches for the LOP.

1.1 Basic definitions

In its graph version the LOP is defined as follows. Let $D_n = (V_n, A_n)$ denote the complete digraph on n nodes, i.e., the directed graph with node set $V_n = \{1, 2, \ldots, n\}$ and the property that for every pair of nodes i and j there is an arc (i, j) from i to j and an arc (j, i) from j to i. A *tournament* (or *spanning tournament*) T in A_n consists of a subset of arcs containing for every pair of nodes i and j either arc (i, j) or arc (j, i), but not both. A (*spanning*) *acyclic tournament* is a tournament without directed cycles, i.e., not containing an arc set of the form $\{(v_1, v_2), (v_2, v_3), \ldots, (v_k, v_1)\}$ for some $k > 1$ and distinct nodes $v_1, v_2, \ldots, v_k$.

A *linear ordering* of the nodes $\{1, 2, \ldots, n\}$ is a ranking of the nodes given as linear sequence, or equivalently, as a permutation of the nodes. We denote the linear ordering that ranks node v_1 first, v_2 second, etc., and v_n last by $\langle v_1, v_2 \ldots, v_n \rangle$ and write $v_i \prec v_j$ if node v_i is ranked before node v_j. If σ denotes a linear

ordering, then $\sigma(i)$ gives the position of node i in this ordering. We will also consider *partial orderings* where only a subset of the nodes is ranked or only some pairs are compared.

It is easy to see that an acyclic tournament T in A_n corresponds to a linear ordering of the nodes of V_n and vice versa: the node ranked first is the one without entering arcs in T, the node ranked second is the one with one entering arc (namely from the node ranked first), etc., and the node ranked last is the one without leaving arcs in T.

Usually, ordering relations are weighted and we have weights c_{ij} giving the benefit or cost resulting when node i is ranked before node j or, equivalently, when the arc (i, j) is contained in the acyclic tournament. The (weighted) *linear ordering problem* is defined as follows.

Linear ordering problem

Given the complete directed graph $D_n = (V_n, A_n)$ with arc weights c_{ij} for every pair $i, j \in V_n$, compute a spanning acyclic tournament T in A_n such that $\sum_{(i,j) \in T} c_{ij}$ is as large as possible.

Alternatively, the LOP can be defined as a matrix problem, the so-called *triangulation problem*.

Triangulation problem

Let an (n,n)-matrix $H = (H_{ij})$ be given. Determine a simultaneous permutation of the rows and columns of H such that the sum of superdiagonal entries becomes as large as possible.

Obviously, by setting arc weights $c_{ij} = H_{ij}$ for the complete digraph D_n, the triangulation problem for H can be solved as a linear ordering problem in D_n. Conversely, a linear odering problem for D_n can be transformed to a triangulation problem for an (n,n)-matrix H by setting $H_{ij} = c_{ij}$ and the diagonal entries $H_{ii} = 0$.

Consider as an example the (5,5)-matrix

$$H = \begin{pmatrix} 0 & 16 & 11 & 15 & 7 \\ 21 & 0 & 14 & 15 & 9 \\ 26 & 23 & 0 & 26 & 12 \\ 22 & 22 & 11 & 0 & 13 \\ 30 & 28 & 25 & 24 & 0 \end{pmatrix}.$$

The sum of its superdiagonal elements is 138. An optimum triangulation is obtained if the original numbering $(1,2,3,4,5)$ of the rows and columns is changed to $(5,3,4,2,1)$, i.e., the original element H_{12} becomes element $H_{\sigma(1)\sigma(2)} = \tilde{H}_{54}$ in the permuted matrix. Thus the optimal triangulation of H is

$$\tilde{H} = \begin{pmatrix} 0 & 25 & 24 & 28 & 30 \\ 12 & 0 & 26 & 23 & 26 \\ 13 & 11 & 0 & 22 & 22 \\ 9 & 14 & 15 & 0 & 21 \\ 7 & 11 & 15 & 16 & 0 \end{pmatrix}.$$

Now the sum of superdiagonal elements is 247.

1.2 Applications of the Linear Ordering Problem

We review some of the many applications of the linear ordering problem.

1.2.1 Equivalent Graph Problems

The *acyclic subdigraph problem* (ASP) is defined as follows. Given a directed graph $D = (V, A)$ with arc weights d_{ij}, for all $(i, j) \in A$, determine a subset $B \subseteq A$ which contains no directed cycles and has maximum weight $d(B) = \sum_{(i,j) \in B} d_{ij}$.

It can easily be seen that this problem is equivalent to the LOP. For a given ASP define a LOP on D_n, where $n = |V|$, by setting for $1 \leq i, j \leq n$, $i \neq j$:

$$c_{ij} = \begin{cases} \max\{0, d_{ij}\}, & \text{if } (i, j) \in A, \\ 0, & \text{otherwise.} \end{cases}$$

If T is a tournament of maximum weight, then $B = \{(i, j) \in T \cap A \mid c_{ij} > 0\}$ is an acyclic subdigraph of D of maximum weight. In the opposite direction, by adding a suitably large constant, we can transform a given LOP into an equivalent one where all weights are strictly positive. Then an acyclic subdigraph of maximum weight is a tournament.

The *feedback arc set problem* (FBAP) in a weighted digraph $D = (V, A)$ consists of finding an arc set B of minimum weight such that $A \setminus B$ is acyclic, i.e., such that B is a so-called *feedback arc set* intersecting every dicycle of D. Obviously, FBAP and ASP are equivalent because they are complementary questions.

Fig. 1.1 shows a digraph on 9 nodes where the arcs of a minimum feedback arc set are drawn as dotted lines. If the six arcs of the feedback arc set are removed, we obtain an acyclic arc set.

Fig. 1.1 A digraph with
minimum feedback arc set

1.2.2 Related Graph Problems

There are some further problems dealing with acyclic subdigraphs. The *node in-
duced acyclic subdigraph problem* asks for a node set $W \subseteq V$ such that the subdi-
graph $(W, A(W))$ is acyclic. (Here $A(W)$ denotes the set of arcs with both end nodes
in W.) The problem can be defined either with node weights d, and $d(W)$ is to be
maximized, or with arc weights c where $c(A(W))$ has to be maximum. Analogously,
the *feedback node set problem* is to find a set $W \subseteq V$ such that $(V \setminus W, A(V \setminus W))$ is
acyclic. Here, sums of node weights or arc weights have to be minimized.

The request that solution digraphs have to be node induced adds a further com-
plexity. These problems cannot be transformed to a pure linear ordering problem
and are even more difficult.

1.2.3 Aggregation of Individual Preferences

Linear ordering problems may occur whenever rankings of some objects are to
be determined. Consider for example the following situation. A set of n objects
$O_1, O_2, \ldots, O_n$ is given which have to be rated by m persons according to their in-
dividual preferences. Then a ranking of these objects is to be found which reflects
these single rankings as closely as possible. The first question to be answered is
how the individual rankings can be obtained. One solution is a pairwise comparison
experiment. For any pair O_i and O_j, $1 \le i < j \le n$, of objects each person decides
whether O_i should be preferred to O_j or vice versa. The results of these $m\binom{n}{2}$ com-
parisons are stored in an (n, n)-matrix $H = (H_{ij})$ where H_{ij} = number of persons
preferring object O_i to object O_j. A ranking of these objects which infers as few
contradictions to the individual rankings as possible can be obtained by triangulat-
ing H. It should be remarked that there are various statistical methods to aggregate
single preference relations to one relation.

This area of application is the oldest one of the LOP. In 1959 Kemeny [77] posed
the following problem (*Kemeny's problem*). Suppose that there are m persons and

each person i, $i = 1,\dots,m$, has ranked n objects by giving a linear ordering T_i of the objects. Which common linear ordering aggregates the individual orderings in the best possible way? We can solve this problem as a linear ordering problem by setting c_{ij} = number of persons preferring object O_i to object O_j. Note that this is basically the problem stated above, but this time the relative ranking of the objects by each single person is consistent (which is not assumed above).

   Slater [119], in 1961, asked for the minimum number of arcs that have to be reversed to convert a given tournament T into an acyclic tournament. In the context of preferences, the input now is a collection of rankings for all pairs i and j of objects stating whether i should be preferred to j or vice versa and the problem is to find a the maximum number of pairwise rankings without contradiction. Also *Slater's problem* can also be solved as a LOP, namely by setting

$$c_{ij} = \begin{cases} 1, & \text{if } (i,j) \in T, \\ 0, & \text{otherwise.} \end{cases}$$

Questions of this type naturally occur in the context of voting (How should a fair distribution of seats to parties be computed from the votes of the electors?) and have already been studied in the 18th century by Condorcet [37].

1.2.4 Binary Choice Probabilities

Let S_n denote the set of all permutations of $\{1,2,\dots,n\}$ and let P be a probability distribution on S_n.

   Define the *induced (binary choice) probability system p* for $\{1,2,\dots,n\}$ as the mapping $p : \{1,2\dots,n\} \times \{1,2\dots,n\} \setminus \{(i,i) \mid i = 1,2,\dots,n\} \to [0,1]$ where

$$p(i,j) = \sum_{S \in S_n,\, i \prec j \text{ in } S} P(S).$$

The question of whether a given vector p is a vector of binary choice probabilities according to this definition is of great importance in mathematical psychology and the theory of social choice (see [48] for a survey).

   In fact, the set of binary choice vectors is exactly the linear ordering polytope which will play a prominent role later in this book.

1.2.5 Triangulation of Input-Output Tables

One field of practical importance in economics is *input-output analysis*. It was pioneered by Leontief [88, 89] who was awarded the Nobel Prize in 1973 for his fundamental achievements. The central component of input-output analysis is the

so-called *input-output table* which represents the dependencies between the different branches of an economy. To make up an input-output table the economy of a country is divided into sectors, each representing a special branch of the economy. An input-output table shows the transactions between the single sectors in a certain year. To be comparable with each other all amounts are given in monetary values. Input-output analysis is used for forecasting the development of industries and for structural planning (see [69] for an introductory survey).

Triangulation is a means for a descriptive analysis of the transactions between the sectors. In a simple model of production structure the flow of goods begins in sectors producing raw material, then sectors of manufacturing follow, and in the last stage goods for consumption and investments are produced. A real economy, of course, does not show such a strict linearity in the interindustrial connections, here there are flows between almost any sectors. Nevertheless it can be observed that the main stream of flows indeed goes from primary stage sectors via the manufacturing sectors to the sectors of final demand. Triangulation is a method for determining a hierarchy of all sectors such that the amount of flow incompatible with this hierarchy (i.e., from sectors ranked lower to sectors ranked higher) is as small as possible. Such rankings allow interpretations of the industrial structure of a country and comparisons between different countries.

1.2.6 Optimal Weighted Ancestry Relationships

This application from anthropology has been published in [56]. Consider a cemetery consisting of many individual gravesites. Every gravesite contains artifacts made of different pottery types. As gravesites sink over the years and are reused, it is a reasonable assumption that the depth of a pottery type is related to its age. So every gravesite gives a partial ordering of the pottery types contained in it. These partial orderings may not be consistent in the sense that pairs of pottery types may be ranked differently depending on the gravesite. The task of computing a global ordering with as few contradictions as possible amounts to solving a linear ordering problem in the complete directed graph where the nodes correspond to the pottery types and the arc weights are aggregations of the individual partial orderings. In [56] several possibilities for assigning arc weights are discussed and a simple heuristic for deriving an ordering is presented.

1.2.7 Ranking in Sports Tournaments

In many soccer leagues each team plays each other team twice. The winner of a match gets three points, in case of a tie both teams get one point. In the standard procedure, the final ranking of the teams in the championship is made up by adding these points and breaking ties by considering the goals scored. Another

possible method of ranking the teams leads to a linear ordering problem. If there are n teams which have played each other twice we construct an (n,n)-matrix H by setting H_{ij} = number of goals which were scored by team i against team j. A triangulation of this matrix yields a ranking of the teams which takes the number of goals scored and (implicitly) the number of matches won into account. Moreover, transitive relations are important, and winning against a top team counts more than beating an average team.

As an example we compare the official ranking of the English Premier League in the season 2006/2007 with a ranking obtained by triangulation. Table 1.1 shows on the left side the official ranking and on the right side an optimum linear ordering. There are alternate optima, however, which we do not list here, but which would make this type of ranking approach problematical for practical use.

Table 1.1 Premier League 2006/2007 (left: official, right: triangulated)

1 Manchester United	1 Chelsea
2 Chelsea	2 Arsenal
3 Liverpool	3 Manchester United
4 Arsenal	4 Everton
5 Tottenham Hotspur	5 Portsmouth
6 Everton	6 Liverpool
7 Bolton Wanderers	7 Reading
8 Reading	8 Tottenham Hotspur
9 Portsmouth	9 Aston Villa
10 Blackburn Rovers	10 Blackburn Rovers
11 Aston Villa	11 Middlesborough
12 Middlesborough	12 Charlton Athletic
13 Newcastle United	13 Bolton Wanderers
14 Manchester City	14 Wigan Athletic
15 West Ham United	15 Manchester City
16 Fulham	16 Sheffield United
17 Wigan Athletic	17 Fulham
18 Sheffield United	18 Newcastle United
19 Charlton Athletic	19 Watford
20 Watford	20 West Ham United

1.2.8 Corruption Perception

The organisation *Transparency International* [122] releases an annual *corruption perception index* which ranks more than 150 countries by their perceived level of corruption. This index is computed from expert assessments and opinion surveys. The respective assessments and surveys only consider a subset of all countries and can thus also be viewed as a partial ordering. In [1] the linear ordering problem is used to aggregate these partial rankings. It is shown that the solution of the linear

ordering problem agrees with the ranking according to the index to a large extent, but exhibits interesting differences for some countries.

1.2.9 Crossing Minimization

The relatively new research field *graph drawing* is concerned with finding good drawings of graphs for making the relations they represent easier to understand. In one of the many problems, the LOP could be employed. Let $G = (V, E)$ be a bipartite graph with the bipartition $V = V_1 \cup V_2$ of its nodes. A basic problem in graph drawing considers the task of drawing the nodes (linearly ordered) on two opposite horizontal lines and drawing the edges as straight lines such that the number of crossing lines is minimized. In [73] it is observed that the so-called *one sided crossing minimization problem* (where the permutation of one of the two sides is fixed) can be solved as a linear ordering problem with good results in practice. By embedding this procedure into a branch-and-bound algorithm, the two sided crossing minimization (without fixing) can also be solved. Figure 1.2 (taken from [73]) shows the results of a heuristic with 30 crossings and the optimal drawing with only 4 crossings.

Fig. 1.2 Two sided crossing minimization

1.2.10 Linear Ordering with Quadratic Objective Function

Bipartite crossing minimization without fixed nodes can also be modelled directly employing the LOP. However in this case the objective function is not linear anymore, but contains products $x_{ij}x_{kl}$. A successful application of this model (using methods different from the ones presented in this book) is reported in [17].

1.2.11 Scheduling with Precedences

Consider a set of n jobs which have to be processed on a single machine. Each job i has a processing time p_i and a weight w_i. Furthermore, there is a set P of job pairs (i, j), each specifying that job i has to be executed before job j. (It is assumed that there are no contradictions within these precedences.) The task consists of finding a linear ordering $\langle k_1, k_2, \ldots, k_n \rangle$ of the jobs such that the total weighted completion time $\sum_{i=1}^{n} w_i t_i$ is minimized, where

$$t_{k_i} = \sum_{j=1}^{i} p_{k_j}$$

and the precedences given by P are observed. This problem can be modeled as a *linear ordering extension problem* where some relative rankings are already fixed beforehand. Solution methods are discussed in [12].

1.2.12 Linear Ordering with Cumulative Costs

Here, in addition to the arc weights c_{ij} of the standard linear ordering problem, there are node weights p_i, $1 \leq i \leq n$, and the task is to find an ordering $\langle k_1, k_2, \ldots, k_n \rangle$ of the nodes minimizing the cost $\sum_{i=1}^{j} \alpha_i$, where

$$\alpha_{k_i} = p_{k_i} + \sum_{j=i+1}^{n} c_{k_i k_j} \alpha_{k_j}.$$

An application of this problem for optimizing UMTS mobile phone telecommunication and its solution with a mixed-integer programming approach is discussed in [10]. Heuristic algorithms are presented in [43].

1.2.13 Coupled Task Problem

Many combinatorial optimization problems require as one constraint that some subset of elements is linearly ordered. We give a brief account to one of them where we could successfully use linear ordering variables in its optimization model [9]. The *coupled task problem* deals with scheduling n jobs each of which consists of two subtasks and where there is the additional requirement that between the execution of these subtasks an exact delay is required. If $\{J_1, J_2, \ldots, J_n\}$ is the set of jobs and $\{T_1, T_2, \ldots, T_{2n}\}$ the set of tasks, where T_{2i-1} and T_{2i} denote the first and second subtask of J_i, then one requirement for a feasible schedule is that all tasks are linearly ordered. We can model this constraint by introducing binary variables y_{kl} indicating

whether task T_k is scheduled before task T_l or not. Of course, additional constraints are necessary to take processing times and gaps properly into account.

1.2.14 Target Visitation Problem

This optimization problem is a composition of the linear ordering problem and the traveling salesman problem and was proposed in [62]. Suppose that, starting from some origin 0, a set $\{1,\ldots,n\}$ of n targets has to be visited. In addition to the distance traveled, priorities have also to be taken into account. If d_{ij} (d_{0i} and d_{i0}) denotes the distance between two targets (the distance between the origin and a target and the distance between a target and the origin) and c_{ij} the gain when target i is visited before target j, the *target visitation problem* consists of finding a visiting sequence $\langle k_1, k_2, \ldots, k_n \rangle$ of the targets maximizing the objective function

$$\sum_{i=1}^{n-1} \sum_{j=i+1}^{n} c_{k_i k_j} - \left(d_{0k_1} + \sum_{i=1}^{n-1} d_{k_i k_{i+1}} + d_{k_n 0} \right).$$

In the literature, the term "linear ordering problem" is sometimes misused also for some problems where a linear ordering of the nodes of some weighted graph has to be found, but where the objective function is a different one. For example, [3] actually considers the *linear arrangement problem* and [76] requests as the objective that the capacity of a cut between two adjacent nodes is as small as possible, i.e., it deals with the *cut-width problem*.

1.3 Benchmark Problems

We have compiled a set of benchmark problem instances of the LOP. These problem instances consist of real-world as well as randomly generated data sets.

1.3.1 Data Format

For finding the optimum triangulation of a matrix, the diagonal entries are irrelevant. Furthermore, orderings compare in the same way if some constant C is added to both entries H_{ij} and H_{ji}. In particular, the optimality of an ordering is not affected by this transformation. However, the quality of bounds does change. If we take diagonal entries into account and add a large constant to every matrix entry, then every feasible solution is close to optimal and no real comparison of qualities is possible. Therefore we transform every problem instance to a suitable normal form.

Definition 1.1. A quadratic (n,n)-matrix H is in *normal form* if

 (i) all entries of H are integral and nonnegative,
 (ii) $H_{ii} = 0$ for all $i = 1, \ldots, n$,
 (iii) $\min\{H_{ij}, H_{ji}\} = 0$ for all $1 \leq i < j \leq n$.

The following example shows a matrix and its normal form.

$$
\begin{pmatrix}
-17 & 36 & 11 & 45 & 7 \\
21 & 22 & 44 & 15 & 9 \\
26 & 23 & 13 & 26 & 12 \\
22 & 22 & 11 & 0 & 33 \\
30 & 9 & 25 & 24 & -7
\end{pmatrix}
\qquad
\begin{pmatrix}
0 & 15 & 0 & 23 & 0 \\
0 & 0 & 21 & 0 & 0 \\
15 & 0 & 0 & 15 & 0 \\
0 & 7 & 0 & 0 & 9 \\
23 & 0 & 13 & 0 & 0
\end{pmatrix}
$$

For our computations, all matrices H are transformed to their normal form $\overline{H}$. Note that this normal form is unique. Since the "true" value of a linear ordering might be interesting for some application, we compute

$$
s_d = \sum_{i=1}^{n} H_{ii} \text{ and } s_t = \sum_{i=1}^{n-1} \sum_{j=i+1}^{n} \min\{H_{ij}, H_{ji}\}.
$$

Now, if $c_{\mathrm{opt}}(\overline{H})$ is the value of an optimum triangulation of $\overline{H}$, then $c_{\mathrm{opt}}(H) = c_{\mathrm{opt}}(\overline{H}) + s_d + s_t$.

In normal form, a matrix can be seen as the specification of a weighted tournament (take all arcs (i, j) with $H_{ij} > 0$ and, if $H_{ij} = H_{ji} = 0$, then choose one of the arcs (i, j) or (j, i)). The LOP then directly corresponds to finding a minimum weight feedback arc set for this tournament.

Definition 1.2. Let H be a matrix and σ be an optimum linear ordering. Then the number

$$
\lambda(H) = \frac{\sum\limits_{\sigma(i) < \sigma(j)} H_{ij}}{\sum\limits_{i \neq j} H_{ij}}
$$

is called *degree of linearity* of H.

The degree of linearity gives the sum of the superdiagonal entries as the percentage of the total sum of the matrix entries (except for the diagonal elements) and allows for some interpretations in economical analysis. It lies between 0.5 (this is a trivial lower bound) and 1.0 (for a triangular matrix) and is an indicator for the closeness of a matrix to a triangular matrix. Note, that the degree of linearity differs depending on whether a matrix is in normal form or not, so some care has to be given when interpreting it.

From a computational point of view it turns out that problems with smaller degree of linearity tend to become more difficult. This is validated by some experiments with random matrices.

We now describe the problem instances selected for the benchmark library.

1.3.2 Input-Output Matrices

These are real-world data sets taken from input-output tables from various sources. The corresponding linear ordering problems are comparatively easy. They are thus more of interest for economists than for the assessment of heuristics for hard problems. So we have not conducted extensive experiments with them. Only the quality of simple heuristics can be assessed with these matrices. The original entries in these tables were not necessarily integral, but for LOLIB they were scaled to integral values.

European 44-Sector Tables

The Statistical Office of the European Communities (Eurostat) compiles input-output tables for the member states of the EC. Our benchmark set contains 31 matrices of dimension 44 from the years 1959 to 1975 (t59b11xx – t75u11xx).

Belgian 50-Sector Tables

These input-output tables (be75eec, be75np, be75oi and be75tot) of 1975 with 50 sectors were compiled for the Belgian economy.

German 56-Sector Tables

These matrices (tiw56n54 – tiw56r72) for some years between 1954 to 1975 were compiled by *Deutsches Institut für Wirtschaftsforschung (DIW)* for the *Federal Republic of Germany*.

German 60-Sector Tables

These input-output tables (stabu70, stabu74 and stabu75) were compiled by the *Statistisches Bundesamt* of the *Federal Republic of Germany* for the years 1970, 1974 and 1975. (In some publications, these matrices were named stabu1 – stabu3.)

US 79-Sector Table

The matrix usa79 is the input-output table for the economy of the United States for the year 1985. It has been made available by Knuth [80].

1.3.3 Randomly Generated Instances A (Type 1)

This is a set of random problems defined by Martí [93] that has been widely used for experiments. Problems are generated from a (0,100) uniform distribution. Sizes are 100, 150, 200 and 500 and there are 25 instances in each set for a total of 100. The names of these instances are t1dn.i (e.g. t1d200.25), where n is the dimension and i the number within the instances of the same size.

In their original definition, these problem instances are not in normal form. For the experiments in this monograph and for publication in the library LOLIB we only give values with respect to the normalized objective function.

1.3.4 Randomly Generated Instances A (Type 2)

This data set has also been defined by Martí [93]. Problems are generated by counting the number of times a sector appears in a higher position than another in a set of randomly generated permutations. For a problem of size n, $n/2$ permutations are generated.

There are 25 instances with sizes 100, 150 and 200, respectively. The names of these instances are t2dn.i (e.g. t2d150.12), where n is the dimension and i the number within instances of the same size.

1.3.5 Randomly Generated Instances B

In this kind of random problem we tried to influence the difficulty of the problems for computational experiments. To this end we generated integer matrices where the superdiagonal entries are drawn uniformly from the interval $[0, U_1]$ and the subdiagonal entries from $[0, U_2]$, where $U_1 \geq U_2$. The difference $U_1 - U_2$ affects the difficulty. Subsequently, the instances were transformed to normal form and a random permutation was applied.

For $n = 40$, we set $U_1 = 100$ and $U_2 = 100 + 4(i - 1)$ for problems p40-i. For $n = 44$ and $n = 50$ we set $U_1 = 100$ and $U_2 = 100 + 2(i - 1)$ for problems p44-i, p50-i, respectively.

1.3.6 SGB Instances

These instances were used in [86] and are taken from the *Stanford GraphBase* [80]. They are random instances with entries drawn uniformly distributed from $[0, 25000]$. The set has a total of 25 instances with $n = 75$. Instances are named sgb75.01 through sgb75.25.

1.3.7 Instances of Schiavinotto and Stützle

Some further benchmark instances have been created and used by Schiavinotto and Stützle [116]. These instances were generated from the input-output tables by replicating them to obtain larger problems. Thus, the distribution of numbers in these instances somehow reflects real input-output tables, but otherwise they behave more like random problems. This data set has been called XLOLIB, and instances with $n = 150$ and $n = 250$ are available.

1.3.8 Instances of Mitchell and Borchers

These instances have been used by Mitchell and Borchers for their computational experiments [100]. They are random matrices where the subdiagonal entries are uniformly distributed in $[0, 99]$ and the superdiagonal entries are drawn uniformly from $[0, 39]$. Furthermore a certain percentage of the entries was zeroed out.

1.3.9 Further Special Instances

We added some further instances that were used in some publications.

EX Instances

These random problems (EX1–EX6) were used in particular in [34] and [35].

econ Instances

The problems instances econ36 through econ77 were generated from the matrix usa79. They turned out not to be solvable as a linear program using only 3-dicycle inequalities.

Paley Graphs

Paley graphs, or more precisely *Paley tournaments*, have been used by Goemans and Hall [58] to prove results about the acyclic subdigraph polytope. These tournaments are defined as follows. Let $q = 3 \mod 4$ be a prime power. Define the digraph on q nodes corresponding to the elements of the finite field $GF(q)$. This digraph contains the arc (i, j) if and only if $j - i$ is a nonzero square in $GF(q)$. Some of them provide interesting difficult linear ordering problems.

`atp` Instances

These instances were created from the results of the ATP tennis tour 1993/1994. Nodes correspond to a selection of players and the weight of an arc (i, j) is the number of victories of player i against player j.

Table 1.2 summarizes the number of instances in each set described above. Moreover, it specifies the number of instances for which either the optimum or only an upper bound is known. In the computational experiments we call the set of 229 instances for which the optimum is known OPT-I, and the set of 255 instances for which an upper bound is known UB-I.

Table 1.2 Number of instances in each set

Set	#Instances	#Optima	#Upper Bounds
IO	50	50	–
SGB	25	25	–
RandomAI	100	–	100
RandomAII	75	25	50
RandomB	90	70	20
MB	30	30	–
XLOLIB	78	–	78
Special	36	29	6
Total	484	229	255

LOLIB is available at the web site `http://heur.uv.es/optsicom/LOLIB`. Also the currently best known values and upper bounds as well as the constants eliminated by the transformation to normal form can be found there.

Chapter 2
Heuristic Methods

Abstract Since the linear ordering problem is NP-hard, we cannot expect to be able to solve practical problem instances of arbitrary size to optimality. Depending on the size of an instance or depending on the available CPU time we will often have to be satisfied with computing approximate solutions. In addition, under such circumstances, it might be impossible to assess the real quality of approximate solutions. In this and in the following chapter we will deal with the question of how to find very good solutions for the LOP in short or reasonable time. The methods described in this chapter are called *heuristic algorithms* or simply *heuristics*. This term stems from the Greek word *heuriskein* which means to find or discover. It is used in the field of optimization to characterize a certain kind of problem-solving methods. There are a great number and variety of difficult problems, which come up in practice and need to be solved efficiently, and this has promoted the development of efficient procedures in an attempt to find good solutions, even if they are not optimal. These methods, in which the process speed is as important as the quality of the solution obtained, are called heuristics or *approximative algorithms*.

2.1 Introduction

As opposed to *exact methods*, which guarantee to give an optimum solution of the problem, heuristic methods only attempt to yield a good, but not necessarily optimum solution. Nevertheless, the time taken by an exact method to find an optimum solution to a difficult problem, if indeed such a method exists, is in a much greater order of magnitude than the heuristic one (sometimes taking so long that in many cases it is inapplicable). Thus we often resort to heuristic methods to solve real optimization problems.

Perhaps the following comment by Onwubolu and Babu [105] is a little far-fetched: "The days when researchers emphasized using deterministic search techniques to find optimal solutions are gone.". But it is true that in practice an engineer,

an analyst or a manager sometimes might have to make a decision as soon as possible in order to achieve desirable results.

Recent years have witnessed a spectacular growth in the development of heuristic procedures to solve optimization problems. This fact is clearly reflected in the large number of articles published in specialized journals. 1995 saw the first issue of the *Journal of Heuristics*, dedicated solely to the publication of heuristic procedures. In the same year the first international congress dealing with these methods, called the *Metaheuristic International Conference (MIC)*, was held in Breckenridge, Colorado (USA).

In addition to the need to find good solutions of difficult problems in reasonable time, there are other reasons for using heuristic methods, among which we want to highlight:

–   No method for solving the problem to optimality is known.
–   Although there is an exact method to solve the problem, it cannot be used on the available hardware.
–   The heuristic method is more flexible than the exact method, allowing, for example, the incorporation of conditions that are difficult to model.
–   The heuristic method is used as part of a global procedure that guarantees to find the optimum solution of a problem.

A good heuristic algorithm should fulfil the following properties:

–   A solution can be obtained with reasonable computational effort.
–   The solution should be near optimal (with high probability).
–   The likelihood for obtaining a bad solution (far from optimal) should be low.

There are many heuristic methods that are very different in nature. Therefore, it is difficult to supply a full classification. Furthermore, many of them have been designed to solve a specific problem without the possibility of generalization or application to other similar problems. The following outline attempts to give wide, non-excluding categories, under which to place the better-known heuristics:

Decomposition Methods

The original problem is broken down into sub-problems that are simpler to solve, bearing in mind, be it in a general way, that subproblems belong to the same problem class.

Inductive Methods

The idea behind these methods is to generalize the smaller or simpler versions to the whole case. Properties or techniques that have been identified in these cases which are easier to analyze, can be applied to the whole problem.

Reduction Methods

These involve identifying properties that are mainly fulfilled by the good solutions and introduce them as boundaries to the problem. The objective is to restrict the space of the solutions by simplifying the problem. The obvious risk is that the optimum solutions of the original problem may be left out.

Constructive Methods

These involve building a solution to the problem literally step by step from scratch. Usually they are deterministic methods and tend to be based on the best choice in each iteration. These methods have been widely used in classic combinatorial optimization.

Local Search Methods

In contrast to the methods previously mentioned, local improvement or local search starts with some feasible solution of the problem and tries to progressively improve it. Each step of the procedure carries out a movement from one solution to another one with a better value. The method terminates when, for a solution, there is no other accessible solution that improves it.

Even though all these methods have contributed to expanding our knowledge of solving real problems, the constructive and local search methods form the foundations of the meta-heuristic procedures [4], which will be described in the next chapter.

2.1.1 Assessing the Quality of Heuristics

There are diverse possibilities for measuring the quality of a heuristic, among which we find the following.

Comparison with the Optimum Solution

Although one normally resorts to an approximative algorithm, because no exact method exists to obtain an optimum solution or it is too time-consuming, sometimes a procedure is available that provides an optimum for a limited set of examples (usually small sized instances). This set of examples can be used to assess the quality of the heuristic method.

Normally, for each example, the following are measured: the percentaged deviation of the heuristic solution value as compared to the optimum one and the mean

of these deviations. If we denote by c_A the value of the solution delivered by heuristic A and by c_{opt} the optimum value of a given example, in a maximization problem like the LOP, the percentaged deviation, *PerDev*, is given by the expression

$$PerDev = 100 \cdot \frac{c_{opt} - c_A}{c_{opt}}.$$

(We assume that all feasible solutions have a positive value.)

Comparison with a Bound

There are situations when no optimum solution is available for a problem, not even for a limited set of examples. An alternative evaluation method involves comparing the value of the solution provided by the heuristic with a bound for the problem (a lower bound if it is a minimization problem and an upper bound if it is a maximization problem). Obviously the quality of fit will depend on the quality of the bound (closeness to optimal). Thus we must somehow have information about the quality of the aforementioned bound, otherwise the proposed comparison would not be of much interest.

Comparison with a Truncated Exact Method

An enumerative method like branch-and-bound explores very many solutions, even though this may be a fraction of the total, and therefore large-scale problems can be computationally out of reach using these methods. Nevertheless, we can establish a limit on the maximum number of iterations (or on the CPU time) to run the exact algorithm. Moreover, we can modify the criteria to fathom a node in the search tree by adding or subtracting (depending on whether it is a minimization or maximization problem) a value Δ to the bound of the node thus fathoming a larger number of nodes and speeding up the method. In this way it guarantees that the value of the best solution provided by the procedure is no further than distance Δ from the optimal value to the problem. In any case, the best solution found with these truncated procedures establishes a bound against which the heuristic can be measured.

Comparison with Other Heuristics

This is one of the most commonly used methods for difficult problems which have been worked on for a long time and for which some good heuristics are known. Similarly to what happens with the bound comparisons, the conclusion of this comparison deals with the quality of fit of the chosen heuristic.

Given that the LOP has been studied in-depth from both the exact viewpoint and that of a heuristic, we have a value of the optimum solution for small and

medium-scale examples, which enables us to establish the optimal deviation in the solution obtained by the heuristics. Furthermore, we can compare the values obtained between the different heuristics to solve the same examples of any size.

Worst Case Analysis

One method that was well-accepted for a time concerns the behavioral analysis of the heuristic algorithm in the worst case; i.e., consider the examples that most disfavor the algorithm and set analytical bounds to the maximal deviation in terms of the optimum solution to the problem. The best aspect of this method is that it established the limits of the algorithm's results for any example. However, for the same reason, the results tend not to be representative of the average behavior of the algorithm. Furthermore, the analysis can be very complicated for more sophisticated heuristics.

An algorithm A for dealing with a maximization problem is called *ε-approximative* if there is a constant $\varepsilon > 0$ such that for every problem instance the algorithm guarantees that a feasible solution can be found with value c_A and the property

$$c_A \geq (1 - \varepsilon)c_{\text{opt}}.$$

The analogous definition for minimization problems is $c_A \leq (1 + \varepsilon)c_{\text{opt}}$.

Concerning the approximability of the LOP the following results are known. Suppose that all objective function coefficients are nonnegative and take some arbitrary ordering. Then either this ordering or its reverse version contains at least half of the sum of all coefficients. So $\frac{1}{2}$-approximation of the LOP is trivial, but nothing better is known.

2.2 Construction Heuristics

We will now review some of the construction heuristics, i.e., methods which follow some principle for successively constructing a linear ordering. The principle should somehow reflect that we are searching for an ordering with high value.

2.2.1 The Method of Chenery and Watanabe

One of the earliest heuristic methods was proposed by Chenery and Watanabe [32]. These authors did not formulate an algorithm, but just gave some ideas of how to obtain plausible rankings of the sectors of an input-output table. Their suggestion is to rank those sectors first which show a small share of inputs from other sectors and of outputs to final demand. Sectors having a large share of inputs from other

industries and of final demand output should be ranked last. Chenery and Watanabe defined coefficients taking these ideas into account to find a preliminary ranking. Then they try to improve this ranking in some heuristic way which is not specified in their paper. The authors admit that their method does not necessarily lead to good approximate solutions of the triangulation problem.

2.2.2 Heuristics of Aujac & Masson

This method [6] is based on so-called *output coefficients*. The output coefficient of a sector i with respect to another sector j is defined as

$$b_{ij} = \frac{c_{ij}}{\sum_{k \neq i} c_{ik}} .$$

Then it is intended to rank sector i before sector j whenever $b_{ij} > b_{ji}$ ("better customer principle"). This is impossible in general. So it is heuristically tried to find a linear ordering with few contradictions to this principle. Subsequently local changes are performed to achieve better triangulations. Similarly an *input coefficient* method can be formulated based on the input coefficients

$$a_{ij} = \frac{c_{ij}}{\sum_{k \neq j} c_{kj}} .$$

2.2.3 Heuristics of Becker

In [8] two further methods are described. The first one is related to the previous ones in that it calculates special quotients to rank the sectors. For each sector i the number

$$q_i = \frac{\sum_{k \neq i} c_{ik}}{\sum_{k \neq i} c_{ki}}$$

is determined. The sector with the largest quotient q_i is then ranked highest. Its corresponding rows and columns are deleted from the matrix, and the procedure is applied to the remaining sectors.

Heuristic of Becker (1)

(1) Set $S = \{1, 2, \ldots, n\}$.
(2) For $k = 1, 2, \ldots, n$:

    (2.1) For each $i \in S$ compute $q_i = \dfrac{\sum\limits_{j \in S \setminus \{i\}} c_{ij}}{\sum\limits_{j \in S \setminus \{i\}} c_{ji}}$.

    (2.2) Let $q_j = \max\{q_i \mid i \in S\}$.
    (2.3) Set $i_k = j$ and $S = S \setminus \{j\}$.

The second method starts with an arbitrarily chosen linear ordering, w.l.o.g. $\langle 1, 2, \ldots, n \rangle$. Then for every $m = 1, 2, \ldots, n-1$ the objective function values of the orderings $\langle m+1, m+2, \ldots, n, 1, \ldots, m \rangle$ are evaluated. The best one among them is chosen, and the procedure is repeated as long as improvements are possible.

Heuristic of Becker (2)

(1) Generate a random ordering.
(2) Let $\langle i_1, i_2, \ldots, i_n \rangle$ denote the current ordering.
(3) Evaluate all of the orderings $\langle i_{m+1}, i_{m+2}, \ldots, i_n, 1, 2, \ldots, i_m \rangle$, for $m = 1, 2, \ldots, n-1$.
(4) If the best one among these orderings is better than the current one, take it as the new current ordering and goto (3).

2.2.4 Best Insertion

This is a simple heuristic which builds an ordering by inserting the next objects at positions which are locally optimal.

Best Insertion

(1) Select an arbitrary object j and set $S = \{1, 2, \ldots, n\} \setminus \{j\}$. Let $\langle j \rangle$ be the current ordering.
(2) For $k = 1, 2, \ldots, n-1$:

    (2.1) Let $\langle i_1, i_2, \ldots, i_k \rangle$ denote the current ordering and choose some $l \in S$.
    (2.2) For every t, $1 \leq t \leq k+1$, compute $q_t = \sum_{j=1}^{t-1} c_{i_j l} + \sum_{j=t}^{k} c_{l i_j}$ and let $q_p = \max\{q_t \mid 1 \leq t \leq k\}$.
    (2.3) Insert l at position p in the current ordering and set $S = S \setminus \{l\}$.

An alternative version of step (2.2) computes

$$q_t = \sum_{j=1}^{t-1} c_{ijl} + \sum_{j=t}^{k} c_{li_j} - \sum_{j=1}^{t-1} c_{li_j} - \sum_{j=t}^{k} c_{ijl}$$

to account for the sum of entries which are "lost" when l is inserted at position t.

Table 2.1 Constructive methods on OPT-I instances

	CW	AM-O	AM-I	Bcq	Bcr	BI1	BI2
IO							
Dev(%)	19.07	32.94	31.45	4.07	30.19	3.24	4.18
Score	231	291	266	101	289	89	104
#Opt	0	0	0	0	0	0	0
SGB							
Dev(%)	12.83	26.15	26.15	3.57	31.56	3.89	3.03
Score	100	125	125	54	175	56	40
#Opt	0	0	0	0	0	0	0
RandomAII							
Dev(%)	2.60	36.50	36.55	1.57	37.75	1.09	1.26
Score	100	135	136	68	162	34	48
#Opt	0	0	0	0	0	0	0
RandomB							
Dev(%)	10.13	24.69	24.69	7.04	26.41	5.24	4.87
Score	276	368	368	194	454	124	106
#Opt	0	0	0	0	0	0	0
MB							
Dev(%)	8.40	43.37	43.37	2.90	40.30	2.49	2.27
Score	120	178	178	80	154	52	48
#Opt	0	0	0	0	0	0	0
Special							
Dev(%)	0.02	0.57	0.14	3.10	0.40	0.01	0.17
Score	64	178	113	210	149	41	83
#Opt	0	0	0	0	0	4	3
OPT-I							
Avg. Dev(%)	10.85	32.97	32.55	3.95	32.35	3.49	3.50
Sum #Opt	0	0	0	0	0	4	3

Table 2.1 reports on our results for 7 constructive heuristics on the OPT-I set (the set of 229 instances with optimum known). In this experiment we compute for each instance and each method the relative deviation *Dev* (in percent) between the best solution value *Value* obtained with the method and the optimal value for that instance. For each method, we also report the number of instances *#Opt* for which an optimum solution could be found. In addition, we calculate the so-called

score statistic [114] associated with each method. For each instance, the *nrank* of method *M* is defined as the number of methods that produce a better solution than the one found by *M*. In the event of ties, the methods receive the same *nrank*, equal to the number of methods strictly better than all of them. The value of *Score* is the sum of the *nrank* values for all the instances in the experiment. Thus the lower the *Score* the better the method. We do not report running times in this table because these methods are very fast and their running times are extremely short (below 1 millisecond). Specifically, Table 2.1 shows results for:

– `CW`: Chenery and Watanabe algorithm
– `AM-O`: Aujac and Masson algorithm (output coefficients)
– `AM-I`: Aujac and Masson algorithm (input coefficients)
– `Bcq`: Becker algorithm (based on quotients)
– `Bcr`: Becker algorithm (based on rotations)
– `BI1`: Best Insertion algorithm (variant 1)
– `BI2`: Best Insertion algorithm (variant 2)

Results in Table 2.1 clearly indicate that OPT-I instances pose a challenge for the simple heuristics with average percentage deviations ranging from 3.49% to 32.97%. In most of the cases none of the methods is able to match the optimum solution (with the exception of `BI1` and `BI2` with 4 and 3 optima respectively in the Special instances). These results show that only `Bcq`, `BI1` and `BI2` can be considered reasonable construction heuristics (with an average percent deviation lower than 5%).

2.3 Local Search

After having constructed some ordering with one of the heuristics above, it is reasonable to look for improvement possibilities. In this section we will describe fairly simple (deterministic) local improvement methods that are able to produce acceptable solutions for the LOP. The basic philosophy that drives local search is that it is often possible to find a good solution by repeatedly increasing the quality of a given solution, making small changes at a time called *moves*. The different types of possible moves characterize the various heuristics. Starting from a solution generated by a construction heuristic, a typical local search performs steps as long as the objective function increases.

Local search can only be expected to obtain optimum or near-optimum solutions for easy problems of medium size, but it is a very important and powerful concept for the design of meta-heuristics, which are the topic of the next chapter.

2.3.1 Insertion

This heuristic checks whether the objective function can be improved if the position of an object in the current ordering is changed. All possibilities for altering the position of an object are checked and the method stops when no further improvement is possible this way.

In problems where solutions are represented as permutations, insertions are probably the most direct and efficient way to modify a solution. Note that other movements, such as swaps, can be obtained by composition of two or more insertions. We define $move(O_j, i)$ as the modification which deletes O_j from its current position j in permutation O and inserts it at position i (i.e., between the objects currently in positions $i - 1$ and i).

Now, the insertion heuristic tries to find improving moves examining eventually all possible new positions for all objects O_j in the current permutation O. There are several ways for organizing the search for improving moves. For our experiments we proceeded as follows:

Insertion

(1) Compute an initial permutation $O = \langle O_1, O_2, \ldots, O_n \rangle$.
(2) For $j = 1, 2, \ldots, n$:

    (2.1) Evaluate all possible insertions $move(O_j, i)$.
    (2.2) Let $move(O_k, i^*)$ be the best of these moves.
    (2.3) If $move(O_k, i^*)$ is improving then perform it and update O.

(3) If some improving move was found, then goto (2).

In [86] two neighborhoods are studied in the context of local search methods for the LOP. The first one consists of permutations obtained by switching the positions of contiguous objects O_j and O_{j+1}. The second one involves all permutations resulting from executing general insertion moves, as defined above. The conclusion from the experiments is that the second neighborhood clearly outperforms the first one, which is much more limited. Furthermore two strategies for exploring the neighborhood of a solution were studied. The *best* strategy selects the move with the largest *move value* among all the moves in the neighborhood. The *first* strategy, on the other hand, scans the list of objects (in the order given by the current permutation) searching for the first object whose movement gives a strictly positive move value. The computations revealed that both strategies provide similar results but the *first* involved lower running times.

2.3.2 The Heuristic of Chanas & Kobylanski

The method developed by Chanas and Kobylanski [32], referred to as the CK method in the following, is based on the following symmetry property of the LOP. If the permutation $O = \langle O_1, O_2, \ldots, O_n \rangle$ is an optimum solution to the maximization problem, then an optimum solution to the minimization problem is $O^* = \langle O_n, O_{n-1}, \ldots, O_1 \rangle$. In other words, when the sum of the elements above the main diagonal is maximized, the sum of the elements below the diagonal is minimized. The CK method utilizes this property to escape local optimality. In particular, once a local optimum solution O is found, the process is re-started from the permutation O^*. This is called the REVERSE operation.

In a global iteration, the CK method performs insertions as long as the solution improves. Given a solution, the algorithm explores the insertion move $move(O_j, i)$ of each element O_j in all the positions i in O, and performs the best one. When no further improvement is possible, it generates a new solution by applying the REVERSE operation from the last solution obtained, and performs a new global iteration. The method finishes when the best solution found cannot be improved upon in the current global iteration.

It should be noted that the CK method can be considered to be a generalization of the second heuristic of Becker described above. The latter evaluates the orderings that can be obtained by rotations of a solution, while the CK method evaluates all insertions. Since these rotations are basically insertions of the first elements to the last positions, we can conclude that Becker's method explores only a fraction of the solutions explored by CK.

2.3.3 k-opt

The k-opt improvement follows a principle that can be applied to many combinatorial optimization problems. Basically, it selects k elements of a solution and locally optimizes with respect to these elements. For the LOP, a possible k-opt heuristic would be to consider all subsets of k objects $O_{i_1}, \ldots, O_{i_k}$ in the current permutation and find the best assignment of these objects to the positions $i_1, \ldots, i_k$. Since the number of possible new assignments grows exponentially with k, we have only implemented 2-opt and 3-opt.

2.3.4 Kernighan-Lin Type Improvement

The main problem with local improvement heuristics is that they very quickly get trapped in a local optimum. Kernighan and Lin proposed the idea (originally in [78] for a partitioning problem) of looking for more complicated moves that are composed of simpler moves. In contrast to pure improvement heuristics, it allows that

some of the simple moves are not improving. In this way the objective can decrease locally, but new possibilities arise for escaping from the local optimum. This type of heuristic proved particularly effective for the traveling salesman problem (where it is usually named *Lin-Kernighan heuristic*).

We only describe the principle of the Kernighan-Lin approach. For practical applications on large problems, it has to be implemented carefully with appropriate data structures and further enhancements like restricted search or limited length of combined moves to speed up the search for improving moves. We do not elaborate on this here.

We devised two heuristics of this type for the LOP. In the first version, the basic move consists of interchanging two objects in the current permutation.

Kernighan-Lin 1

(1) Compute some linear ordering O.
(2) Let $m = 1$, $S_m = \{1, 2, \ldots, n\}$.
(3) Determine objects $s, t \in S_m$, $s \neq t$, the interchange of which in the current ordering leads to the largest increase g_m of the objective function (increase may be negative).
(4) Interchange s and t in the current ordering. Set $s_m = s$ and $t_m = t$.
(5) If $m < \lfloor n/2 \rfloor$, set $S_{m+1} = S_m \setminus \{s, t\}$ and $m = m + 1$. Goto (3).
(6) Determine $1 \leq k \leq m$, such $G = \sum_{i=1}^{k} g_i$ is maximum.
(7) If $G \leq 0$ then Stop, otherwise, starting from the original ordering O, successively interchange s_i and t_i, for $i = 1, 2, \ldots, k$. Let O denote the new ordering and goto (2).

The second version builds upon insertion moves.

Kernighan-Lin 2

(1) Compute some linear ordering O.
(2) Let $m = 1$, $S_m = \{1, 2, \ldots, n\}$.
(3) Among all possibilities for inserting an object of S_m at a new position determine the one leading to the largest increase g_p of the objective function (increase may be negative). Let s be this object and p the new position.
(4) Move s to position p in the current ordering. Set $s_m = s$ and $p_m = p$.
(5) If $m < n$, set $S_{m+1} = S_m \setminus \{s\}$ and $m = m + 1$. Goto (3).
(6) Determine $1 \leq k \leq m$, such $G = \sum_{i=1}^{k} g_i$ is maximum.
(7) If $G \leq 0$ then Stop, otherwise, starting from the original ordering O, successively move s_i to position p_i, for $i = 1, 2, \ldots, k$. Let O denote the new ordering and goto (2).

2.3.5 Local Enumeration

This heuristic chooses windows $\langle i_k, i_{k+1}, \ldots, i_{k+L-1} \rangle$ of a given length L of the current ordering $\langle i_1, i_2, \ldots, i_n \rangle$ and determines the optimum subsequence of the respective objects by enumerating all possible orderings. The window is moved along the complete sequence until no more improvements can be found. Of course, L cannot be chosen too large because the enumeration needs time $O(L!)$.

Local Enumeration

(1) Compute some linear ordering O.

(2) For $i = 1, \ldots, n - L + 1$:

  (2.1) Find the best possible rearrangement of the objects at positions $i, i+1, \ldots, i+L-1$.

(3) If an improving move has been found in the previous loop, then goto (2).

Table 2.2 reports on our results for 7 improving heuristics on the OPT-I set of instances. As in the construction heuristics, we report, for each instance and each method, the relative percent deviation *Dev*, the number of instances *#Opt* for which an optimum solution is found, and the *score* statistic. Similarly, we do not report running times in this table because these methods are fairly fast. Specifically, the results obtained with the following improvement methods (started with a random initial solution) are given:

- `LSi`: Local Search based on insertions
- `2opt`: Local Search based on 2-opt
- `3opt`: Local Search based on 3-opt
- `LSe`: Local Search based on exchanges
- `KL1`: Kernighan-Lin based on exchanges
- `KL2`: Kernighan-Lin based on insertions
- `LE`: Local enumeration

As expected, the improvement methods are able to obtain better solutions than the construction heuristics, with average percentage deviations (shown in Table 2.2) ranging from 0.57% to 2.30% (the average percentage deviations of the construction heuristics range from 3.49% to 32.97% as reported in Table 2.1). We have not observed significant differences when applying the improvement method from different initial solutions. For example, as shown in Table 2.2 the LSi method exhibits a *Dev* value of 0.16% on the RandomAII instances when it is started from random solutions. When it is run from the CW or the Bcr solutions, it obtains a *Dev* value of 0.17% and 0.18% respectively.

Table 2.2 Improvement methods on OPT-I instances

	LSi	2opt	3opt	LSe	KL1	KL2	LE
IO							
Dev(%)	1.08	0.64	0.23	1.73	1.35	4.24	0.01
Score	243	181	125	295	239	232	49
#Opt	0	1	4	0	1	0	43
SGB							
Dev(%)	0.16	0.81	0.53	1.35	0.63	0.28	1.09
Score	42	122	84	154	100	63	135
#Opt	1	0	0	0	0		0
RandomAII							
Dev(%)	0.16	0.77	0.38	0.62	0.61	0.09	0.54
Score	46	161	81	134	134	29	112
#Opt	0	0	0	0	0	0	0
RandomB							
Dev(%)	0.79	4.04	2.13	3.78	3.51	0.61	3.56
Score	124	400	232	387	359	95	362
#Opt	1	0	0	0	0	1	0
MB							
Dev(%)	0.02	0.57	0.14	3.10	0.40	0.01	0.17
Score	64	178	113	210	149	41	83
#Opt	0	0	0	0	0	4	3
Special							
Dev(%)	1.19	3.30	2.05	3.21	2.40	0.89	3.52
Score	69	144	82	138	120	49	156
#Opt	4	2	2	2	3	3	3
OPT-I							
Avg. Dev(%)	0.57	1.69	0.91	2.30	1.49	1.02	1.48
Sum #Opt	5	3	6	2	4	8	49

2.4 Multi-Start Procedures

Multi-start procedures were originally conceived as a way of exploiting a local or neighborhood search procedure, by simply applying it from multiple random initial solutions. It is well known that search methods based on local optimization, aspiring to find global optima, usually require certain diversification to overcome local optimality. Without this diversification, such methods can become reduced to tracing paths that are confined to a small area of the solution space, making it impossible to find a global optimum. *Multi-start algorithms* can be considered to be a bridge between simple (classical) heuristics and complex (modern) meta-heuristics. The *re-start mechanism* of multi-start methods can be super-imposed on many different search methods. Once a new solution has been generated, a variety of options can be used to improve it, ranging from a simple greedy routine to a complex meta-heuristic. This section focuses on the different strategies and methods that can be

used to generate solutions to launch a succession of new searches for a global optimum.

The principle layout of a multi-start procedure is the following.

Multi-Start

(1) Set $i=1$.

(2) While the stopping condition is not satisfied:

    (2.1) Construct a solution x_i. **(Generation)**

    (2.2) Apply local search to improve x_i and let x_i' be the solution obtained. **(Improvement)**

    (2.3) If x_i' improves the best solution, update it. Set $i = i+1$. **(Test)**

The computation of x_i in (2.1) is typically performed with a constructive algorithm. Step (2.2) tries to improve this solution, obtaining x_i'. Here, a simple improvement method can be applied. However, this second phase has recently become more elaborate and, in some cases, is performed with a complex meta-heuristic that may or may not improve the initial solution x_i (in this latter case we set $x_i' = x_i$).

2.4.1 Variants of Multi-Start

We will first review some relevant contributions on multi-start procedures.

Early papers on multi-start methods are devoted to the Monte Carlo random re-start in the context of nonlinear unconstrained optimization, where the method simply evaluates the objective function at randomly generated points. The probability of success approaches 1 as the sample size tends to infinity under very mild assumptions about the objective function. Many algorithms have been proposed that combine the Monte Carlo method with local search procedures [115]. The convergence for random re-start methods is studied in [120], where the probability distribution used to choose the next starting point can depend on how the search evolves. Some extensions of these methods seek to reduce the number of complete local searches that are performed and increase the probability that they start from points close to the global optimum [96].

In [13] relationships among local minima from the perspective of the best local minimum are analyzed, finding convex structures in the cost surfaces. Based on the results of that study, they propose a multi-start method where starting points for greedy descent are adaptively derived from the best previously found local minima. In the first step, *adaptive multi-start heuristics* generate random starting solutions and run a greedy descent method from each one to determine a set of corresponding random local minima. In the second step, *adaptive starting solutions* are constructed based on the local minima obtained so far and improved with a greedy descent method. This improvement is applied several times from each adaptive starting

solution to yield corresponding *adaptive local minima*. The authors test this method for the traveling salesman problem and obtain significant speedups over previous multi-start implementations.

Simple forms of multi-start methods are often used to compare other methods and measure their relative contribution. In [7] different genetic algorithms for six sets of benchmark problems commonly found in the genetic algorithms literature are compared: traveling salesman problem, job-shop scheduling, knapsack and bin packing problem, neural network weight optimization, and numerical function optimization. The author uses the *multi-start method* (*multiple restart stochastic hill-climbing*) as a basis for computational testing. Since solutions are represented with strings, the improvement step consists of a local search based on random flipping of bits. The results indicate that using genetic algorithms for the optimization of static functions does not yield a benefit, in terms of the final answer obtained, over simpler optimization heuristics.

One of the most well known multi-start methods is the *greedy adaptive search procedure (GRASP)*. The GRASP methodology was introduced by Feo and Resende [45] and was first used to solve set covering problems [44]. We will devote a section in the next chapter to describe this methodology in detail.

A multi-start algorithm for unconstrained global optimization based on *quasi-random samples* is presented in [67]. Quasi-random samples are sets of deterministic points, as opposed to random, that are evenly distributed over a set. The algorithm applies an inexpensive local search (steepest descent) on a set of quasi-random points to concentrate the sample. The sample is reduced, replacing worse points with new quasi-random points. Any point that is retained for a certain number of iterations is used to start an efficient complete local search. The algorithm terminates when no new local minimum is found after several iterations. An experimental comparison shows that the method performs favorably with respect to other global optimization procedures.

An open question in order to design a good search procedure is whether it is better to implement a simple improving method that allows a great number of global iterations or, alternatively, to apply a complex routine that significantly improves a few generated solutions. A simple procedure depends heavily on the initial solution but a more elaborate method takes much more running time and therefore can only be applied a few times, thus reducing the sampling of the solution space. Some meta-heuristics, such as GRASP, launch limited local searches from numerous constructions (i.e., starting points). In other methods, such as tabu search, the search starts from one initial point and, if a restarting procedure is also part of the method, it is invoked only a limited number of times. In [94] the balance between restarting and search-depth (i.e., the time spent searching from a single starting point) is studied in the context of the matrix bandwidth problem. Both alternatives were tested with the conclusion that it was better to invest the time searching from a few starting points than re-starting the search more often. Although we cannot draw a general conclusion from these experiments, the experience in the current context and in previous projects indicates that some meta-heuristics, like tabu search, need

to reach a critical search depth to be effective. If this search depth is not reached, the effectiveness of the method is severely compromised.

2.4.2 Experiments with the LOP

In this section we will describe and compare 10 different constructive methods for the LOP. It should be noted that, if a constructive method is completely deterministic (with no random elements), its replication (running it several times) will always produce the same solution. Therefore, we should add random selections in a constructive method to obtain different solutions when replicated. Alternatively, we can modify selections from one construction to another in a deterministic way by recording and using some frequency information. We will look at both approaches, which will enable us to design constructive methods for the LOP that can be embedded in a multi-start procedure.

Above we have described the construction heuristic of Becker [8] in which for each object i the value q_i is computed. Then, the objects are ranked according to the q-values $q_i = \sum_{k \neq i} c_{ik} / \sum_{k \neq i} c_{ki}$.

We now compute two other values that can also be used to measure the attractiveness of an object to be ranked first. Specifically, r_i and c_i are, respectively, the sum of the elements in the row corresponding to object i, and the sum of the elements in the column of object i, i.e., $r_i = \sum_{k \neq i} c_{ik}$ and $c_i = \sum_{k \neq i} c_{ki}$.

Constructive Method G1

This method first computes the r_i values for all objects. Then, instead of selecting the object with the largest r-value, it creates a list with the most attractive objects, according to the r-values, and randomly selects one among them. The selected object is placed first and the process is repeated for n iterations. At each iteration the r-values are updated to reflect previous selections (i.e., we sum the c_{ik} across the unselected elements) and the candidate list for selection is computed with the highest evaluated objects. The method combines the *random selection* with the *greedy evaluation*, and the size of the candidate list determines the relative contribution of these two elements.

Constructive method G1

(1) Set $S = \{1, 2, \ldots, n\}$. Let $\alpha \in [0, 1]$ be the percentage for selection and O be the empty ordering.

(2) For $t = 1, 2, \ldots, n$:

    (2.1) Compute $r_i = \sum\limits_{k \in S, k \neq i} c_{ik}$ for all $i \in S$.

    (2.2) Let $r^* = \max\{r_i \mid i \in S\}$.

    (2.3) Compute the candidate list $C = \{i \in S \mid r_i \geq \alpha r^*\}$.

    (2.4) Randomly select $j^* \in C$ and place j^* at position t in O and set $S = S \setminus \{j^*\}$.

Constructive Methods G2 and G3

Method G2 is based on the c_i-values computed above. It works in the same way as G1 but the attractiveness of object i is now measured with c_i instead of r_i. Objects with large c-values are placed now in the last positions.

Constructive method G2

(1) Set $S = \{1, 2, \ldots, n\}$. Let $\alpha \in [0, 1]$ be the percentage for selection and O be the empty ordering.

(2) For $t = n, n - 1, \ldots, 1$:

    (2.1) Compute $c_i = \sum\limits_{k \in S, k \neq i} c_{ki}$ for all $i \in S$.

    (2.2) Let $c^* = \max\{c_i \mid i \in S\}$.

    (2.3) Compute the candidate list $C = \{i \in S \mid c_i \geq \alpha c^*\}$.

    (2.4) Randomly select $j^* \in C$ and place j^* in position t in O and set $S = S \setminus \{j^*\}$.

In a similar way, constructive method G3 measures the attractiveness of object i for selection with q_i and performs the same steps as G1. Specifically, at each iteration the q-values are computed with respect to the unselected objects, a restricted candidate list is formed with the objects with largest q-values, and one of them is randomly selected and placed first.

Constructive Methods G4, G5 and G6

These methods are designed analogously to G1–G3, except that the selection of objects is from a candidate list of the least attractive and the solution is constructed starting from the last position of the permutation. We give the specification of G6 which is modification of G3.

Constructive method G6

(1) Set $S = \{1, 2, \ldots, n\}$. Let $\alpha \geq 0$ be the percentage for selection and O be the empty ordering.

(2) For $t = 1, 2, \ldots, n$:

  (2.1) For all $i \in S$, compute

$$q_i = \frac{\sum\limits_{k \in S, k \neq i} c_{ik}}{\sum\limits_{k \in S, k \neq i} c_{ki}}.$$

  (2.2) Let $q^* = \min\{q_i \mid i \in S\}$.

  (2.3) Compute the candidate list $C = \{i \in S \mid q_i \leq (1 + \alpha)q^*\}$.

  (2.4) Randomly select $j^* \in C$ and place j^* in position $n - t + 1$ in O and set $S = S \setminus \{j^*\}$.

Constructive Method MIX

This is a mixed procedure derived from the previous six. The procedure generates a fraction of solutions from each of the previous six methods and combines these solutions into a single set. That is, if n solutions are required, then each method Gi, $i = 1, \ldots, 6$, contributes $n/6$ solutions.

Constructive Method RND

This is a random generator. This method simply generates random permutations. We use it as a basis for our comparisons.

Constructive Method DG

This is a general purpose diversification generator suggested in [55] which generates diversified permutations in a systematic way without reference to the objective function.

Constructive Method FQ

This method implements an algorithm with frequency-based memory, as proposed in tabu search [52] (we will see this methodology in the next chapter). It is based on modifying a measure of attractiveness with a frequency measure that discourages

objects from occupying positions that they have frequently occupied in previous
solution generations.

The constructive method FQ (proposed in [19]) is based on the notion of
constructing solutions employing modified *frequenciesfrequency*. The generator
exploits the permutation structure of a linear ordering. A frequency counter is
maintained to record the number of times an element i appears in position j. The
frequency counters are used to penalize the "attractiveness" of an element with re-
spect to a given position. To illustrate this, suppose that the generator has created
30 solutions. If 20 out of the 30 solutions have element 3 in position 5, then the
frequency counter $freq(3,5) = 20$. This frequency value is used to bias the potential
assignment of element 3 in position 5 during subsequent constructions, thus induc-
ing *diversification* with respect to the solutions already generated.

The attractiveness of assigning object i to position j is given by the greedy func-
tion $fq(i,j)$, which modifies the value of q_i to reflect previous assignments of object i
to position j, as follows:

$$fq(i,j) = \frac{\sum\limits_{k \neq i} c_{ik}}{\sum\limits_{k \neq i} c_{ki}} - \beta \frac{\max_q}{\max_f} freq(i,j),$$

where $\max_f = \max\{freq(i,j) \mid i = 1,\ldots,n, j = 1,\ldots,n\}$ and $\max_q = \max\{q_i \mid i = 1,\ldots,n\}$.

Constructive method FQ

(1) Set $S = \{1,2,\ldots,n\}$. Let $\beta \in [0,1]$ be the percentage for diversifica-
tion and $freq(i,j)$ be the number of times object i has been assigned
to position j in previous constructions.

(2) For $t = 1,2,\ldots,n$:

(2.1) For all $i,j \in S$ compute $fq(i,j) = \frac{\sum\limits_{k \neq i} c_{ik}}{\sum\limits_{k \neq i} c_{ki}} - \beta \frac{\max_q}{\max_f} freq(i,j)$.

(2.2) Let i^* and j^* be such that $fq(i^*,j^*) = \max\{fq(i,j) \mid i,j \in S\}$.

(2.3) Place i^* at position j^* in O and set $S = S \setminus \{i^*\}$.

(2.4) $freq(i^*,j^*) = freq(i^*,j^*) + 1$.

It is important to point out that $fq(i,j)$ is an adaptive function since its value
depends on attributes of the unassigned elements at each iteration of the construction
procedure.

In our first experiment we use the instance `stabu75` from `LOLIB`. We have
generated a set of 100 solutions with each of the 10 generation methods. Figure 2.1
shows in a box-and-whisker-plot representation, the value of the 100 solutions gen-
erated with each method. Since the LOP is a maximization problem, it is clear that
the higher the value, the better the method. We can therefore say that constructive
method G3 obtains the best results. Other methods, such as FQ and MIX also obtain

Fig. 2.1 Objective function value box-plot for each method

solutions with very good values, but their box-plot representation indicates that they also produce lower quality solutions. However, if the construction is part of a global method (as is the case in multi-start methods), we may prefer a constructive method able to obtain solutions with different structures rather than a constructive method that provides very similar solutions. Note that if every solution is subjected to local search, then it is preferable to generate solutions scattered in the search space as starting points for the local search phase rather than good solutions concentrated in the same area of the solution space. Therefore, we need to establish a trade off between *quality* and *diversity* when selecting our construction method.

Given a set of solutions P represented as permutations, in [95] a diversity measure d is proposed which consists of computing the distances between each solution and a "center" of P. The sum (or alternatively the average) of these $|P|$ distances provides a measure of the diversity of P. The diversity measure d is calculated as follows:

(1) Calculate the median position of each element i in the solutions in P.
(2) Calculate the *dissimilarity* (distance) of each solution in the population with respect to the median solution. The dissimilarity is calculated as the sum of the absolute difference between the position of the elements in the solution under consideration and the median solution.
(3) Calculate d as the sum of all the individual dissimilarities.

For example, assume that P consists of the orderings $\langle A,B,C,D \rangle$, $\langle B,D,C,A \rangle$, and $\langle C,B,A,D \rangle$. The median position of element A is therefore 3, since it occupies positions 1, 3 and 4 in the given orderings. In the same way, the median positions of B,C and D are 2, 3 and 4, respectively. Note that the median positions might not

induce an ordering, as in the case of this example. The diversity value of the first
solution is then calculated as $d_1 = |1 - 3| + |2 - 2| + |3 - 3| + |4 - 4| = 2$.

In the same way, the diversity values of the other two solutions are obtained as
$d_2 = 4$ and $d_3 = 2$. The diversity measure d of P is then given by $d = 2 + 4 + 2 = 8$.

We then continue with our experiment to compare the different constructive
methods for the LOP. As described above, we have generated a set of 100 solutions
with each of the 10 generation methods. Figure 2.2 shows the box-and-whisker plot
of the diversity values of the solution set obtained with each method.

Fig. 2.2 Diversity value box-plot for each method

Figure 2.2 shows that MIX and FQ obtain the highest diversity values (but also
generate other solutions with low diversity values). As expected, the random con-
structive method RND consistently produces high diversity values (always generat-
ing solutions with an associated d-value over 800 in the diagram).

As mentioned, a good method must produce a set of solutions with high quality
and high diversity. If we compare, for example, generators MIX and G3 we observe
in Fig. 2.1 that G3 produces slightly better solutions in terms of solution quality, but
Fig. 2.2 shows that MIX outperforms G3 in terms of diversity. Therefore, we will
probably consider MIX as a better method than G3. In order to rank the methods we
have computed the average of both measures across each set.

Figure 2.3 shows the average of the diversity values on the x-axis and the average
of the quality on the y-axis. A point is plotted for each method.

As expected, the random generator RND produces a high diversity value (as mea-
sured by the dissimilarity) but a low quality value. DG matches the diversity of RND
using a systematic approach instead of randomness, but as it does not use the value
of solutions, it also presents a low quality score. The mixed method MIX provides a

Fig. 2.3 Quality and diversity for each method

good balance between dissimilarity and quality, by uniting solutions generated with methods G1 to G6.

We think that quality and diversity are equally important, so we have added both averages. To do so, we use two relative measures $\triangle C$ for quality, and $\triangle d$ for diversity. They are basically standardizations to translate the average of the objective function values and diversity values respectively to the [0,1] interval. In this way we can simply add both quantities.

Fig. 2.4 Quality and Diversity for each method

Figure 2.4 clearly shows the following ranking of the 10 methods, where the overall best is the FQ generator: G5, G4, G2, G1, DG, RND, G6, G3, MIX and FQ. These results are in line with previous works which show the inclusion of memory

structures (frequency information) to be effective within the multi-start framework. However, one should note that this method of ranking has been obtained considering both quality and diversity with equal weight. If we vary this criterion, the ranking would also change.

It should be noted that, unlike other well-known methods that we will review in the next chapter, multi-start procedures have not yet become widely implemented and tested as a meta-heuristic itself for solving complex optimization problems. We have shown new ideas that have recently emerged within the multi-start area that add a clear potential to this framework which has yet to be fully explored.

Chapter 3
Meta-Heuristics

Abstract In this chapter we elaborate on meta-heuristics for optimization from a beginner's perspective. Basically, we start from scratch to describe the different methodologies and provide the reader with the elements and strategies to build and implement them successfully. Although we describe the adaptation of these methods to the linear ordering problem, we do not limit our descriptions to this problem. Instead, we present all methods in their generic form and then adapt them to the specific case of the LOP. This way, the reader can easily apply these methods to a wide range of combinatorial optimization problems.

3.1 Introduction

In the last decade a series of methods have appeared under the name of *meta-heuristics*, which aim to obtain better results than those obtained with "traditional heuristics". The term meta-heuristic was coined by Glover [54] in 1986. In this book we shall use the term heuristics to refer to the classical methods as opposed to meta-heuristics, which we shall reserve to refer to the most recent and complex ones. In some texts one can find the expression *modern heuristics* to refer to meta-heuristics [110]. Osman and Kelly [106] introduce the following definition:

> *"A meta-heuristic is an iterative generation process which guides a subordinate heuristic by combining intelligently different concepts for exploring and exploiting the search spaces using learning strategies to structure information in order to find near-optimal solutions."*

Meta-heuristic procedures are conceptually ranked "above" heuristics in the sense that they guide their design. Thus, facing an optimization problem, we can employ any of these procedures to design a specific algorithm for computing an approximate solution.

Heuristic methods have been developed to such an extent, with new methodologies appearing "almost daily", that to offer an exhaustive survey of all of them lies far outside of the the scope of this monograph. Furthermore, our prime interest lies

in those that have been used successfully for the LOP. We can list as better established ones the following approaches:

– *Estimation distribution algorithms*
– *Evolutionary algorithms*
– *Fuzzy adaptive neighborhood search*
– *Genetic algorithms*
– *Greedy randomized adaptive search procedure*
– *Guided local search*
– *Heuristic concentration*
– *Memetic algorithms*
– *Multi-start methods*
– *Path relinking*
– *Scatter search*
– *Simulated annealing*
– *Tabu search*
– *Variable neighborhood search*

It is important to highlight two aspects. The first one is that the above list is not exhaustive; as we mentioned previously this would be very difficult to establish. The second one is that some methods can be considered specializations or adaptations of other more generic ones (a consideration that is not free from controversy). For example, "ant colony optimization" is motivated by the trails left by ants in search for optimum paths. In our view, this could as well be interpreted as a short term memory structure in tabu search. Researchers in this field increasingly propose combinations and hybridizations among the different methods, which makes the boundary between them less well defined, in many cases taking some ingredients from others.

All the meta-heuristic methodologies have many degrees of freedom and the user must take several decisions in order to design the final algorithm. Hertz and Widmer [66] identified two categories on which meta-heuristics are based: *local search* and *population search*. In the former the method iterates over a single solution, while in the later the method iterates over a set of solutions. TS, SA or VNS are local search based methods, while SS or GA are population search based methods (in which a set of solutions evolves during the search process). In their study the authors provide the following guidelines to help designing a good method for a given problem within each category.

Local search

– It should be easy to generate solutions in the search space under consideration.
– The solutions in the neighborhood of a solution should be close to this solution in some way.
– The topology induced by the objective function in the neighborhood space should not be too "flat".
– Each solution in the search space should be linked with the optimum solution by a (short) sequence of moves.

Population search

–   Pertinent information should be transmitted during the combination phase (in which new solutions are obtained with the combination of old ones).
–   The combination of two equivalent *parent solutions* should not produce a new solution that is different from the parents.
–   Diversity should be preserved in the population.

In this chapter we are going to describe the application of the main meta-heuristic procedures to the LOP. We will consider only those procedures that are relatively consolidated and that have proved efficient for a representative collection of problems. Specifically, we shall describe the adaptation of GRASP, TS, SA, VNS, SS, GA and MA to the LOP.

3.2 GRASP

The GRASP methodology was developed in the late 1980s, and this acronym for *greedy randomized adaptive search procedures* was coined by Feo and Resende [45]. It was first used in [44] to solve computationally difficult set covering problems. Each GRASP iteration consists of constructing a trial solution and then applying an improvement procedure to find a local optimum (i.e., the final solution for that iteration). The construction phase is *iterative*, *greedy* and *adaptive*. It is iterative because the initial solution is built considering one element at a time. It is greedy because the addition of each element is guided by a greedy function. It is adaptive because the element chosen at any iteration in a construction is a function of those chosen previously and thus relevant information is updated from one construction step to the next. The improvement phase typically entails a local search procedure. The GRASP method is similar in layout to the general multi-start methods described in the previous chapter.

GRASP

(1) Set $i=1$.
(2) While $i < MaxIter$:

   (2.1) Construct a solution x_i. (**Construction phase**)
   (2.2) Apply a local search method to improve x_i. Let x_i' be the solution obtained. (**Improvement phase**)
   (2.3) If x_i' improves the best solution found, update it. (**Test**)
   (2.4) $i = i + 1$.

We now describe the two phases in more detail.

3.2.1 Construction Phase

At each iteration of the construction phase, GRASP maintains a candidate list CL of elements which can be added to the partial solution under construction to obtain a feasible complete solution. All candidate elements are evaluated according to a greedy function in order to select the next element to be added in the construction. The greedy function usually represents the marginal increase in the cost function by adding the element to the partial solution. Element evaluation is used to create a restricted candidate list RCL with the best elements in CL, i.e., those with the largest incremental cost in a maximization problem (as in the LOP). The element to be added to the partial solution is randomly selected from those in the RCL. Once the selected element is added to the partial solution, the candidate list is updated and the evaluations (incremental costs) are recalculated.

A particularly appealing characteristic of GRASP is its ease of implementation. We only need to define a construction mechanism according to the description above and a local search procedure. Moreover, the construction usually has one parameter related to the quality of the elements in the restricted candidate list. GRASP typically performs a pre-established number of iterations (construction + improvement) and returns the best solution found overall.

The first step in designing a GRASP construction is to define a greedy function in order to evaluate the relative contribution of adding an element to the partial solution under construction. In [20] the following three evaluators are proposed for the LOP, inspired by the simple heuristics described in the previous chapter:

- The evaluation $e_1(i)$ of object i is the sum of the elements in the matrix in its corresponding row:

$$e_1(i) = \sum_{j=1}^{n} c_{ij}.$$

- The evaluation $e_2(i)$ of object i is the difference between the maximum of the sum of columns, *colmax*, and the sum of the elements in the matrix in its corresponding column:

$$e_2(i) = colmax - \sum_{j=1}^{n} c_{ji}.$$

- The evaluation $e_3(i)$ of object i is the sum of the elements in the matrix in its corresponding row divided by the sum of the elements in its corresponding column:

$$e_3(i) = \frac{\sum_{j=1}^{n} c_{ij}}{\sum_{j=1}^{n} c_{ji}}.$$

The goal of the constructive method is to obtain a permutation of size n, so in each iteration we select an object from an RCL to be placed in the next available position. The permutations are constructed from left to right, meaning that we start from position 1 and end at position n. Initially, all objects are in the unassigned

list U. For each object $i \in U$, we calculate its associated evaluation $e_1(i)$ (or alternatively $e_2(i)$ or $e_3(i)$) with respect to the unassigned objects (i.e., we only compute in the expressions above the c-values associated to unassigned objects). This measures the attractiveness of each object. The largest attractiveness value e_1^* of all the unassigned objects is multiplied by a parameter α, $0 \leq \alpha \leq 1$. This value represents a threshold that is used to build the RCL. In particular, the RCL comprises all the objects in U with an attractiveness measure that is at least as large as the threshold value and we set RCL $= \{i \in U \mid e_1(i) \geq \alpha e_1^*\}$, where $e_1^* = \max\{e_1(i) \mid i \in U\}$.

Then, the next object to be assigned is randomly selected from the RCL. After the assignment has been made, the list of unassigned objects is updated and the measure of attractiveness for the objects in the updated set U is recalculated. The process is repeated n times and the outcome is a feasible solution of the LOP.

The example in Figure 3.1 shows the value of 100 solutions generated with the construction described above. We apply the evaluator e_1 and create the RCL with α set to 0.4 to solve the LOLIB instance be75eec. The best solution found in these 100 iterations has value 199072. Note that GRASP performs an independent sampling of the solution space, and therefore there is no connection between different constructions. In other words, we cannot predict the value of a solution by considering the value of the previously generated solutions. This is why we cannot find any pattern in Figure 3.1.

Fig. 3.1 Value of GRASP constructions

The design of different evaluation functions for the GRASP construction phase leads us to an interesting point: How can we compare different constructive methods? This is not a specific question about GRASP, but it is a general question about constructive methods. Obviously we want to obtain good solutions, so the larger the value of the solutions obtained (in a maximization problem), the better the constructive method. However, as discussed in the previous chapter in the context of general

multi-start methods, we need to establish a trade-off between quality and diversity when selecting our construction method.

Let GC_1, GC_2 and GC_3 be the three GRASP constructive methods obtained with the greedy functions e_1, e_2 and e_3 respectively. In order to compare their relative performance when solving the LOLIB instance be75eec, we generate 100 solutions with each method and then compute a measure of both the quality and diversity of all 100 solutions. To evaluate quality we can simply compute the average objective function value of the 100 solutions generated with each method. To evaluate diversity we can compute the distance between each pair of solutions and then calculate the average of these distances. The distance between two orderings or permutations $p = \langle p_1, p_2, ..., p_n \rangle$ and $q = \langle q_1, q_2, ..., q_n \rangle$ is given by

$$d(p,q) = \sum_{i=1}^{n} |p_i - q_i|.$$

Table 3.1 shows both measures for the three constructive methods. We can see that GC_3 provides the best results in terms of quality, since it obtains an average objective function value of 222013.87, which compares favorably with 178080.9 obtained with GC_1 and 161885.1 obtained with GC_2. On the other hand, GC_2 provides the best results in terms of diversity, since it obtains an average distance value of 802.5, which compares favorably with 743.6 obtained with GC_1 and 632.9 obtained with GC_3. However, if we consider both criteria together, quality and diversity, GC_1 is probably the best constructive method, providing a good trade-off between them.

Table 3.1 Comparison of constructive methods

Method	Quality	Diversity
GC_1	178080.90	743.6
GC_2	161885.10	802.5
GC_3	222013.87	632.9

GRASP may be viewed as a repetitive sampling technique, producing a sample solution in each iteration from an unknown distribution, whose mean and variance are functions of the restrictive nature of the RCL. For example, if RCL is restricted to a single element, then only one solution will be produced, the variance of the distribution will be zero and the mean will equal the value of the greedy solution. If the RCL has more than one element, then many different solutions will be generated, implying a larger variance. Since greediness plays a less important role in this case, the mean solution value may be worse. However, the value of the best solution found outperforms the mean value and can be better than the greedy solution (and in some cases may be optimal).

In the definition of the restricted candidate list RCL given above, we can see that only the elements with a value exceeding αe_1^* are included in this list. We can therefore interpret the parameter α as the degree of greediness (between 0 and 1)

in the selection of the element to be added to the partial solution. If α takes the value 0, all the candidate elements will be included in RCL and the method will actually perform a random selection (not greedy at all). On the contrary, if α takes the value 1, only the most highly evaluated elements will be included in RCL, and the method will perform a greedy selection (with no randomization at all). It is then expected that low values of α will produce diverse solutions of low quality, while large values of α will produce similar solutions (low diversity) of good quality.

To test this hypothesis, we consider the LOLIB instance be75eec and generate 100 solutions with the constructive method GC_1 with each particular value of α in the set $\{0.0, 0.1, 0.2, ..., 1.0\}$. Figure 3.2 shows the average value of the quality and diversity of each set of 100 solutions, obtained as in the previous experiment, for each value of α. In order to compare both values, quality and diversity, we represent them as a percentage and then we add them up to obtain a global evaluation of GC_1 with each α value considered.

Figure 3.2 shows that, as expected, when α is set to 0, the average quality obtained is the lowest (0% in the figure) but the diversity is the largest (100% in the figure). Symmetrically, when α is set to 1, the average quality obtained is close to the largest one (95% in the figure) but the diversity is the lowest one (0% in the figure). If we consider the addition of both percentage values (which means that we consider quality and diversity as equally important), we can see that when α is set to 0.4 we obtain the best solutions (185% in the figure). Therefore we will use GC_1 with α set to 0.4 as the constructive method in our GRASP algorithm for the LOP.

Fig. 3.2 Performance according to α

In [33] an alternative GRASP construction is proposed. It is based on a variant of the evaluation function $e_3(i)$ in which, instead of ratios, differences are computed. This way, the evaluation $e(i)$ of object i is the sum of the elements in the matrix in

its corresponding row minus the sum of the elements in its corresponding column, i.e., $e(i) = \sum_{j=1}^{n} c_{ij} - \sum_{j=1}^{n} c_{ji}$.

The authors compared the use of $e(i)$ with $e_3(i)$ and did not observe significant differences in the quality of solutions, although they stated that adapting the greedy function is computationally more expensive when ratios are used.

The construction method in [33] follows the same steps as the other GRASP methods described in this section (constructing first a candidate list of elements and computing then a restricted list RCL from which random selections are made). However, instead of inserting the selected element in the next available position, it inserts the element in the best position according to the value of the elements already selected. This strategy, introduced in [26], can be considered a local search in a reduced (partial) neighborhood.

3.2.2 Improvement Phase

The local search of our GRASP algorithm is based on the neighborhood search developed for the LOP presented in [86]. Here insertions are used as the primary mechanism to move from one solution to another. We again define $move(p_j, i)$ as the local modification deleting p_j from its current position j in permutation p and inserting it at position i (i.e., between the objects currently at positions $i-1$ and i).

A key to designing an efficient local search procedure is the incremental computation of the solution value. In other words, when we move from one solution p to another q, we need to calculate the objective function value of q from the value of p quickly and efficiently. In the LOP, when we apply $move(p_j, i)$ to solution p, obtaining solution q, it is easy to check that the value $z(q)$ of q can be directly obtained from the value $z(p)$ by considering only the elements in row and column j in the data matrix C. More precisely, if $i < j$ then

$$z(q) = z(p) + \sum_{k=i}^{j-1} (c_{p_j p_k} - c_{p_k p_j}).$$

Starting from a solution generated in the construction phase, the improvement phase of our GRASP algorithm performs iteration steps as long as the objective function increases. At each step, it considers an element p_j in p and scans all positions i from 1 to n, in search for the position i^* with the best associated move $move(p_j, i^*)$. If this is an improving move, it is performed; otherwise it is discarded and the method resorts to the next element p_{j+1} in the solution.

Fig. 3.3 Construction and local search values

Figure 3.3 shows the value of 100 solutions constructed with GC_1 (with α set to 0.4) and the value of the improved solutions with this local search procedure on the `LOLIB` instance `be75eec`. The effectiveness of the local search procedure is clearly shown, since it is able to improve the solutions by 48.5% on average. It also shows a large variability in the objective function of the constructed solutions, with a minimum value of 154728 and a maximum of 199072. However, on the contrary, it shows very similar values in the improved solution, with a minimum of 262387 and a maximum of 264940. Therefore, we cannot conclude that constructed solutions with low quality produce low quality improved solutions. This would indicate that diversity is more important than quality in the constructive method.

In our final experiment in this section, we apply the GRASP algorithm for 10 seconds to the 229 instances in the OPT-I set (with optimum known). Table 3.2 shows the number of instances in each set, the average percentage relative deviation *Dev* between the best solution value found with this method and the optimal value, and the number of instances, *#Opt*, for which an optimum solution is found.

Table 3.2 GRASP on OPT-I instances

	IO	SGB	RandAII	Rand B	MB	Special	**Total**
#Instances	50	25	25	70	30	29	229
Dev(%)	0.00	0.00	0.01	0.00	0.00	0.05	0.01
#Opt	49	23	5	70	21	14	182

The results in Table 3.2 clearly indicate that OPT-I instances do not pose a challenge for the GRASP methodology, since it is able to obtain 182 optimum solutions out of 229 instances in 10 seconds of running time. We will see in the next sections that most of the meta-heuristic are able to obtain high quality solutions on these instances within short running times.

3.3 Tabu Search

Tabu search [52] is a meta-heuristic that guides a local heuristic search procedure
to explore the solution space beyond local optimality by allowing non-improving
moves. This, in turn, requires some additional mechanisms based on memory struc-
tures to avoid cycling, which creates flexible search behavior. The name *tabu* comes
from the fact that move selection is limited because certain moves or solutions are
forbidden (declared tabu) as a result of the recorded information. In this section we
will see how to implement both simple and complex memory structures that allow
us to record useful information (called *attributes*) about the solutions visited during
the search. Based on this information we will design a local search method with
attributive and adaptive memory.

Tabu search (TS) uses *attributive memory* for guiding purposes. This type of
memory records information about solution attributes that change in moving from
one solution to another. For example, in a graph or network setting, attributes can
consist of nodes or arcs that are added, dropped or repositioned by the moving mech-
anism. In production scheduling, the index of jobs may be used as attributes to in-
hibit or encourage the method to follow certain search directions.

We can compare the memory of tabu search with the *explicit memory* of branch-
and-bound methods (see Chap. 4) in which an exhaustive memory is applied to know
the solutions we have already examined during the search process. In contrast, TS
incorporates *attributive* and *adaptive* memory. The attributive term means that we
are not going to record all the solutions generated (because we are not interested in
examining all the solutions of the problem) but we are going to record some of their
properties (attributes). The term "adaptive" means that the rules and properties that
we are going to apply may evolve during the search process.

In this section we start with an adaptation of the so-called *short term memory*
design to the LOP. In this design we only record attributes of recently visited solu-
tions. Then, we will see how we can add *long term memory* strategies to this basic
design to obtain a competitive solution method for the LOP. Both memory strategies
together, short and long term, constitute the core of most tabu search implementa-
tions.

If we see the search space as a huge set of solutions and think that we are only
able to explore a tiny part of it, we easily understand the rationale behind the two
search strategies *intensification* and *diversification*. Roughly speaking, the first one
favors the exploration of promising areas of the solution space, while the second one
drives the search to new regions of the solution space. Therefore, intensification and
diversification have complementary objectives. Intensification strategies are based
on modifying choice rules to encourage move combinations and solution features
that have historically been found to be good. In some settings, they consist of revis-
iting attractive regions to search them more thoroughly. Diversification strategies, on
the other hand, are based on visiting unexplored regions, increasing the effective-
ness in exploring the solution space of search methods based on local optimization.
Both strategies interact with the two memory structures mentioned above, short and
long term, to create an efficient search algorithm.

3.3.1 Short Term Memory

Tabu search begins in the same way as an ordinary local or neighborhood search, proceeding iteratively from one solution to another until a chosen termination criterion is satisfied. When TS is applied to an optimization problem with the objective of minimizing or maximizing $f(x)$ subject to $x \in X$, each solution x has an associated neighborhood $\mathcal{N}(x)$, and each solution $y \in \mathcal{N}(x)$ can be reached from x by an operation called a *move*.

When contrasting TS with a simple descent method where the goal is to minimize $f(x)$, we must point out that such a method only permits moves to neighbor solutions that improve the current objective function value and ends when no improving solutions can be found. Tabu search, on the other hand, permits moves that deteriorate the current objective function value but the moves are chosen from a modified neighborhood $\mathcal{N}^*(x)$. Short and long term memory structures are responsible for the specific composition of $\mathcal{N}^*(x)$. In other words, the modified neighborhood is the result of maintaining a selective history of the states encountered during the search.

In short term memory we usually consider attributes of the solutions recently visited, or the moves performed, in the last iterations. These attributes are used to exclude some elements in the neighborhood of the solutions in the next iterations. $\mathcal{N}^*(x)$ is typically a subset of $\mathcal{N}(x)$, and the *tabu classification* serves to identify elements of $\mathcal{N}(x)$ excluded from $\mathcal{N}^*(x)$. Let $\mathcal{T}(x)$ be this set of solutions in $\mathcal{N}(x)$, labeled as tabu, that we do not consider admissible for selection, i.e.,

$$\mathcal{N}^*(x) = \mathcal{N}(x) \setminus \mathcal{T}(x).$$

Recency based memory, as its name suggests, keeps track of attributes of solutions that have changed during the recent past. To exploit this memory, selected attributes that occur in solutions recently visited are labeled *tabu-active*, and solutions that contain tabu-active elements, or particular combinations of these attributes, are those that become tabu, thus being included in $\mathcal{T}(x)$. This prevents certain solutions from the recent past from belonging to $\mathcal{N}^*(x)$ and hence from being revisited.

The basic tabu search algorithm for the LOP proposed in [86] implements a short term memory structure alternating two phases: intensification and diversification. It is based on insertions as the improvement phase of the GRASP algorithm described above. However, instead of scanning the objects in the search for a move in their original order (from 1 to n), they are randomly selected in the intensification phase based on a measure of influence.

Short term tabu search

(1) Generate an initial solution x.
(2) Determine the attributes to establish the tabu status of solutions.
(3) While the stopping condition is not satisfied:

   (3.1) Compute $\mathcal{T}(x)$ and set $\mathcal{N}^*(x) = \mathcal{N}(x) \setminus \mathcal{T}(x)$.
   (3.2) Let y be the best solution in $\mathcal{N}^*(x)$.
   (3.3) Set $x = y$.
   (3.4) Update the tabu status of solutions.

For each object, there are at most m, $m \leq 2n - 2$, relevant elements in the matrix, i.e., those elements that may contribute to the objective function value. The elements in the main diagonal of the matrix are excluded because their sum does not depend on the ordering of the objects. This indicates that objects should not be treated equally by a procedure that selects an object for a local search (i.e., for search intensification). We define w_j as the weight of object j by setting

$$w_j = \sum_{i \neq j} (c_{ij} + c_{ji}).$$

Note that weight values do not depend on the permutation p, and therefore they can be calculated off-line, i.e., before the search begins. The weight values will be used to bias the selection of objects during the tabu search intensification phase.

Short Term Intensification Phase

An iteration in the *intensification phase* begins by randomly selecting an object. The probability of selecting object j is proportional to its weight w_j. The move $move(p_j, i)$ with the largest move value is selected. (Note that this rule may result in the selection of a non-improving move.) The move is executed even when the move value is not positive, resulting in a deterioration of the current objective function value. The moved object becomes tabu-active for *TabuTenure* iterations, and therefore it cannot be selected for insertions during this time.

The number of times that object j has been chosen to be moved is accumulated in the value *freq*(j). This frequency information is used for diversification purposes. The intensification phase terminates after *MaxIter* consecutive iterations without improvement.

Short Term Diversification Phase

The *diversification phase* is performed for *MaxDiv* iterations. In each iteration, an object is randomly selected, where the probability of selecting object j is inversely

proportional to the frequency count $freq(j)$. The chosen object is placed in the best position, as determined by the move values associated with the insert moves.

The basic tabu search procedure stops when *MaxGlo* global iterations are performed without improving the value of the best solution p^* found. A global iteration is an application of the intensification phase followed by the application of the diversification phase.

3.3.2 Long Term Memory

In TS strategies based on short term memory, $\mathcal{N}^*(x)$ characteristically is a subset of $\mathcal{N}(x)$, and the tabu classification serves to identify elements of $\mathcal{N}(x)$ excluded from $\mathcal{N}^*(x)$. In TS strategies that include longer term considerations, $\mathcal{N}^*(x)$ may also be expanded to include solutions not ordinarily found in $\mathcal{N}(x)$, such as solutions visited and evaluated in a past search, or identified as high quality neighbors of these past solutions.

The most common attributive memory approaches are *recency* based memory and *frequency* based memory. Recency, as its name suggests, keeps track of solutions' attributes that have changed during the recent past and we usually apply it in short term memory as described in the previous section. Frequency based memory provides a type of information that complements the information provided by recency based memory, broadening the foundation for selecting preferred moves. Like recency, frequency often is weighted or decomposed into subclasses. Also, frequency can be integrated with recency to provide a composite structure for creating penalties and inducements that modify move evaluations.

Frequencies typically consist of ratios, whose numerators represent the number of iterations (that we will refer to as *transition measures*) where an attribute of the solutions visited changes, and the denominators generally represent the total number of associated iterations. Alternatively, the numerators can represent the number of iterations where an attribute belongs to solutions visited on a particular trajectory (*residence measures*). Therefore, the ratios can produce transition frequencies that keep track of how often attributes change, or residence frequencies that keep track of how often attributes are members of solutions generated.

A long term diversification phase to complement the basic tabu search algorithm for the LOP is implemented in [86]. The long term diversification is applied after a pre-established number of global iterations have elapsed without improving the value of p^*.

For each object p_j, a rounded average position $\alpha(p_j)$ is calculated using the positions occupied by this object in the set of elite solutions and the solutions visited during the last intensification phase. The set of elite solutions consists of the best solutions found during the entire search. Instead of only recording (and updating) the best solution p^* found during the search, the long term memory records (and updates) a set of best solutions. Then, it calculates the average position occupied by each object in this set.

As mentioned above, in the computation of $\alpha(p_j)$ we also include the position occupied by the object during the last application of the intensification phase. Roughly speaking, we want to obtain a new solution far from those in the elite set but also far from those visited in the last iterations.

Long Term Diversification Phase

The long term diversification phase performs n steps, scanning the objects in their original order for $j = 1$ to n. In step j we insert object p_j in its complementary position $n - \alpha(p_j)$ according to the average position computed above. In mathematical terms, we apply $move(p_j, n - \alpha(p_j))$.

This strategy is inspired by the REVERSE operation of [26]. We, however, incorporate information about solutions that have been recently visited (during the last intensification phase) and solutions of high quality that have been found during the search (elite solutions). Purposefully constructing solutions that are "far away" from those in the elite set constitutes a diversifying element that also complements the diversification described in the short term memory.

In [86] series of experiments are performed to first establish the value of the key search parameters, and then to compare their method with the state of the art procedures. The first experiment has the goal of finding appropriate values for the three critical search parameters: *TabuTenure*, *MaxIter*, and *MaxDiv*. With a full factorial design with 3 levels for each parameter on the input-output matrices of LOLIB, it is determined that the best values are: *TabuTenure* $= 2\sqrt{n}$, *MaxIter* $= n$ and *MaxDiv* $= 0.5n$. In the second computational experiment the authors compared:

–   the short term memory tabu search method, STS (without the long-term diversification,
–   the complete tabu search method, TS,
–   the method of Chanas and Kobylanski [26],
–   the method of Becker [8], and
–   the improvement method of the GRASP algorithm, LS, described in Sect. 3.2.

We refer to Chanas and Kobylanski's method as CK, and as CK-10 for the application of the method from 10 randomly generated initial solutions. In a similar way, LS-10 refers to the application of the improvement local search method from 10 different starting solutions. The experimentation in [86] is limited to three sets of instances: the input-output instances IO of LOLIB, the instances SGB from the Stanford GraphBase [80] and randomly generated instances, RandomAI. A uniform distribution with parameters (0, 25000) was used to generate the random instances of sizes 75, 150 and 200 (25 instances per size). Table 3.3 shows, for each method, the average percentaged deviation from optimality and/or best known solutions, the number of optimum solutions, and the average CPU time (seconds on a 166 MHz Pentium). (Since the optimum solutions were not known for the random instances, we list in this case the number of best solutions found.)

Table 3.3 Comparison of methods

	LS	LS-10	Becker	CK	CK-10	STS	TS
Input-output tables							
Deviation	0.15	0.02	8.95	0.15	0.02	0.04	0.00
Normal deviation	0.18	0.03	10.38	0.18	0.03	0.05	0.00
No. of opt. solutions	11	22	0	11	27	30	47
CPU time	0.01	0.08	0.02	0.10	1.06	0.33	0.93
SGB instances							
Deviation	0.18	0.01	2.04	0.08	0.01	0.02	0.00
No. of opt. solutions	3	20	0	4	22	35	70
CPU time	0.06	0.55	0.20	1.45	16.33	1.16	4.09
Random instances							
Deviation	0.47	0.23	3.08	0.53	0.28	0.06	0.00
No. of best solutions	0	2	0	0	0	40	73
CPU time	0.12	1.21	0.80	10.67	108.44	10.79	20.19

Table 3.3 shows that Becker's procedure is clearly inferior in terms of solution quality although, given its simplicity, its performance is quite acceptable. The performance of the LS and CK methods is very similar across the three problem sets. TS outperforms all other methods in terms of solution quality, specially in the Random set, where it provides 73 of the 75 best known solutions. However, in terms of computational effort, TS consumes more running time than the other methods, with the exception of CK-10. The short term memory implementation, STS, performs very well considering its short running time. For the input-output matrices, we report in Table 3.3 both the deviation reported in [86] (see row "Deviation") and the deviation computed for the instances transformed to normal form (see row "Normal deviation"). As expected the latter is slightly larger (we include both deviations as a baseline for future comparisons).

Fig. 3.4 Performance graph

As a final result of the study in [86] the performance graph shown in Figure 3.4 is given in which the best three methods identified in the previous experiment were run in a way that the best solution found was reported every 0.5 seconds. These data points were used to generate the graph. The difference in quality between CK and the LS procedure may be due to the number of initial solutions used for each method. While the LS procedure can be applied 150 times during 5 seconds, the CK method can be applied only 10 times during the same amount of time. The superior performance of TS is made evident by Figure 3.4.

This study reveals that, although ignored in some tabu search implementations, the long term diversification by itself is an important component of the methodology, since in all cases it enhanced the performance of the basic procedure (as can be seen in Table 3.3 comparing the results of STS and TS).

We now report on our own experiment with tabu search for the LOP. Specifically, we apply the tabu search algorithm for 10 seconds to the 229 instances in the OPT-I set (with optimum known). Table 3.4 shows the number of instances in each set, the average percentage relative deviation, *Dev*, between the best solution value found with this method and the optimal value, and the number of instances, *#Opt*, for which an optimum solution is found.

Table 3.4 Tabu search on OPT-I instances

	IO	SGB	RandAII	Rand B	MB	Special	**Total**
#Instances	50	25	25	70	30	29	229
Dev(%)	0.00	0.00	0.00	0.00	0.00	0.02	0.00
#Opt	50	25	25	70	30	15	215

The results in Table 3.4 clearly show that the tabu search method is able to obtain the optimal solution for most of the instances in the OPT-I set. Specifically, it obtains 215 optimum out of 229 instances.

3.4 Simulated Annealing

The *simulated annealing* methodology [79] is based on a presumed analogy between the physical process known as annealing of solids and the algorithmic process of solving an optimization problem.

Annealing refers to a physical process in which a solid, placed in a heat bath, is first heated up by increasing the temperature of the bath, and then cooling down by slowly lowering this temperature. In this way, the particles of the solid arrange themselves in the low energy ground state (called crystal). If the cooling is done very quickly, irregularities are locked into the particles' structure, thus obtaining an amorphous solid (with a higher trapped energy level than in a perfectly structured crystal). Therefore, to obtain a perfect crystal, the temperature must be lowered gradually, following what Kirpatrick et al. [79] called *careful annealing*, where

temperature descends slowly through a series of levels. At each temperature value t, the solid is then allowed to reach *thermal equilibrium*, which can be characterized by a probability given by the Boltzmann distribution [41].

From the optimization viewpoint, SA proceeds in the same way as ordinary local search but incorporates some randomization in the move selection process. Specifically, it avoids getting trapped in a local optimum by means of nonimproving moves. These moves are accepted according to certain probabilities taken from the analogy with the annealing process. To simulate the change to thermal equilibrium of a solid for a given value of t, Metropolis et al. [97] proposed a *Monte Carlo method*. This simulation, introduced in the late 1940's to model the behavior of gases in a heat bath, is implemented in SA to determine whether a change in the physical system will be accepted or not.

Simulated annealing claims that an optimization problem or more precisely, the set of its feasible solutions, can be interpreted as the states of a physical system and that the process of careful annealing can be re-interpreted for the optimization problem to find an optimum solution. An optimum solution thus corresponds to a ground state. This analogy is shown in Table 3.5.

Table 3.5 Simulated annealing analogy

Physical system	Optimization problem
State	Feasible solution
Energy	Objective function
Ground state	Optimum solution
Careful annealing	Simulated annealing

At a given state of the solid (solution), the Monte Carlo method applies a random small perturbation and computes the difference in energy ΔE between the current state and the perturbed state. If this difference is negative ($\Delta E < 0$), the change is accepted and the perturbed solution becomes the current one because it has lower energy. However, if this difference is positive ($\Delta E > 0$), the probability of the perturbed state being accepted is computed according to the Boltzmann distribution. In mathematical terms

$$P(\text{new solution is accepted}) = e^{-\Delta E / t k_B},$$

where k_B is the *Boltzmann constant* [41]. The Monte Carlo method (also called the *Metropolis algorithm*) is applied to generate moves (states) and accept all improving moves, but in addition also worsening moves with a certain probability. It should be noted that, when run long enough, Monte Carlo simulation indeed generates states according to the Boltzmann distribution.

As is tabu search (TS) in its basic form, simulated annealing is also based on local search. When comparing SA and TS with a simple descent method to maximize $f(x)$, we know that such a method only permits improving moves to neighbor solutions. In physical terms, this would be *rapid quenching* instead of careful

annealing. Simulated annealing and tabu search, on the other hand, permit moves that deteriorate the current objective function value.

In the previous section we described how the moves are chosen in TS from a modified neighborhood $\mathcal{N}^*(x)$ according to the tabu status. Non-improving moves are accepted in SA according to the Boltzman distribution. Since the LOP is a maximization problem, a move from x to y is an improving one if $\Delta = f(y) - f(x) > 0$.

The basic outline of a simulated annealing method is the following. The algorithm starts at some temperature t which is decreased in the course of the algorithm. For each temperature, a number of moves according to the Metropolis algorithm is executed to simulate getting to the thermal equilibrium. Various schemes for guiding this process are given in the literature. For certain schemes one can prove that this algorithm indeed yields an optimum solution with probability 1, but these schemes lead to much too slow a convergence to be useful for practical applications. But at least, this result indicates the potential of the approach.

Simulated Annealing

(1) Generate an initial solution x.
(2) Choose an initial temperature t, a cooling factor c, $0 < c < 1$, and a repetition factor r.
(3) While not *frozen*:

    (3.1) Repeat the following steps r times:

        (3.1.1) Choose $y \in \mathcal{N}(x)$ at random.
        (3.1.2) Compute $\Delta = f(y) - f(x)$.
        (3.1.3) If $\Delta \geq 0$, set $x = y$.
        (3.1.4) If $\Delta < 0$, compute a random number p, $0 \leq p \leq 1$. If $p \leq \frac{e^{\Delta}}{t}$, set $x = y$.

(4) Update temperature $t = ct$.

Johnson et al. [71] give some guidelines to set the values of the parameters in the SA algorithm shown above: the initial temperature t, the cooling factor c, the repetition factor r and how to establish the stopping criterion (when it is *frozen*). The authors declare themselves to be skeptical about the relevance of the details of the analogy to the actual performance of SA algorithms in practice (and we share their point of view). In line with this, they examine and test the values of this key search parameters from the performance viewpoint, regardless of their meaning in the physical process. They actually translate the computation of these parameters into others that are more closely related to the optimization process:

–   A new parameter, *initprob*, is used to determine an initial temperature t. Based on an abbreviated trial annealing run, a temperature is found at which the fraction of accepted moves is approximately *initprob*, and this is used as a starting temperature. They recommend setting *initprob* to 0.4 in the problem tested.

- The repetition factor r is set as $r = 16s$ where s is the expected neighborhood size.
- The *minpercent* parameter is introduced to determine whether the annealing run is *frozen* or not. A counter is kept which is incremented each time a temperature is completed for which the percentage of accepted moves is at most *minpercent*, and is reset to 0 each time a solution is found to be better than the incumbent one. If the counter reaches 5 the process is declared *frozen* and the SA algorithm stops.

As far as we know there is no previous implementation of the SA methodology to the LOP. However, there are several approximations with similar methods or variants of the LOP. In [11] a SA method is proposed for the *chromosome reconstruction problem*. This problem arises when creating maps in genetics to reconstruct the DNA sequence. Although the authors called the mathematical formulation of this problem *optimal linear ordering*, it should be noted that it involves computing an ordering (clone ordering) that minimizes the sum of differences between successive clones. However, in this computation they do not include the differences between noncontiguous elements (clones) and only consider consecutive elements. Thus, it is not the same objective function as in the LOP. On the other hand, as will be shown in a subsequent chapter, SA has been implemented to solve an auxiliary problem that appears when studying the convex hull of feasible solutions (the polytope) of the LOP. Specifically, we will see how this method is applied to separate small facets defining inequalities for the LOP.

Charon and Hudry [28] describe the application of the *noising method* to some combinatorial optimization problems, including the LOP. This method proceeds in the same way as local search, performing a move at each iteration, from the current solution x to a neighbor solution $y \in \mathcal{N}(x)$. However, when they compute the *move value*, $\Delta = f(y) - f(x)$, they add a perturbation or *noise* and, instead, consider $\Delta_{noise} = f(y) - f(x) + \sigma$ where σ is a random value drawn from the interval $[-r, r]$ according to a pre-established probability distribution. As the search progresses the *noise rate r* decreases and the method finishes when no improvement move is available. The authors point out that different methodologies, including SA, can be considered an instance of the noising method. In other words, if we choose the parameters properly, especially the probability distribution, the noising method generalizes these methodologies. In their computational experiments, the authors test their adaptation of the noising method to the LOP with different parameters, including noising rates, descent variants and restarting mechanisms, and conclude that the method provides high quality solutions.

As in the previous sections, we report the results obtained when applying this method for 10 seconds to the 229 instances in the OPT-I set. Specifically, we implemented the SA method described in the basic outline above. Table 3.6 shows the number of instances in each set, the average percentage deviation, *Dev*, and the number of instances, *#Opt*, for which SA obtains an optimum solution.

Table 3.6 Simulated Annealing on OPT-I instances

	IO	SGB	RandAII	Rand B	MB	Special	**Total**
#Instances	50	25	25	70	30	29	229
Dev(%)	0.03	0.03	0.08	0.25	1.33	0.39	0.35
#Opt	16	0	0	10	0	6	32

The results in Table 3.6 clearly show that the simulated annealing obtains good solutions on the OPT-I set, since the average percentage deviation overall is 0.35%. However, these results are inferior in quality to those obtained with the tabu search method (as shown in Table 3.4, TS presents an average percent deviation of 0.00%). At the end of this chapter we will show an extensive comparison among all the methods described.

3.5 Variable Neighborhood Search

Variable neighborhood search (VNS) is rapidly becoming a method of choice for designing solution procedures for hard combinatorial optimization problems [63]. VNS is based on a simple and effective idea: a systematic change of the neighborhood within a local search algorithm. In this section we follow the description given in [49] to adapt the VNS methodology to the LOP.

As stated in [63], VNS is based on the following three principles, where principle (2) is true for all optimization problems, but principles (1) and (3) may or may not hold depending on the problem at hand.

(1)   A local minimum with respect to one neighborhood is not necessarily so with another.
(2)   A global minimum is a local minimum with respect to all possible neighborhood structures.
(3)   For many problems local minima with respect to one or several neighborhoods are relatively close to each other.

To apply the VNS methodology we first need to define different neighborhoods for our problem. Let again $move(p_j, i)$ denote the local modification where p_j is deleted from its current position j in permutation p and inserted at position i (i.e., between the objects currently in positions $i - 1$ and i). In [49] García et al. proposed the following neighborhoods to adapt the VNS methodology to the LOP:

$$\mathcal{N}_1(p) = \{p' \mid p' \text{ is obtained by } move(p_j, i), \text{ for some } j = 1, \ldots, n,$$
$$\text{and } i \in \{j - 1, j + 1\}\},$$
$$\mathcal{N}_2(p) = \{p' \mid p' \text{ is obtained by } move(p_j, i), \text{ for some } j = 1, \ldots, n,$$
$$\text{and } i = 1, \ldots, n, i \neq j\}.$$

The neighborhood $\mathcal{N}_1(p)$ consists of permutations p' that are reached by switching the positions of contiguous objects in p. $\mathcal{N}_2(p)$ consists of all permutations p' resulting from executing general insertion moves in p. Based on $\mathcal{N}_2$, composed neighborhoods $\mathcal{N}_k$, $k = 3, \ldots, n$, can be defined. The neigborhood $\mathcal{N}_k(p)$ is the set of solutions that are obtained when we apply the general insertion move $k - 1$ times from p. Obviously, $\mathcal{N}_k$ can be defined in terms of $\mathcal{N}_2$. E.g., for $k = 3$, a solution r is in $\mathcal{N}_3(p)$ if, for some $q \in \mathcal{N}_2(p)$, we have $r \in \mathcal{N}_2(q)$.

In the following subsections we will describe five different implementations of variable neighborhood search (and its variants) for the LOP. We will first describe them in detail and then report some computational experiments to compare them.

3.5.1 Variable Neighborhood Descent

The *variable neighborhood descent (VND)* method is obtained when the change between neighbors is performed in a deterministic way.

The implementation for the LOP performs a local search for the best solution in $\mathcal{N}_1$ and only resorts to performing one move in $\mathcal{N}_2$ when the search is trapped in a local optimum found in $\mathcal{N}_1$. If the local search in $\mathcal{N}_2$ is able to improve the best solution found so far, the search continues in $\mathcal{N}_1$, otherwise it terminates.

3.5.2 Restricted Variable Neighborhood Search

This restricted version RVNS is also limited to $\mathcal{N}_1$ and $\mathcal{N}_2$. It repeatedly performs three steps combining stochastic and deterministic strategies (assume $k \in \{1, 2\}$):

(1) A solution p' is randomly generated in $\mathcal{N}_k(p)$. (**Shaking**)
(2) Apply a local search from p' to obtain a local optimum p''. (**Improving**)
(3) If p'' is better than p, then p is replaced by p'' and k is set to 1; otherwise, k is switched (from 1 to 2 or from 2 to 1). (**Updating**)

As in VND, k is initially set to 1 and the method resorts to $\mathcal{N}_2$ when $\mathcal{N}_1$ (now in combination with local search) fails to improve the current solution. However, if $\mathcal{N}_2$ also fails to improve on the incumbent solution, instead of stopping the search, RVNS sets $k = 1$ and randomly selects another trial solution in $\mathcal{N}_1$, repeating the three steps again. The sequence is repeated until a *MaxIter* number of consecutive iterations is performed with no further improvement. In step (2) we apply the local search method based on $move(p_j, i)$ ($\mathcal{N}_2(p)$) described in the previous chapter.

3.5.3 Basic Variable Neighborhood Search

The basic variable neighborhood search method (BVNS) follows the same scheme as RVNS based on the three steps shaking, improving and updating. However, in this version the method uses k_{max} neighborhoods.

Basic variable neighborhood search

(1) Generate an initial solution p.
(2) Define neighborhoods $\mathcal{N}_k$ with $k=1,\ldots,k_{max}$.
(3) Set $counter = 0$ and $k = 1$.
(4) While $counter < MaxIter$:

    (4.1) Randomly generate a solution p' in $\mathcal{N}_k(p)$.
    (4.2) Apply a local search from p' to obtain a local optimum p''.
    (4.3) If p'' is better than p, then set $p = p''$, $counter = 0$ and $k = 1$.
    (4.4) Else set $k = k + 1$ (if $k = k_{max}$, then set $k = 1$) and $counter = counter + 1$.

Initially, k is set to 1 and in the shaking step a solution $p' \in \mathcal{N}_k(p)$ is randomly generated. Then, a local search method is applied from p' to obtain a local optimum p''. In the update step, if p'' is better than p, p is replaced by p'' and k is set to 1; otherwise, k is incremented. The method repeats these three phases until a *MaxIter* number of consecutive iterations is performed with no further improvement. As in the previous version, we consider the $\mathcal{N}_2$ descent procedure as the local search phase.

3.5.4 Frequency Variable Neighborhood Search

This variant combines VNS and tabu search. As we have seen in the previous variants, VNS is mainly based on random sampling of selected neighborhoods in combination with local search. On the other hand, tabu search is based on the notion of recording information (attributes) to perform a guided and deterministic (i.e., not random) search of the solution space. We therefore can say that to some extent they represent opposite approaches. However, in the frequency variable neighborhood search, FVNS, both approaches are integrated into a single design.

As in BVNS, the FVNS method repeats shaking, improving and updating steps until a *MaxIter* number of consecutive iterations is performed with no further improvement. However, in FVNS, the shaking step is guided by frequency information recorded in previous iterations. This is made to diversify the search in a controlled way (instead of a completely random and uncontrolled way). As previously shown (see Sect. 3.3), diversification is the notion of expanding the search to unexplored regions in the solution space (and it does not necessarily means randomization).

Diversification strategies are generally based on either encouraging the incorporation of new elements or discouraging elements visited frequently.

In Chap. 2 we compared several different constructive methods for the LOP. Among them, in particular, FQ is based on a frequency counter that records the number of times an element appears in a specific position. This counter is used to penalize the "attractiveness" of an element with respect to a given position. In a similar way, we now use a frequency counter $freq(i)$ to record the number of times object i has been moved. Therefore, each time object i is moved from one position to another in the shaking or the local search phase, we increment $freq(i)$ by one unit. We use this frequency counter to generate a new solution in the shaking step. Since we want to diversify, we select objects j with a small frequency value $freq(j)$. Specifically, in the shaking step, we randomly select k_{max} objects in the incumbent solution p to be moved. The probabilistic selection rule is inversely proportional to the frequency count. The selected objects are moved to the best available position.

3.5.5 Hybrid Variable Neighborhood Search

As mentioned in the introduction of this chapter, researchers in the meta-heuristic field increasingly propose combinations and hybridizations among the different methodologies, which makes the boundary between them less well defined, in many cases taking some elements from others. Actually we described in the previous subsection FVNS in which a simple memory structure based on frequencies (a tabu search element indeed) plays an important role in the shaking step of VNS.

In this subsection we describe two straightforward hybridizations of VNS in which we simply replace the local search of the improving step with a meta-heuristic method, thus obtaining the following three steps:

(1) A solution p' is randomly generated in $\mathcal{N}_k(p)$. (**Shaking**)
(2) Apply a meta-heuristic from p'. Let p'' be the output solution. (**Searching**)
(3) If p'' is better than p, then p is replaced by p'' and k is set to 1, otherwise, k is incremented by 1 (if $k = k_{max}$, k is set to 1). (**Updating**)

We considered the two previous VNS variants, BVNS and FVNS, and replaced the local search with the short term memory tabu search algorithm described in the previous section (in which no longer term structures are present) according to the scheme above. We called the resulting methods BVNS-TS and FVNS-TS, respectively. In their computational experiments, García et al. [49] tested these procedures on the input-output instances from LOLIB, the Stanford GraphBase instances, and the random instances A of type I and type II.

In a preliminary experiment of [49] simple local search methods are compared to find appropriate values of the key search parameters. Specifically, they compare VND with the local search based on the $\mathcal{N}_2$ neighborhood on both real and random instances. The experiment confirms what is well known for the LOP: random instances are more difficult than the real input-output instances. It also shows that

there are small variations in the results of these procedures. However, if we run them from different initial solutions it becomes clear that VND saves time since it only resorts to $\mathcal{N}_2$ when the search is trapped in $\mathcal{N}_1$.

In the first experiment of [49] VNS variants that do not incorporate long term strategies are compared with two previous methods on the input-putput instances and on the random instances of type I. Specifically, they compare RVNS, BVNS, FVNS with the short term tabu search, ST-TS [86], described in the previous section, and the method by Chanas and Kobilansky, CK-10 [26], run from 10 randomly generated initial solutions. Table 3.7 reports the number of best solutions found (optimum solutions for the LOLIB) and the percentaged deviation of each method on each set of instances.

Table 3.7 shows that the best solution quality is obtained by the variants of the VNS methodology, which are able to match a larger number of optimum and best known solutions than the short term TS and CK methods. In particular, for the input-output instances, RVNS finds 39 optimum solutions, BVNS 35, FVNS 34, ST-TS 30 and CK-10 finds 27. On the other hand, on the random instances, RVNS determines only 10 best known solutions out of 25, BVNS 19, FVNS 17, ST-TS 14 and CK-10 finds 4. All methods are extremely fast considering that their running times are below 0.02 seconds. The performance of CK is clearly inferior with a lower number of optimum solutions than those achieved by the other approaches. However, as mentioned, it is a simple heuristic and its results are quite acceptable considering its simplicity.

Table 3.7 Comparison of basic methods

Instances	RVNS	BVNS	FVNS	ST-TS	CK-10
Input-output					
No. of opt. solutions	39	35	34	30	27
Percentaged deviation	0.02	0.03	0.05	0.04	0.02
Random A type I					
No. of best solutions	10	19	17	14	4
Percentaged deviation	0.15	0.03	0.04	0.05	0.12

The final experiments in [49] compare basic and hybrid VNS with the best heuristics for the LOP, in particular BVNS, BVN-TS, FVN-TS and the previous approaches TS (the complete tabu search approach described in the previous section) and SS (the scatter search approach described in the next section).

Table 3.8 shows, for each procedure, the average deviation from optimality in percent, the number of optimum solutions, and the average CPU time in seconds for each set of instances. Since optimum solutions are not known for the large random instances, the deviation for these problems is given with respect to the best solution found during each experiment. Also the number of best solutions found is reported instead of the number of optimum solutions. We have set the stopping parameter *MaxIter* in the VNS versions to 100 to approximate the running time consumed by the LT-TS method.

For the input-output instances the long term tabu search algorithm TS is able to obtain 47 optimum solutions in 0.024 seconds while the VNS variants obtain a number of optimum solutions that ranges between 39 and 46 in less computation time than LT-TS. The performance of the SS method in this experiment is clearly inferior in terms of quality considering its running time.

For large and more difficult instances, as in the previous experiments, SS obtains very good solutions but it needs longer running times than the other methods. FVN-TS and LT-TS are clearly the best methods in terms of solution quality achieved within small running times. Both obtain the same number of optimum solutions for the SGB instances, although LT-TS presents a smaller percentaged deviation on average and a larger computational time than FVN-TS. On the other hand, FVN-TS obtains 6 best solutions and 0.16 average percent deviation in the random type I instances, while LT-TS obtains 10 best solutions and 0.06 average percent deviation. Results in random type II instances are different since FVN-TS is able to obtain 29 best solutions in 0.22 seconds of running time, which compares favorably with all the other methods considered.

It is interesting to see that, although the frequency VNS version (FVNS) does not improve the memory less variant (VNS) as shown in previous tables, when we coupled the VNS methods with tabu search it seems that the use of frequency based memory improves the basic VNS in solving the LOP.

Table 3.8 Comparison of best methods

	BVNS	BVN-TS	FVN-TS	TS	SS
Input-output instances					
Deviation	0.0208	0.0370	0.0082	0.0007	0.0133
No. of opt. solutions	40	41	46	47	42
CPU time	0.015	0.013	0.018	0.024	0.04
SGB instances					
Deviation	0.0251	0.0087	0.0104	0.0018	0.0023
No. of opt. solutions	7	11	14	14	15
CPU time	0.067	0.039	0.052	0.090	0.153
Random A type I					
Deviation	0.1870	0.1615	0.1600	0.0615	0.0130
No. of best solutions	4	5	6	10	48
CPU time	1.020	0.289	0.305	0.417	0.709
Random A type II					
Deviation	0.0053	0.0029	0.0019	0.0014	0.0017
No. of best solutions	16	20	29	28	18
CPU time	0.607	0.336	0.220	0.269	0.457

Figure 3.5 shows in a box-and-whisker plot the value of the best solution obtained with LT-TS and FVN-TS on the largest random instances (types I and II with $n=200$).

Fig. 3.5 Best value box plot

Note that the higher the value, the better the method (since we are maximizing the objective function). This figure clearly shows that both methods present similar results (in terms of the best solution found) although LT-TS presents a marginal improvement over FVN-TS in these instances since it is able to obtain, in some cases, better solutions (the higher *whisker* is larger in LT-TS than in FVN-TS).

As in previous sections, we report the results obtained when applying this method for 10 seconds to the 229 instances in the OPT-I set. Specifically, we consider the basic VNS method described above. Table 3.9 shows the number of instances in each set, the average percentaged deviation *Dev*, and the number of instances *#Opt*, for which SA obtains an optimum solution.

Table 3.9 BVNS on OPT-I instances

	IO	SGB	RandAII	Rand B	MB	Special	**Total**
#Instances	50	25	25	70	30	29	229
Dev(%)	0.00	0.00	0.00	0.02	0.00	0.11	0.02
#Opt	50	25	19	64	30	20	208

Results in Table 3.9 clearly show that the VNS methodology obtains high quality solutions on the OPT-I set. The average percentaged deviation overall is 0.02% and the number of optimum solutions is 208 out of 229 instances.

3.6 Scatter Search

Scatter search (SS) is an evolutionary or population based method in the sense that it operates on a set of solutions, combining them to obtain new and hopefully better solutions [85]. Nowadays it is a well established method within the meta-heuristic community and has been successfully applied to a wide range of optimization problems. However, general awareness of the method still lags behind that of other population based methods such as genetic algorithms or well established meta-heuristics like tabu search.

There are three elements that we need to define in any evolutionary method: a way to generate solutions, a way to combine solutions and a way to maintain a set (population) of solutions. When we design these elements for the problem we are faced with, we follow the guidelines given by the meta-heuristic. In this section we will see how to define these elements for the LOP and how they interact according to the SS methodology.

Scatter search was first introduced in [53] as a heuristic for integer programming. In the original proposal, solutions are purposely (i.e., non-randomly) generated to take characteristics of various parts of the solution space into account. Scatter search orients its explorations systematically relative to a set of reference points that typically consist of good solutions obtained by prior problem solving efforts, where the criteria for "good" are not restricted to objective function values, and may apply to subcollections of solutions rather than to a single solution, as in the case of solutions that differ from each other according to certain specifications.

The scatter search methodology is very flexible, since each of its elements can be implemented in a variety of ways and degrees of sophistication. In this section we give a basic design to implement scatter search based on the following five methods:

(1) A *diversification generation method* to generate a collection of diverse trial solutions, using an arbitrary trial solution (or seed solution) as an input.

(2) An *improvement method* to transform a trial solution into one or more enhanced trial solutions. (Neither the input nor the output solutions are required to be feasible, though the output solutions will usually be expected to be feasible. If no improvement of the input trial solution results, the enhanced solution is considered to be the same as the input solution.)

(3) A *reference set update method* to build and maintain a reference set consisting of the b best solutions found (where the value of b is typically small, e.g., not greater than 20), organized to provide efficient access by other parts of the method. Solutions gain membership to the reference set according to their quality or their diversity.

(4) A *subset generation method* to operate on the reference set and to produce a subset of its solutions as a basis for creating combined solutions.

(5) A *solution combination method* to transform a given subset of solutions produced by the subset generation method into one or more combined solution vectors.
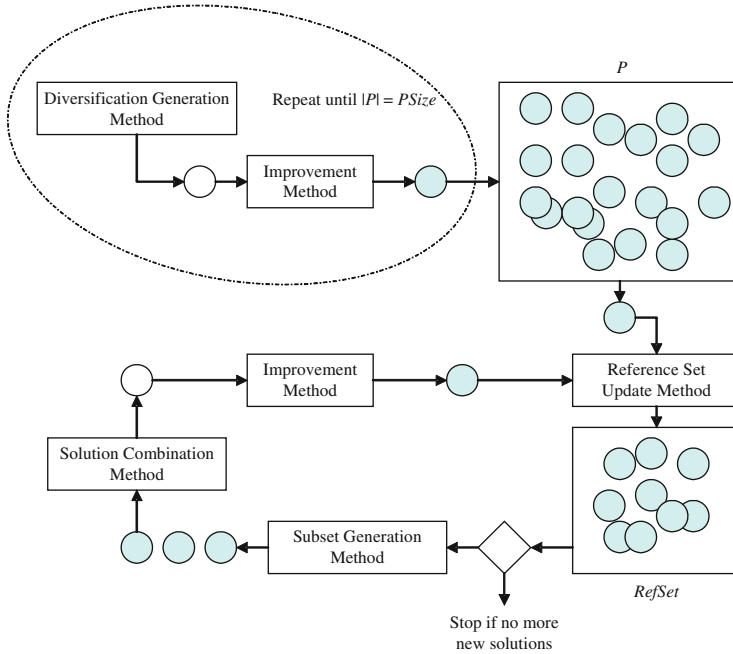
Fig. 3.6 Schematic flow diagram of scatter search

In Figure 3.6, diversification generation and improvement methods are initially applied, adding improved solutions to *P*, until the cardinality of *P* reaches *PSize* solutions that are different from each other. The darker circles represent improved solutions resulting from the application of the improvement method. The main search loop appears to the left of the box containing the reference solutions (labeled *RefSet*). The subset generation method takes reference solutions as input to produce solution subsets to be combined. Solution subsets contain two or more solutions. The new trial solutions resulting from the application of the combination method are subjected to the improvement method and handed to the reference set update method. This method applies rules regarding the admission to the reference set of solutions coming from *P* or from the application of the combination and improvement methods. Of the five methods in scatter search, only four are strictly required. The improvement method is usually needed if high quality outcomes are desired, but a scatter search procedure can be implemented without it. The advanced features of scatter search are related to the way the five methods are implemented. That is, the sophistication comes from the implementation of the SS methods instead of the decision to include or exclude some elements (like in the case of tabu search).

Scatter search

(1) Set $P = \emptyset$ and $PSize = 100$.
(2) While $|P| < PSize$:

    (2.1) Generate solution x with the diversification generation method.
    (2.2) Apply the improvement method to x. Let x' be the output.
    (2.3) If $x' \notin P$, then add x' to P.

(3) Build the reference set $RefSet = \{x_1, \ldots, x_b\}$ containing b "good" (w.r.t. quality and diversity) solutions in P.
(4) Set $NewSolutions = \texttt{True}$.
(5) While $NewSolutions$:

    (5.1) Generate a set $NewSubsets$ which consists of all pairs of solutions in $RefSet$ that include at least one new solution.
    (5.2) Set $NewSolutions = \texttt{False}$.
    (5.3) While $NewSubsets \neq \emptyset$:

        (5.3.1) Select the next set $S \in NewSubsets$.
        (5.3.2) Apply the solution combination method to S to obtain one or more new solutions x.
        (5.3.3) Apply the improvement method to x. Let x' be the output.
        (5.3.4) If $x' \notin RefSet$ and improves its worst solution, then replace this with x' and set $NewSolutions = \texttt{True}$.
        (5.3.5) Delete S from $NewSubsets$.

(6) Output the best solution in $RefSet$.

The core of the SS method is the *reference set* in which good solutions are stored (where good refers not only to quality but also to diversity). The scatter search algorithm has three main stages according to the status of the reference set: reference set creation, reference set update, and reference set rebuild.

In the following subsections we will describe how these three stages can be implemented to build a scatter search algorithm for a given problem and how they are adapted in [19] to the LOP.

Similarities and contrasts between SS and the original proposals for genetic algorithms (GA) are observed in [85]. Both are instances of what are sometimes called "population based" or "evolutionary" approaches. Both incorporate the idea that a key aspect of producing new elements is to generate some form of combination of existing elements. However, GA approaches are predicated on the idea of choosing parents randomly to produce offsprings, and further on introducing randomization to determine which components of the parents should be combined. In contrast, the SS approach does not emphasize randomization, particularly in the sense of being indifferent to choices among alternatives. Instead, the approach is designed to incorporate strategic responses, both deterministic and probabilistic, that take evaluations and history into account. SS focuses on generating relevant outcomes without losing

the ability to produce diverse solutions, due to the way the generation process is implemented.

3.6.1 Reference Set Creation

The diversification generation method is used to build a large set P of diverse solutions. The size $PSize$ of P is typically at least 10 times the size of $RefSet$. The initial reference set is built taking solutions from P according to the reference set update method.

The reference set $RefSet$ is a collection of both high quality solutions and diverse solutions, which are used to generate new solutions by way of applying the combination method. In this basic design we can use a simple mechanism to construct an initial reference set based on a distance function between solutions, and then update it during the search. The size of the reference set is denoted by $b = b_1 + b_2$. The construction of the initial reference set starts with the selection of the best b_1 solutions from P. These solutions are added to $RefSet$ and deleted from P. For each solution in $P \setminus RefSet$, the minimum of the distances to the solutions in $RefSet$ is computed. Then, the solution with the maximum of these minimum distances is selected. This solution is added to $RefSet$ and deleted from P and the minimum distances are updated. (In applying this max-min criterion, or any criterion based on distances, it can be important to scale the problem variables, to avoid a situation in which a particular variable or subset of variables dominates the distance measure and distorts the appropriate contribution of the vector components.) The process is repeated b_2 times, thus the resulting reference set has b_1 high quality solutions and b_2 diverse solutions.

In Chap. 2 we reviewed constructive methods for the LOP: G1, G2, G3, G4, G5, G6, MIX, RND, DG, and FQ. Any of them can be used as a diversification generation method within a SS algorithm. Moreover, we have seen that FQ, which is based on frequency information, performs best, considering quality and diversity as equally important. Therefore, we will employ this generator in our SS method for the LOP.

In the standard SS design [85] the improvement method is applied to all the solutions in P. However, in some problems where local search methods are usually extremely time consuming, the improvement method is not applied across the board but rather in a selective manner. In the LOP we follow the standard design and apply the improvement method to all the solutions in P. In Sect. 3.2 we describe a local search procedure using insertions. Based on moves $move(p_j, i)$, the local search method performs steps as long as the objective function increases. We apply this local search as the improvement method of our SS algorithm.

3.6.2 Reference Set Update

The search is initiated in this second stage by applying the subset generation method that, in its simplest form, involves generating all pairs of reference solutions. The pairs of solutions in *RefSet* are selected one at a time and the solution combination method is applied to generate one or more trial solutions. These trial solutions are subjected to the improvement method. The reference set update method is applied once again to build the new *RefSet* with the best solutions, according to the objective function value, from the current *RefSet* and the set of trial solutions. The basic procedure terminates after all the subsets generated are subjected to the combination method and none of the improved trial solutions are admitted to *RefSet* under the rules of the reference set update method.

Solution combination methods in scatter search typically are not limited to combining just two solutions and therefore the subset generation method in its more general form consists of creating subsets of different sizes. The scatter search methodology is such that the set of combined solutions (i.e., the set of all combined solutions that the implementation intends to generate) may be produced in its entirety at the point where the subsets of reference solutions are created. Therefore, once a given subset is created, there is no merit in creating it again. This creates a situation that differs noticeably from those considered in the context of genetic algorithms, where the combinations are typically determined by the spin of a roulette wheel (see next section for a description).

The procedure for generating subsets of reference solutions in advanced SS applications uses a strategy to expand pairs into subsets of larger size while controlling the total number of subsets to be generated. In other words, the mechanism does not attempt to create all subsets of size 2, then all the subsets of size 3, and so on until reaching the subsets of size $b-1$ and finally the entire *RefSet.* This approach would not be practical because there are 1023 subsets in a reference set of a typical size $b = 10$. Even for a smaller reference set, combining all possible subsets is not effective, because many subsets will be almost identical. For example, a subset of size four containing solutions 1, 2, 3, and 4 is almost the same as all the subsets with four solutions for which the first three solutions are solutions 1, 2 and 3. And even if the combination of subset $\{1,2,3,5\}$ were to generate a different solution than the combination of subset $\{1,2,3,6\}$, these new trial solutions would likely converge to the same local optimum after the application of the improvement method.

The following approach selects representative subsets of different sizes by creating subset types:

– **Type 1**: all 2-element subsets,
– **Type 2**: 3-element subsets derived from the 2-element subsets by augmenting each 2-element subset to include the best solution not in this subset,
– **Type 3**: 4-element subsets derived from the 3-element subsets by augmenting each 3-element subset to include the best solutions not in this subset, and
– **Type 4**: the subsets consisting of the best i elements, for $i = 5$ to b.

We will use this approach in our subset generation method for the LOP. In [19] an experiment is designed with the goal of assessing the contribution of combining subset types 1 to 4 in the context of the LOP. The experiment tried to identify how often, across a set of benchmark problems, the best solutions came from combinations of reference solution subsets of various sizes. Since subset types 1 to 4, respectively, generate solutions from 2 to up to b reference solutions, it is sufficient to keep a 4-element array for each solution generated during the search. The first element of the array is the counter corresponding to subset type 1; the second element corresponds to subset type 2, etc. The array for each solution in the initial reference set starts as (0,0,0,0), meaning that there are no sources. The array then counts the number of times the different subset types are used. E.g., suppose that three solutions in a subset of type 2 with arrays $[2,0,0,1]$, $[5,1,0,0]$ and $[0,1,0,0]$ are combined. Then a new solution resulting from this combination is accompanied by the array $[7,3,0,1]$, the sum of the other arrays, plus 1 added to position 2.

In an experiment with 15 input-output instances from LOLIB tracking arrays are used in [19] to find the percentage of times that each subset type produces solutions that become members of the reference set. The same experiment was also conducted employing 15 randomly generated instances with 100 objects and with weights drawn uniformly distributed between 0 and 100. The percentages are shown in Figure 3.7, where the bars labeled "LOLIB" represent the results from the experiments with the input-output instances and the bars labelled "Random" correspond to the results from the randomly generated instances.

As noted before, the combination method is an element of scatter search whose design depends on the problem context. Although it is possible in some cases to design context independent combination procedures, it is generally more effective to base the design on specific characteristics of the problem setting.

The combination method for the LOP employs a min-max construction based on votes. The method scans (from left to right) each solution in a subset (called the reference permutation), and uses the rule that each reference permutation in the combination subset votes for its first object that so far has not been included in the combined permutation (referred to as the "incipient object"). The voting determines the object to be assigned to the next free position in the combined permutation (where the incipient object with more votes is assigned). This is a min-max rule in the sense that, if any object of the reference permutation is chosen other than the incipient object, then it would increase the deviation between the reference and the combined permutations. Similarly, if the incipient object were placed later in the combined permutation than its next available position, this deviation would also increase. So the rule attempts to minimize the maximum deviation of the combined solution from the reference solution, subject to the fact that other reference solutions in the subset are also competing to contribute.

Fig. 3.7 Performance of subset types

This voting scheme can be implemented using a couple of variations that depend on the way votes are modified:

–   The vote of a given reference solution is weighted according to the incipient object's position (referred to as the "incipient position"). A smaller incipient position gets a higher vote. For example, if the object in the first position of some reference permutation is not assigned to the combined permutation during the first 4 assignments, then the vote is weighted more heavily to increase the chances of having that object assigned to position 5 of the combined permutation. The rule emphasizes the preference of this assignment to one that introduces an (incipient) object that occurs later in some other reference permutation.
–   A bias factor gives more weight to the vote of a reference permutation with higher quality. Within the current organization of the scatter search implementation in this tutorial such a factor should only have weak influence because it is expected that high quality solutions will be strongly represented anyway.

We chose to implement the first variant with a tie-breaking rule based on solution quality. The tie-breaking rule is used when more than one object receives the same votes. Then the object with highest weighted vote is selected, where the weight of a vote is directly proportional to the objective function value of the corresponding reference solution.

3.6.3 Reference Set Rebuild

In basic scatter search implementations, the reference set is updated by replacing the reference solution with the worst objective function value with a new trial solution that has a better objective function value. Since we always assume that *RefSet* is ordered, the best solution is x_1 and the worst solution is x_b. So, when a new trial solution x is generated as a result of the application of the combination and

improvement methods, the objective function value of the new trial solution x is used to determine whether *RefSet* needs to be updated (it is updated by replacing x_b with x and reordering its elements). If the reference set remains unchanged after the update (no trial solution improves the worst solution x_b in the set), then a rebuild step is performed.

A reference set $\{x_1, x_2, \ldots, x_{b_1}, x_{b_1+1}, \ldots, x_{b_1+b_2}\}$, where $b = b_1 + b_2$, is partially rebuilt with the following diversification update when no new trial solutions are admitted to it:

(1) Solutions $x_{b_1+1}, \ldots, x_{b_1+b_2}$ are deleted from *RefSet*.
(2) The diversification generation method is used to construct a set P of new solutions.
(3) We sequentially select b_2 solutions from P and move them to *RefSet*. We apply the same min-max criterion, which is part of the reference set update method, as we did when *RefSet* was constructed the first time.

Figure 3.8 shows a schematic representation of the rebuilding mechanism. A series of experiments is performed to assess the quality of the SS implementation in [19] for the LOP. We report here on three of them. The first one is designed to find the best values for the key parameters of the scatter search algorithm. For this experiment, again 15 input-output instances and 15 randomly generated instances with 100 objects (as described above) are used.

The following values were tested during these experiments:

– *PSize*: 50, 100, 150,
– b: 10, 20, 40,
– (b_1, b_2): $(5,5)$, $(10,10)$, $(5,15)$, $(15,5)$, $(20,20)$.
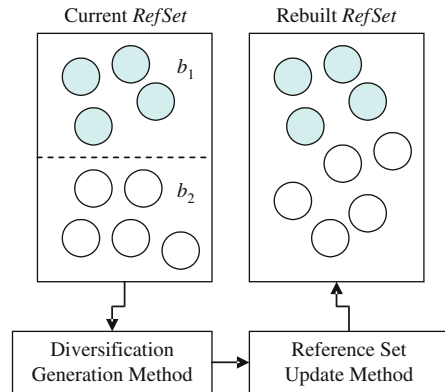
Fig. 3.8 *RefSet* rebuilding

The experiments revealed that a significant change in the solution quality is due to the increase in *PSize*. The experiments were inconclusive about the advantage

of increasing the size of the reference set (i.e., the value of b) beyond 20 when *PSize* is not greater than 100. The experiments showed that the best results are obtained when $b_1 = b_2$. Therefore the key parameters were set as *PSize* $= 100$, $b = 20$ and $b_1 = b_2 = 10$.

The second experiment was performed to learn about the ranks of the reference solutions that generated the best solutions found in the search process. To this end it was tracked from which rank or position in the reference set the best solutions came from. Say e.g., the overall best solution came from combining the 3rd and 5th best solutions, where one of these came from combining the 1st, 2nd and 6th best solutions, and the other came from ... etc. This trace would give an idea of which solutions are important as components of others.

Figure 3.9 shows the results of this experiment. Its interpretation is as follows. Consider the line associated with rank 1. Then, the count of (almost) 18 in the first point of this line indicates that rank 1 solutions were generated (approximately) 18 times from other rank 1 solutions. Similarly, rank 1 solutions were generated 8 times from rank 2 solutions. The decaying effect exhibited by all the lines indicate that high quality solutions tend to generate new solutions that are admitted to the reference set. This is evident by the counts corresponding to rank 1 in the *x*-axis. During the search, on average, rank 1 solutions generated 18 rank 1 solutions, 4 rank 10 solutions and 1 rank 20 solution.

Fig. 3.9 Source ranks of reference solutions

In the third experiment the SS algorithm is compared with some of the state-of-the-art meta-heuristics for the LOP on the input-output matrices, SGB instances and random problems A (types I and II). In particular, we report about a comparison of SS (with its parameters set as specified above) with the method CK of [26] and the improved GRASP method LS (described in Sect. 3.2). Again, CK-10 denotes the application of CK from 10 randomly generated initial solutions.

Table 3.10 shows, for each method, the average percent deviation from optimality (or from the best known solutions), the number of optimum (or best) solutions, and the average CPU time (seconds on a Pentium 166 MHz).

Table 3.10 Comparison of best methods

	LS	CK	CK-10	TS	SS
Input-output instances					
Deviation	0.15	0.15	0.02	0.04	0.01
No. of opt. solutions	11	11	27	33	42
CPU time	0.01	0.10	1.06	0.49	2.35
SGB instances					
Deviation	0.31	0.05	0.01	0.01	0.01
No. of opt. solutions	0	0	0	5	18
CPU time	0.10	2.73	31.93	3.03	14.05
Random A type I					
Deviation	0.45	0.48	0.29	0.11	0.02
No. of best solutions	0	0	0	5	21
CPU time	0.08	6.90	67.12	12.74	63.82
Random A type II					
Deviation	0.02	0.02	0.01	0.00	0.00
No. of best solutions	0	0	0	41	14
CPU time	0.07	4.30	44.51	7.91	43.36

Table 3.10 shows that the local search procedure LS is clearly inferior in terms of solution quality, although the simplicity of the approach remains appealing. The performance of LS and CK is very similar within each of the four problem sets, but their deviation from the optimum (or best) solutions is significantly higher in the case of the random instances type I. Both LS and CK were started from a randomly generated solution. TS and SS, on the other hand, are quite robust, as is evident by the negligible change in the deviation values across problem sets.

It is difficult to measure solution quality in terms of percentaged deviation, since TS and SS have very small average deviations from optimality. In terms of the number of optima (or best solutions), TS is very competitive, considering that it is able to find 33 optima for the input-output instances and 41 best solutions for random type II problems. The most robust method is SS in terms of number of optima or best solutions found. However, this is achieved at the expense of higher computation times.

Finally, we report the results obtained when applying SS for 10 seconds to the 229 instances in the OPT-I set. Table 3.11 reports the #Instances, *Dev* and *#Opt* statistics in each subset of instances.

Table 3.11 Scatter Search on OPT-I instances

	IO	SGB	RandAII	Rand B	MB	Special	**Total**
#Instances	50	25	25	70	30	29	229
Dev(%)	0.01	0.00	0.01	0.01	0.00	0.09	0.02
#Opt	42	20	13	62	16	12	165

The results in Table 3.11 clearly show that the Scatter Search methodology obtains high quality solutions on the OPT-I set. The average percentaged deviation overall is 0.02% and the number of optimum solutions is 165 out of 229 instances.

3.7 Genetic Algorithms

The idea of applying the biological principle of natural evolution to artificial systems, introduced more than three decades ago, has seen impressive growth in the past few years [98].

Usually grouped under the term *evolutionary algorithms* or *evolutionary computation*, we find the domains of *genetic algorithms*, *evolution strategies*, *evolutionary programming*, and *genetic programming*. Evolutionary algorithms have been successfully applied to numerous problems from different domains, including optimization, automatic programming, machine learning, economics, ecology, population genetics, studies of evolution and learning, and social systems.

A genetic algorithm is an iterative procedure that consists of a constant size population of individuals, each one represented by a finite string of symbols, known as the genome, encoding a possible solution in a given problem space. This space, referred to as the search space, comprises all possible solutions to the problem at hand. Generally speaking, genetic algorithms are applied to spaces that are too large to be exhaustively searched (such as those in combinatorial optimization). Solutions to a problem were originally encoded as binary strings due to certain computational advantages associated with such encoding. Also the theory about the behavior of algorithms was based on binary strings. Because in many instances it is impractical to represent solutions using binary strings, the solution representation has been extended in recent years to include character based encoding, real-valued encoding, and tree representations.

The standard genetic algorithm [39] proceeds as follows: an initial population of individuals is generated at random or heuristically. In every evolutionary step, denoted as a generation, the individuals in the current population are decoded and evaluated according to some predefined quality criterion, referred to as their fitness (evaluated by a fitness function). To form a new population (the next generation), individuals are selected according to their fitness. Many selection procedures are currently in use, one of the simplest being Holland's original fitness-proportionate selection, where individuals are selected with a probability proportional to their relative fitness. This ensures that the expected number of times an individual is chosen

is approximately proportional to its relative performance in the population. Thus, high fitness ("good") individuals stand a better chance of "reproducing", while low fitness ones are more likely to disappear.

The *roulette wheel selection* [36] is a common implementation of a proportional selection mechanism. In this selection process, each individual in the population is assigned a portion of the wheel proportional to the ratio of its fitness and the population's average fitness.

Fig. 3.10 One point crossover

Tournament and *ranking* are two other popular selection techniques. Tournament selection consists of choosing q individuals from a population of n individuals and selecting the best according to the fitness value to survive into the next generation. Hence, n tournaments are necessary to build the population for the next generation. Binary tournaments, for which $q = 2$, are the most common implementation of this selection technique. Ranking ignores the fitness values and assigns selection probabilities based exclusively on rank. Genetically inspired operators are used to introduce new individuals into the population, i.e., to generate new points in the search space. The best known such operators are *crossover* and *mutation*. Crossover is performed, with a given probability (the crossover probability or crossover rate), between two selected individuals, called parents, by exchanging parts of their genomes (i.e., encoding) to form two new individuals, called offspring; in its simplest form, substrings are exchanged after a randomly selected crossover point. This operator tends to enable the evolutionary process to move toward promising regions of the search space. Figure 3.10 depicts a one point crossover operation.

The mutation operator is introduced to prevent premature convergence to local optima by randomly sampling new points in the search space. Mutation entails flipping bits at random, with some (small) probability. Figure 3.11 shows a graphical representation of the mutation operator.

Fig. 3.11 Mutation operation

The principle layout of a genetic algorithm is the following.

Genetic algorithm

(1) Create the initial population P by randomly generating n solutions.
(2) While the stopping condition is not met:

    (2.1) Evaluate the solutions in P and update the best solution found if necessary. (**Evaluate**)
    (2.2) Calculate the probability of surviving based on solution quality. Evolve P by choosing n solutions according to their probability of surviving. (**Survival of the fittest**)
    (2.3) Select a fraction pc of the solutions in P to be combined. Selection is at random with the same probability for each element of P. The selected elements are randomly paired for combination, with each pair generating one or more offsprings that are added to P. (**Combine**)
    (2.4) A fraction pm of the solutions in P is selected for mutation. The mutated solution is improved and added to P. (**Mutate**)

(3) Output the best solution in P.

Genetic algorithms are stochastic iterative processes that are not guaranteed to converge in practice; the termination condition may be specified as some fixed, maximal number of generations or as the attainment of an acceptable fitness level for the best individual. However, contrary to other meta-heuristics, theoretical studies [39], based on Markov chains and the so-called *schema theorem*, have been developed to establish the convergence conditions to a global optimum with respect to the selection strategies and operators of the method.

In [27] a GA for the LOP based on the following three elements is proposed:

(1) A selection procedure. The n individuals in the current population are sorted according to their objective value (where the best one comes first). Then, the probability to choose an individual for combination is proportional to its rank in this ordering (i.e., the best individual has a selection probability n times the probability of the worst one).
(2) A crossover operator. Two selected solutions (individuals) are "crossed" or combined to obtain a new one. Specifically, given $u = \langle u_1, u_2, \ldots, u_n \rangle$ and

$v = \langle v_1, v_2, \ldots, v_n \rangle$, we first compute the auxiliary array $a = (a_1, a_2, \ldots, a_n)$ where $a_i = u_i + v_i$. Then, we order the elements in a from the lowest to the highest: $a_{\sigma(1)} \leq a_{\sigma(2)} \leq \ldots \leq a_{\sigma(n)}$. The offspring w is constructed as

$$w = \langle a_{\sigma(1)}, a_{\sigma(2)}, \ldots, a_{\sigma(n)} \rangle.$$

(3) A mutation operator. Regardless of the objective function, a selected individual is mutated by applying a simple transformation.

Table 3.12 reports the results obtained with our implementation of a classical GA for the LOP according to the three elements above. This table shows the #Instances, *Dev* and *#Opt* statistics in each subset of instances in the OPT-I set. As in our previous experiments, we run the method for 10 seconds on each instance.

Table 3.12 Genetic Algorithm on OPT-I instances

	IO	SGB	RandAII	Rand B	MB	Special	**Total**
#Instances	50	25	25	70	30	29	229
Dev(%)	0.38	0.76	21.44	0.75	35.53	1.38	10.04
#Opt	9	0	0	1	0	4	14

The results in Table 3.12 show that the GA obtains relatively low quality results on the OPT-I instances. It is only able to obtain 14 optimum solutions out of 229 instances. Note that the methods reported in the previous sections obtain a larger number of optimum solutions (215 in the case of the tabu search).

In [116] a genetic algorithm coupled with a local search procedure is developed for the LOP. This hybrid method is called *memetic algorithm*. In the initialization, a population of individuals is obtained by first generating a set of random permutations (solutions) and then applying a local search procedure to each of them. The local procedure is based on insertions where the neighborhood is examined in random order and the first improving move is performed. This method is very similar to the local search procedure implemented for the LOP in other meta-heuristics described in previous sections.

In each iteration of the algorithm, called *generation*, new solutions are generated by applying crossover and mutation to randomly selected solutions in the population (according to a uniform distribution). The crossover operator takes two individuals of the current population and combines them into a new individual, while the mutation operator introduces a perturbation into an individual. Local search is applied again to improve each new solution. The new population is created by merging the best solutions in the population and the new improved solutions. It is worth mentioning that the authors consider four different crossover operators: DPX (similar distance from parents), CX (classical crossover), OB (order based crossover) and Rank (computing the average ranking of the elements). In computational experiments CX and OB performed best.

Finally, if the average of the objective function value of the population does not change in a certain number of consecutive generations, the population is rebuilt in a similar way as the Reference Set is rebuilt in scatter search: the best solutions are kept and the worst solutions are replaced with new ones.

Table 3.13 shows the results obtained with the memetic algorithm [116] on the OPT-I instances. We run the method for 10 seconds and report the #Instances, *Dev* and *#Opt* statistics in each subset of instances.

Table 3.13 Memetic Algorithm on OPT-I instances

	IO	SGB	RandAII	Rand B	MB	Special	**Total**
#Instances	50	25	25	70	30	29	229
Dev(%)	0.00	0.00	0.00	0.00	0.00	0.00	0.00
#Opt	50	25	25	70	30	28	228

Table 3.13 clearly shows that the MA obtains the best results so far on the OPT-I instances (it only misses one optimum solution out of 229 instances). Comparing these results with those obtained with a classical GA method (reported in Table 3.12) we can conclude that the inclusion of a local search makes an important difference in a meta-heuristic procedure.

In [70] a similar method, based on combining a classical GA with a local search is presented. Is is called *hybrid genetic algorithm* and it is very similar to the method in [116]. The local search is also based on exchanges and it also applies the CX and OB crossover operators. However, instead of DPX and Rank, it applies PMX (partially matched crossover). In their experimentation the authors conclude that GA without a local search produces low quality results, but its hybridization with local search is able to match state-of-the-art methods.

We can establish that both scatter search and genetic algorithms belong to the family of population based meta-heuristics. Moreover, they were both proposed in the seventies: while Holland [68] introduced genetic algorithms and the notion of imitating nature and the "survival of the fittest" paradigm, Glover [53] introduced scatter search as a heuristic for integer programming that expanded on the concept of surrogate constraints. Both methods are based on maintaining and evolving a population of solutions throughout the search. Although the population based approach makes SS and GA similar in nature, as described in [85], there are fundamental differences between these two methodologies:

– The population in genetic algorithms is about one order of magnitude larger than the reference set in scatter search. A typical GA population size is 100, while a typical SS reference set size is 10.
– A probabilistic procedure is used to select parents to apply crossover and mutation operators in GAs while the combination method is applied to a predetermined list of subsets of reference solutions in scatter search.
– The evolution of a GA population follows the "survival of the fittest" philosophy, which is implemented using probabilistic rules. In scatter search, changes

in the reference set are controlled by the deterministic rules in the reference set update method.

– The use of local search procedures is an integral part of scatter search, while it was added to GAs in order to create hybrid implementations that would yield improved outcomes.

– The subset generation method considers combinations of more than two solutions. GAs are typically restricted to combining two solutions.

– Full randomization is the typical mechanism used in GAs to build the initial population. Diversity in scatter search is achieved with the application of the diversification generation method, which is designed to balance diversity with solution quality.

Although GA and SS have contrasting views about searching a solution space, it is possible to create a hybrid approach without entirely compromising the scatter search framework. Specifically, if we view the crossover and mutation operators as an instance of a combination method, it is then straightforward to design a scatter search procedure that employs genetic operators for combination purposes.

3.8 Empirical Comparison

In this section we compare the meta-heuristics described in the previous sections. Specifically, we consider the following methods:

– TS: Tabu Search
– MA: Memetic Algorithm
– VNS: Variable Neighbourhood Search
– SA: Simulated Annealing
– SS: Scatter Search
– GRASP: Greedy ramdomized adaptive search procedure
– GA: Genetic Algorithm

We divide our experimentation into two parts according to the classification of the instances introduced in Chap. 1. Table 3.14 reports the results on the 229 OPT-I instances and Table 3.15 reports those on the 255 UB-I instances.

In each experiment we compute for each instance and each method the relative deviation *Dev* (in percent) between the best solution value obtained with the method and the optimal value for that instance (in the UB-I instances we do not know the optimal value and therefore we instead consider the best known value). For each method, we also report the number of instances *#Opt* for which an optimum solution could be found (*#Best* in the case of UB-I instances). In addition, we calculate the so-called *score statistic* [114] associated with each method. For each instance, the *nrank* of method *M* is defined as the number of methods that found a better solution than the one found by *M*. In the event of ties, the methods receive the same *nrank*, equal to the number of methods strictly better than all of them. The value of *Score* is

the sum of the *nrank* values for all the instances in the experiment. Thus the lower the *Score* the better the method.

Table 3.14 Meta-heuristics on OPT-I instances

	TS	MA	VNS	SA	SS	GRASP	GA
IO							
Dev(%)	0.00	0.00	0.00	0.03	0.01	0.00	0.38
Score	50	50	50	270	90	57	343
#Opt	50	50	50	16	42	49	9
SGB							
Dev(%)	0.00	0.00	0.00	0.03	0.00	0.00	0.76
Score	25	25	25	197	47	31	220
#Opt	25	25	25	0	20	23	0
RandAII							
Dev(%)	0.00	0.00	0.00	0.08	0.01	0.01	21.44
Score	25	25	47	197	80	116	242
#Opt	25	25	19	0	13	5	0
RandB 70							
Dev(%)	0.00	0.00	0.02	0.25	0.01	0.00	0.75
Score	70	70	103	486	115	70	607
#Opt	70	70	64	10	62	70	1
MB							
Dev(%)	0.00	0.00	0.00	1.33	0.00	0.00	35.53
Score	30	30	30	240	98	71	295
#Opt	30	30	30	0	16	21	0
Special							
Dev(%)	0.02	0.00	0.11	0.39	0.09	0.05	1.38
Score	64	30	66	166	109	84	226
#Opt	15	28	20	6	12	14	4
OPT-I							
Dev(%)	0.00	0.00	0.02	0.35	0.02	0.01	10.04
#Opt	215	228	208	32	165	182	14

Table 3.14 shows that most of the meta-heuristics considered are able to obtain all the optimal solutions within the time limit of 10 seconds considered (they actually obtain it in around 1 second). We therefore conclude that instances in OPT-I are easy for the best meta-heuristics and therefore not adequate to compare them.

In our second experiment we target the UBI-instances for which the optimum is not known but we have an upper bound for comparison. We therefore compute for each instance and each method the relative deviation *D.Best* (in percent) between the best solution value *Value* obtained with the method and the best known value *BestValue* as well as the relative deviation *D.UB* (in percent) between *Value* and the upper bound. For each method, we also report the number of instances *#Best* for which the value of the solution is equal to *BestValue*. As in the previous experiment we calculate the score statistic. Table 3.15 reports the values of these four statistics on the UB-I instances when running the 7 meta-heuristics for 10 seconds.

Table 3.15 Meta-heuristics on UB-I instances

	TS	MA	VNS	SA	SS	GRASP	GA
RandAI							
D.Best	0.12	0.05	0.47	1.77	0.26	0.42	10.59
D.UB	17.81	17.75	18.10	18.88	17.92	18.05	26.28
Score	201	105	482	641	326	461	936
#Best	5	33	0	0	1	0	0
RandAII							
D.Best	0.01	0.00	0.01	0.07	0.02	0.04	35.97
D.UB	0.38	0.38	0.39	0.44	0.40	0.41	36.21
Score	63	25	74	175	109	151	191
#Best	3	39	8	0	0	0	0
RandB							
D.Best	0.00	0.00	0.00	0.31	0.04	0.00	0.91
D.UB	3.20	3.20	3.26	3.51	3.24	3.20	4.08
Score	20	20	67	160	50	20	175
#Best	20	20	11	0	11	20	0
XLOLIB							
D.Best	0.62	0.12	0.42	0.53	0.68	1.14	23.99
D.UB	3.21	2.72	3.01	3.13	3.27	3.72	25.96
Score	307	87	200	266	320	460	758
#Best	0	2	0	0	0	0	0
Special							
D.Best	0.43	0.03	0.50	2.05	0.32	0.65	9.27
D.UB	9.61	9.26	9.67	11.04	9.52	9.81	17.35
Score	21	7	27	57	17	28	63
#Best	3	4	2	0	3	3	0
UB-I							
D.Best	0.23	0.04	0.28	0.95	0.26	0.45	16.15
D.UB	6.84	6.66	6.89	7.40	6.87	7.04	21.98
#Best	31	98	21	0	17	23	0

The results in Table 3.15 show that MA is able to obtain the largest number of best solutions (98 of a total of 255 instances) in short runs (10 seconds). No other method is able to obtain more than 31 best solutions, which clearly indicates the superiority of MA. On the other hand, considering average percentage deviations with respect to the best solutions, the differences among the methods appear to be very small. MA presents on average a deviation of 0.04% while TS, SS and VNS present averages deviations of 0.23% 0.26% and 0.28%, respectively. This indicates that although these methods are not able to match the best solution values, they obtain solutions with values very close to the best.

According to the differences among methods observed in Table 3.15, where the deviations w.r.t. the best solution known range from 0.04% to 17.72%, we can conclude that the instances in set UB-I are more difficult to solve than those in OPT-I (where the deviations range from 0.00% to 15.31%).

Chapter 4
Branch-and-Bound

Abstract We now turn to the discussion of how to solve the linear ordering problem
to (proven) optimality. In this chapter we start with the branch-and-bound method
which is a general procedure for solving combinatorial optimization problems. In
the subsequent chapters this approach will be realized in a special way leading to
the so-called *branch-and-cut method*. There are further possibilities for solving the
LOP exactly, e.g. by formulating it as *dynamic program* or as *quadratic assignment
problem*, but these approaches did not lead to the implementation of practical algo-
rithms and we will not elaborate on them here.

4.1 Introduction

Combinatorial optimization deals with a special type of mathematical optimization
problem with the property that the set of feasible solutions is finite. In its most
general form such a problem is defined on a finite set $\mathscr{I}$ (set of feasible solutions)
and a function $f : \mathscr{I} \to \mathbf{R}$ has to be optimized. Since the set of feasible solutions $\mathscr{I}$
is finite, the problem could in principle be solved by enumeration. However, the
number of feasible solutions can be very large, thus prohibiting this approach in
general.

Branch-and-bound tries to deal with these many feasible solutions in a systematic
way. Basically, it is a divide-and-conquer approach that tries to solve the original
problem by splitting it into smaller problems for which upper and lower bounds
are computed and may be employed to exclude large parts of the solution set from
further consideration.

Of course, the general definition of a combinatorial optimization problem given
above is of no use unless we have a reasonable characterization of $\mathscr{I}$ and an algo-
rithmic way of evaluating the objective function.

For many problems the objective function can be defined in a simple way and
they can be formulated as follows. (2^E denotes the power set of E.)

Definition 4.1. Let the finite set $E = \{e_1, e_2 \ldots, e_n\}$ (ground set) and $\mathscr{I} \subseteq 2^E$ (set of feasible solutions) be given. Assume that there is a function $c : E \to \mathbf{R}$ (objective function) such that the value of a feasible solution $F \in \mathscr{I}$ is given as $c(F) = \sum_{e \in F} c(e)$. The *linear combinatorial optimization problem* $(E, \mathscr{I}, c)$ consists of finding $F \in \mathscr{I}$ such that $c(F)$ is as large as possible.

The LOP fits into this scheme by setting $E = A_n$, $\mathscr{I} = \{T \subset A_n \mid T$ is an acyclic tournament$\}$, and $c((i,j)) = c_{ij}$. By complementing the function c we can also deal with minimization problems.

The crucial part of a successful branch-and-bound algorithm is the computation of upper bounds for subproblems. Here one uses the fundamental concept of relaxation.

Definition 4.2. Suppose that two combinatorial problems $(E, \mathscr{I}, f)$, $(E', \mathscr{I}', f')$ and an injective function $\varphi : E \to E'$ are given. The problem $(E', \mathscr{I}', f')$ is a *relaxation* of $(E, \mathscr{I}, f)$, if $\varphi(F) \in \mathscr{I}'$ and $f(F) = f'(\varphi(F))$, for all $F \in \mathscr{I}$.

(More generally one can define that a problem $\max\{f(x) \mid x \in T\}$ is a relaxation of the problem $\max\{c(x) \mid x \in X\}$ if $X \subseteq T$ and $f(x) \geq c(x)$, for all $x \in X$.)

Hence a solution of the relaxed problem gives an upper bound on the optimum objective function value of the problem it was derived from. The tighter the relaxation, the better this bound will be. Of course, a relaxation is only useful if it can be treated efficiently by optimization algorithms.

Branch-and-bound can be outlined as follows.

Branch-and-Bound Algorithm

(1) Initialize the list of active subproblems with the original problem.

(2) If the list of active subproblems is empty, `Stop` (the best feasible solution found so far is optimal).

(3) Otherwise, choose some subproblem from the list of active problems and "solve" it as follows:

    (3.1) Find an optimal solution for the subproblem, or

    (3.2) prove that the subproblem has no feasible solution, or

    (3.3) prove that there is no feasible solution for the subproblem that has larger objective value than the best feasible solution that is already known, or

    (3.4) split the subproblem into further subproblems and add them to the list of active problems.

(4) Goto step (2).

The splitting of problems into subproblems can be represented by the so-called *branch-and-bound tree*, the root of which represents the original problem.

It is crucial for the efficiency of a branch-and-bound algorithm that the branch-and-bound tree does not grow too large. Already, if problems are only split into

two subproblems, the number of subproblems grows exponentially fast. Therefore subproblems have to be solved if possible by alternatives (3.1), (3.2) or (3.3). Alternative (3.1) rarely occurs, and relaxations are important for (3.2) and (3.3). With respect to (3.2), if a relaxation of the subproblem is already infeasible, then the subproblem itself is also infeasible. To be able to fathom the subproblem using (3.3), good lower and upper bounds must be available. Lower bounds are obtained by finding feasible solutions. These are computed either by solving some subproblem to optimality or more often by determining good feasible solutions using one of the many heuristics we have discussed in the previous chapters. Upper bounds can be computed by using relaxations where in principle any type of relaxation can be employed. It is clear that the stronger the bounds obtained by a relaxation are, the better the performance of the algorithm will be. Without a suitable relaxation branch-and-bound tends to completely enumerate the set of feasible solutions and thus becomes unsuitable for practical computations.

Besides *bounding*, the second component of this approach is *branching* which denotes the splitting of the current subproblem into a collection of new subproblems whose union of feasible solutions contains all feasible solutions of the current subproblem. For 0/1-problems, the simplest branching rule consists of defining two new subproblems: in the first subproblem some chosen variable is required to have the value 1 in every feasible solution and in the second one to have the value 0. Other branching strategies are possible. There are also several heuristics for choosing the next subproblem to be considered.

4.2 Branch-and-Bound with Partial Orderings

One of the earliest branch-and-bound algorithms for the LOP was proposed by de Cani [23] in 1972. He encountered the linear ordering problem when studying procedures to obtain a ranking of n objects on the basis of a number of pairwise comparisons. The algorithm successively constructs partial rankings of more and more objects and tries to prove by upper bounds that some partial rankings need not be extended for finding optimum solutions.

An upper bound for the optimum solution value of a LOP is clearly given by

$$z_0 = \sum_{i=1}^{n-1} \sum_{j=i+1}^{n} \max\{c_{ij}, c_{ji}\}.$$

If we require that object i is to be ranked before object j then an upper bound is given by

$$z(i,j) = z_0 - \max\{c_{ij}, c_{ji}\} + c_{ij}.$$

Partial rankings are built up, and each branching operation in the tree corresponds to inserting a further object at some position in the partial ranking. For each partial

ranking an upper bound for the best possible extension to a complete ranking can be calculated.

Suppose that an ordering of k objects (w.l.o.g. objects $1, \ldots, k$), say $\langle 1, 2, \ldots, k \rangle$, is given. If we insert object $k+1$ at position l, $1 \le l \le k+1$, of this ordering then, as an upper bound on the objective funtion value for the linear orderings containing the respective partial ordering of $k+1$ objects, we get

$$z(1, \ldots, l-1, k+1, l, l+1, \ldots, k) =$$
$$z(1, \ldots, k) + \sum_{r=1}^{l-1} c_{r,k+1} + \sum_{r=l}^{k} c_{k+1,r} - \sum_{r=1}^{k} \max\{c_{k+1,r}, c_{r,k+1}\}.$$

We start the branch-and-bound algorithm by arbitrarily choosing two of the objects, say i and j, and calculate $z(i, j)$ and $z(j, i)$. We generate two nodes, the first one corresponding to the partial ordering ranking i before j, and the second one to ranking j before i. Then we proceed at the node with the larger upper bound value. Suppose that we are at a node of level k, $1 < k < n$, of the tree (assuming that the root node is on level 1). Then k objects are already partially ordered. From the remaining $n - k$ objects we select one (according to some rule). The $k+1$ upper bounds obtained by inserting the new object at each possible position are calculated, and we proceed in that branch of the tree corresponding to the largest of these values. At level n a complete ranking of the objects is found.

The upper bounds are exploited in the usual way for backtracking and excluding parts of the tree from further consideration.

Figure 4.1 visualizes the development of the branch-and-bound tree with this approach. Nodes are labeled with partial ordering relations. Note that it is not a binary tree, the number of branches increases with the level.

Fig. 4.1 Construction of the branch-and-bound tree

4.3 Lexicographic Search

Another one of the early methods for the solution of linear ordering problems is the *lexicographic search algorithm* proposed in 1968 by Korte and Oberhofer [82, 83]. It is actually not a branch-and-bound algorithm and can better be characterized as an *enumeration scheme*. Korte and Oberhofer were interested not only in an optimum triangulation of an input-output matrix, but also in the number of optima and the number of so-called relatively optimum solutions.

Definition 4.3. A matrix $C = (c_{ij})$ satisfies *Helmstädter's conditions* if

(i) $\sum_{l=i}^{k} c_{il} \geq \sum_{l=i}^{k} c_{li}$, for all $i < k$, and

(ii) $\sum_{l=k}^{i} c_{li} \geq \sum_{l=k}^{i} c_{il}$, for all $k < i$.

Such matrices are also called *relatively optimum*.

These conditions were given by Helmstädter [65]. Of course, every optimally triangulated matrix satisfies *Helmstädter's conditions*. If one of the conditions is violated, then a simple reordering can improve the objective function. If, for example, condition (i) is violated for some i and k, then it would be profitable to change the subsequence $\langle i, i+1, \dots, k \rangle$ to $\langle i+1, \dots, k, i \rangle$.

The lexicographic search algorithm enumerates all permutations of the n objects by fixing at level k of the enumeration tree the k-th position of the permutations. More precisely, if a node at level k is generated, then the first k positions $\sigma(1), \dots, \sigma(k)$ are fixed. Based on this fixing, several of Helmstädter's conditions can be checked. If one is violated then no relatively optimum solution has $\sigma(1), \dots, \sigma(k)$ in the first k positions. The node can be ignored, and a backtracking operation is performed.

Figure 4.2 shows part of the enumeration tree for 4 objects. Nodes at level k show the fixings of the first k positions. (Here the root node is at level 0.)

Since there is no bounding according to objective function values eventually all relatively optimum solutions are enumerated, and the best ones among them are optimum. Korte and Oberhofer applied their algorithm to get information about the distribution of relatively and absolutely optimum solutions for real-world and random problems.

4.4 Extension of Lexicographic Search to Branch-and-Bound

Possiblities for the derivation of bounds are already mentioned in [83], although not implemented. Recall that we may w.l.o.g. assume that all matrix entries c_{ij} are nonnegative.

Fig. 4.2 Lexicographic search tree

Suppose that we have fixed the first k positions $\sigma(1), \ldots, \sigma(k)$ of a permutation. Let $I = \{\sigma(1), \ldots, \sigma(k)\}$ and $J = \{1, \ldots, n\} \setminus I$. If σ is extended to some complete permutation of $\{1, \ldots, n\}$ we have

$$\sum_{i=1}^{n-1} \sum_{j=i+1}^{n} c_{\sigma(i)\sigma(j)} = \sum_{i=1}^{k-1} \sum_{j=i+1}^{k} c_{\sigma(i)\sigma(j)} + \sum_{i\in I} \sum_{j\in J} c_{ij} + \sum_{i=k+1}^{n-1} \sum_{j=i+1}^{n} c_{\sigma(i)\sigma(j)}.$$

The first two terms on the right hand side of this equation do not depend on the permutation of J. An upper bound for the third term (which amounts to determining an upper bound for the triangulation problem with respect to the ground set J) yields an upper bound for all possible extensions.

Kaas [75] developed a lexicographic branch-and-bound scheme based on the algorithm of [83]. He also applies Helmstädter's conditions to rule out solutions which are not relatively optimum, but in addition uses this equation to compute upper bounds (heuristics are used for bounding the third term from above).

Note that the sum of the subdiagonal entries of the matrix obtained by permuting the rows and columns according to σ is calculated by

$$\sum_{i=1}^{n-1} \sum_{j=i+1}^{n} c_{\sigma(j)\sigma(i)} = \sum_{i=1}^{k-1} \sum_{j=i+1}^{k} c_{\sigma(j)\sigma(i)} + \sum_{i\in I} \sum_{j\in J} c_{ji} + \sum_{i=k+1}^{n-1} \sum_{j=i+1}^{n} c_{\sigma(j)\sigma(i)}.$$

The sum of the first two right-hand side terms gives a lower bound on the sum of the subdiagonal entries of any extension and can thus be used to compute an upper bound on the sum of the superdiagonal entries.

There are more authors who have formulated a branch-and-bound method for this problem (e.g. [64, 87, 108]) based on the same ideas as outlined above.

4.5 Branch-and-Bound with Lagrangian Relaxation

This is a more recent branch-and-bound approach of Charon and Hudry [29] where Lagrangian relaxation techniques are used for bound computations.

Define 0/1 variables x_{ij}, $1 \leq i, j \leq n$, $i \neq j$, where $x_{ij} = 1$ if i is ranked before j, and $x_{ij} = 0$ otherwise. Then the LOP can be formulated as the linear 0/1 program

$$\max \sum_{(i,j) \in A_n} c_{ij} x_{ij}$$
$$x_{ij} + x_{jk} + x_{ki} \leq 2, \text{ for all distinct nodes } i, j, k \in V_n,$$
$$x_{ij} + x_{ji} = 1, \text{ for } 1 \leq i < j \leq n,$$
$$x_{ij} \in \{0,1\}, \text{ for } 1 \leq i, j \leq n, i \neq j.$$

This is the canonical IP formulation of the LOP and we will discuss it in more detail in the next chapters.

Lagrangian relaxation removes constraints from the original problem and penalizes their violation in a modified objective function.

Let T be the set of all triples (r, s, t) of three distinct nodes where $r < s$ and $r < t$. The following Lagrangian relaxation is used in [29] for a given vector $\mu \geq 0$ of Lagrangian multipliers.

$$L(\mu) = \max \sum_{(i,j) \in A_n} c_{ij} x_{ij} + \sum_{(r,s,t) \in T} (2 - x_{rs} - x_{st} - x_{tr}) \mu_{rst}$$
$$x_{ij} + x_{ji} = 1, \text{ for } 1 \leq i < j \leq n,$$
$$x_{ij} \in \{0,1\}, \text{ for } 1 \leq i, j \leq n, i \neq j.$$

This problem can be solved trivially (more details below) and obviously gives an upper bound on the optimum objective function value of the LOP. Note that this relaxation does not exactly meet the requirements of Definition 4.2. Here we have $f(F) \leq f'(\varphi(F))$, but the upper bound property holds as well.

The best such bound can be found by solving the *Lagrangian dual problem*

$$\min_{\mu} L(\mu).$$

To this end one uses so-called *subgradient* or *bundle methods* which are able to find good approximations of the best bound. In theory, they could compute this bound exactly, but then a very slow convergence of the step sizes of the subgradient algorithm to 0 is required. In practice, step sizes are decreased faster and the optimum bound is therefore not met exactly.

With the notation $(i, j) \in (r, s, t)$ if (i, j) is one of the edges (r, s), (s, t) or (t, r) the objective function for computing $L(\mu)$ can be rewritten as

$$\max \sum_{(i,j)\in A_n} c_{ij}x_{ij} + \sum_{(i,j,k)\in T} (2 - x_{ij} - x_{jk} - x_{ki})\mu_{ijk}$$

$$= \max \sum_{(i,j)\in A_n} \Big(c_{ij} - \sum_{(r,s,t)\in T,(i,j)\in(r,s,t)} \mu_{rst}\Big)x_{ij} + 2 \sum_{(r,s,t)\in T} \mu_{rst}$$

$$= \max \sum_{(i,j)\in A_n} d_{ij}x_{ij} - C,$$

where d_{ij} is set according to the second-to-last line and C is constant. In this form it is clear that finding the optimum in 0/1 variables with only $x_{ij} + x_{ji} = 1$ as constraints is trivial.

We now describe a prototypical realization of a subgradient approach for approximating $\min_\mu L(\mu)$.

LOPSubgradient(D_n, c)

(1) Let τ be an initial step size and $0.9 \leq \alpha < 1$ a decrement factor.
(2) Set $t_1 = \tau$, $\mu_{rst}^0 = 0$ for every $(r,s,t) \in T$ and $k = 1$.
(3) While $t_k > \varepsilon$:

    (3.1) Compute $L(\mu)$ by setting $x_{ij} = 1$, if $d_{ij} > d_{ji}$, and $x_{ij} = 0$, otherwise, for $i < j$, and set $x_{ji} = 1 - x_{ij}$.

    (3.2) Define d^k by setting $d_{rst}^k = 2 - x_{rs} - x_{st} - x_{tr}$, for $(r,s,t) \in T$.

    (3.3) Set $\mu_{rst}^{k+1} = \mu_{rst}^k + t_k \times d_{rst}^k$, for every $(r,s,t) \in T$. If $\mu_{rst}^{k+1} < 0$, then set $\mu_{rst}^{k+1} = 0$.

    (3.4) Set $t_{k+1} = \alpha t_k$ and increment k.

(4) Return the best bound found.

We do not want to introduce the background of nondifferentiable optimization here. We just note that d^k is a so-called *subgradient* and that with α very close to 1 and ε small, $\min_\mu L(\mu)$ is usually very well approximated at the expense of, however, considerable running time. The decrease of the bounds is not monotone, and therefore in (4) the best bound found is returned. Bundle methods extend this principle and compute combinations of subgradients of several iterations for the direction of the next step. In practice, they are much more powerful.

We illustrate this (simple) subgradient algorithm with two applications for finding an upper bound for the LOLIB instance be75np. In the first run (Fig. 4.3) α is set to 0.999999, while in the second run (Fig. 4.4) it is set to 0.99. The figures show the development of the upper bounds for the first 100 iterations. The first run terminates after 2 hours giving the upper bound 790989 (which as will be seen in the next chapter is the minimum of $L(\mu)$), the second run terminates after only 5 seconds giving the upper bound 799467. The decrease of the upper bound is fast in the first iterations. After 100 iterations we have bound 800077 in the first run and 801963 in the second run, but to reach bounds near the optimum substantial running time has to be invested.

**Fig. 4.3** Simple subgradient method with slow convergence

Fig. 4.4 Simple subgradient method with fast convergence

This is just a simple example for the general behaviour of subgradient methods. In [29] a variant that finds a good compromise between fairly fast convergence and good quality of bounds is developed. With this branch-and-bound method, instances of moderate sizes can be solved successfully. As the heuristic for finding good linear orderings, the meta-heuristic *noising method* is used which was described in Chap. 3.

The algorithm is publically available [29] and we have applied it (with standard parameter settings) to a set of random benchmark instances with $n = 40$. Table 4.1

Table 4.1 Branch-and-bound with Lagrangian relaxation

Problem	Optimum	#nodes	CPU
p40-01	29457	4261	0:06
p40-02	27482	444285	6:13
p40-03	28061	29660	0:31
p40-04	28740	267766	3:51
p40-05	27450	639401	9:26
p40-06	29164	79270	1:09
p40-07	28379	226125	3:27
p40-08	28267	103477	1:31
p40-09	30578	431607	5:47
p40-10	31737	86698	0:59
p40-11	30658	333902	4:16
p40-12	30986	96651	1:27
p40-13	33903	13423	0:09
p40-14	34078	73839	0:57
p40-15	34659	396839	4:45
p40-16	36044	9295	0:07
p40-17	38201	3040	0:04
p40-18	37562	9669	0:07
p40-19	38956	9857	0:07
p40-20	39658	13067	0:09

shows the number of branch-and-bound nodes and the CPU time on a standard PC. Large variations can be observed which are depending on the gap between the Lagrangian bound and optimum value. We will compare these results with a different bounding approach based on linear programming in the next chapter.

Chapter 5
Branch-and-Cut

Abstract This chapter focuses on the approach for solving the LOP to optimality which can currently be seen as the most successful one. It is a branch-and-bound algorithm, where the upper bounds are computed using linear programming relaxations.

5.1 Integer Programming

Linear programming is concerned with the problem of maximizing a linear objective function subject to finitely many linear constraints. Given a matrix $A \in \mathbf{R}^{(m,n)}$ and vectors $b \in \mathbf{R}^m$, $c \in \mathbf{R}^n$, the task is to find a vector $x^\star \in \mathbf{R}^n$ with

$$c^T x^\star = \max\{c^T x \mid Ax \leq b\}.$$

Such a problem is called a *linear programming problem* or *linear program* (*LP*). Note that a minimization problem can be transformed to a maximization problem by complementing the objective function. Equations can be expressed as pairs of inequalities, and bounds on variables can be included into $Ax \leq b$. The set of feasible solutions $\{x \mid Ax \leq b\}$ is a polyhedron and if the problem has a finite optimum solution then it has an optimum vertex solution. We do not go into more details on polyhedra here. They will be a central topic of Chap. 6.

Very effective algorithms for solving linear programming problems have been developed and very large instances can be treated in reasonable time. Important methods are the primal and dual simplex algorithms (as vertex following methods) and the barrier method (as an interior-point method).

For combinatorial optimization, linear programming models are not sufficient. Usually, some or all of the variables have to take integer values. If some variables have to be integral, we speak about a *mixed-integer linear programming problem*. If all variables are required to be integer, then we have an *integer linear program*; if all variables have to take values 0 or 1, the problem is called *linear 0/1*

programming problem (*0/1-IP*). This type of problem is an important one in combinatorial optimization. There has been significant progress in the development of algorithms for solving mixed-integer programming problems. But still, problems with several hundreds of variables and constraints can be difficult depending on their specific structure.

Usually, a linear combinatorial optimization problem $(E, \mathscr{I}, c)$ can easily be turned into a linear 0/1 programming problem. First, we associate with every set $F \subseteq E$ its *characteristic vector* χ^F by setting

$$\chi^F_e = \begin{cases} 1, & \text{if } e \in F, \\ 0, & \text{otherwise.} \end{cases}$$

Then equations and inequalities have to be found which are satisfied by all characteristic vectors corresponding to the feasible sets and are violated by all 0/1 vectors for sets $F \subseteq E$ with $F \notin \mathscr{I}$. In most cases, the combinatorial properties discriminating feasible solutions can be expressed in a straighforward way by linear constraints.

The LOP can be formulated as 0/1-IP as follows. We use 0/1 variables x_{ij}, for $(i, j) \in A_n$, stating whether arc (i, j) is present in the tournament or not. The basic observation is that in a tournament exactly one of the arcs (i, j) and (j, i) is present for every pair of nodes i and j, and that a tournament is acyclic if and only if it does not contain any dicycle of length 3.

The respective 0/1-IP is

$$\max \sum_{(i,j) \in A_n} c_{ij} x_{ij}$$
$$x_{ij} + x_{ji} = 1, \text{ for all } i, j \in V_n, i < j,$$
$$x_{ij} + x_{jk} + x_{ki} \leq 2, \text{ for all } i, j, k \in V_n, i < j, i < k, j \neq k,$$
$$x_{ij} \in \{0, 1\}, \text{ for all } i, j \in V_n.$$

This 0/1-IP can be considered as the canonical IP formulation of the problem. Note that the $\binom{n}{2}$ equations could actually be simply removed by substituting every variable x_{ij}, $j > i$, by $1 - x_{ji}$. More problematical is the large number $2\binom{n}{3}$ of 3-dicycle constraints.

Obviously, if we replace the constraints "$x_{ij} \in \{0,1\}$" by "$0 \leq x_{ij} \leq 1$" then we obtain a linear programming problem, the *canonical LP relaxation*. The optimum objective function value of this relaxation gives an upper bound on the solution value of the LOP.

If, by chance, the optimum solution is integral then the characteristic vector of an optimum acyclic tournament is found. If not, then together with heuristics providing good feasible solutions, we can immediately design a branch-and-bound algorithm for solving the LOP to optimality.

In this branch-and-bound approach problems are split into subproblems by fixing one selected variable either to 1 or to 0. During its execution the algorithm keeps a

list $\mathcal{K}$ of active problems and the value L of the best feasible solution $\bar{x}$ found so far.

LOPBranch-and-Bound

(0) Run a heuristic to provide a feasible ordering $\bar{x}$ with value L.

(1) Let problem P_0 be the canonical IP formulation. Set $\mathcal{K} = \{P_0\}$ and $k = 0$.

(2) If $\mathcal{K} = \emptyset$, Stop (The ordering $\bar{x}$ is optimal with value L).

(3) Select a problem $P_j \in \mathcal{K}$.

(4) Solve the linear programming relaxation LP_j of P_j. If LP_j is infeasible then set $c^* = -\infty$, otherwise let x^* be its optimum solution with value c^*. Distinguish the following three cases:

    (4.1) If $c^* \leq L$, then remove P_j from $\mathcal{K}$.

    (4.2) If $c^* > L$ and x_i^* is integral, then set $L = c^*$, $\bar{x} = x^*$ and remove P_j from $\mathcal{K}$.

    (4.3) If $c^* > L$ and x_i^* is not integral, then select an arc $(i, j) \in A_n$ with fractional value x_{ij}^*. Remove P_j from $\mathcal{K}$ and add the new problems P_{k+1} and P_{k+2} to $\mathcal{K}$, where
$$P_{k+1} = P_j \text{ with additional constraint } x_{ij} = 1,$$
$$P_{k+2} = P_j \text{ with additional constraint } x_{ij} = 0.$$
Set $k = k + 2$.

(5) Goto (2).

Clearly, this algorithm solves the LOP in finite time since in the worst case all possible feasible solutions would be enumerated. In principle, this approach actually splits a problem by putting bounds on a single variable and can also be used for general linear 0/1 or mixed integer programs. It was originally formulated for the latter case by Dakin [38].

The above enumeration scheme essentially fixes $i \prec j$ in one branch and $j \prec i$ in the other one. Of course, other schemes could also be used in this framework.

5.2 Cutting Plane Algorithms

If the LOP is formulated as 0/1-IP as above, instances of moderate size, say up to $n = 30$, can be solved with a comemrcial MIP solver in several hours. The solution time is highly problem dependent, in particular depending on the strength of the LP relaxation. For easy instances, the relaxation can sometimes even already provide an optimum solution. For difficult instances, however, the LP bound can be more than 4% off the optimum value and a large number of branch-and-bound nodes is generated. It is the main purpose of this and the following chapter to exhibit ways of overcoming these difficulties.

A first observation is that the number of dicycle constraints $2\binom{n}{3}$ is fairly large (already 323,400 for $n = 100$). It constitutes a major drawback if such large LPs have to be solved at every node of the branch-and-bound tree. On the other hand, only few of these constraints are binding at the optimum vertex and most of them are (depending on the objective function) actually not needed at all. The cutting plane approach makes use of this fact and tries to incorporate only important inequalities.

Suppose that the LP $\max\{c^T x \mid Ax \le b, 0 \le x \le 1\}$ with very many constraints has to be solved. The *cutting plane approach* solves such an LP as follows.

Cutting plane algorithm

(1) Initialize P as linear program $\max\{c^T x \mid 0 \le x \le 1\}$.
(2) Solve problem P and obtain an optimum solution x^*.
(3) If $Ax^* \le b$, $\mathtt{Stop}$ (x^* solves $\max\{c^T x \mid Ax \le b, 0 \le x \le 1\}$).
(4) Otherwise, choose an inequality $A_i.x \le b_i$ from $Ax \le b$ such that $A_i.x^* > b_i$. Augment P by this inequality and goto (2).

The term "cutting plane method" was chosen because in every iteration the current optimum point is cut off since it is infeasible for the real problem.

In step (4) also more than one violated inequality can be added. In practical computations this is highly advisable. It is also useful to eliminate constraints from the LP which are not binding at the current optimum. For the reoptimization in (2) after the addition of constraints, the dual simplex algorithm is most useful because it can be started immediately since dual feasibility is still given.

As an example we illustrate the algorithm when solving the LP relaxation of the $\mathtt{LOLIB}$ benchmark problem $\mathtt{be75np}$ (here $n = 75$). Table 5.1 shows for every reoptimization the number of inequalities added, removed resp., the current number of constraints in the LP and the current optimum value. The bound given by the relaxation is 790989 (the optimum linear ordering has value 790966). Instead of the possible 39200 only 1781 3-cycle inequalites were generated.

Table 5.1 Solution of $\mathtt{be75np}$ with cutting planes

#added ineq.	#removed ineq.	#current ineq.	Objective value
0	0	0	816602
200	0	200	810126
200	33	367	806134
200	33	534	802379
200	70	664	800896
200	55	809	795674
200	51	958	794413
200	83	1075	791815
200	48	1227	791196
126	71	1282	791036
27	35	1274	790990
28	17	1285	790989

Further issues such as how to select inequalities to be added, will be addressed later. Note the fact that, actually, in steps (3) and (4), it is not required to have an explicit list of the constraints but just to be able to answer the question of whether x^* is feasible and, if not, to provide a violated inequality. Furthermore, and this is what makes this approach even more powerful: in principle, we do not even have to know the system of inequalities explicitly. We want to make this more precise.

Let $\max\{c^T x \mid Ax \leq b, x \in \{0,1\}^n\}$ be a formulation of a combinatorial optimization problem as 0/1-IP. From polyhedral theory, it is known that there exists a finite system $Bx \leq d$ such that the vertices of the polyhedron $\{x \mid Bx \leq d\}$ are exactly the feasible 0/1 solutions of $Ax \leq b$. Therefore, if the system is known, then the LOP can be solved as the linear programming problem $\max\{c^T x \mid Bx \leq d\}$.

An inequality $f^T x \leq f_0$ is said to be *valid* with respect to $\{x \mid Ax \leq b, x \in \{0,1\}^n\}$ if $f\overline{x} \leq f_0$ for all $\overline{x} \in \{x \mid Ax \leq b, x \in \{0,1\}^n\}$.

If, in the cutting plane algorithm, we can find violated valid inequalities w.r.t. the set $\{x \mid Ax \leq b, x \in \{0,1\}^n\}$ as long as the current point is not feasible then we could in principle not only solve the relaxation but also the combinatorial optimization problem if we have access to all necessary inequalities from the system $Bx \leq d$. This algorithm would be correct under the assumptions that all occuring LPs can be solved, that a cutting plane can always be generated, and that it terminates after a finite number of iterations.

Therefore, at the core of applying linear programming to combinatorial optimization we have the following problem.

Separation Problem
Given a description $\{x \mid Ax \leq b, x \in \{0,1\}^n\}$ of a combinatorial optimization problem and a vector y^*, either prove that y^* is feasible or find an inequality which is valid for $\{x \mid Ax \leq b, x \in \{0,1\}^n\}$, but violated by y^*.

An algorithm which solves the separation problem is called an *exact separation algorithm*.

Usually, we have the following situation when attempting to design a cutting plane algorithm for a combinatorial optimization problem.

- A linear characterization $\{x \mid Ax \leq b, x \in \{0,1\}^n\}$ is available and the feasibility test just amounts to checking whether x^* is binary.
- The system $Bx \leq d$ is not known, but some partial information is available and there are types (called *classes*) of inequalities which have similar structure or logic behind them. Classes may (and will) contain many inequalities (usually at least expontially many in n).
- For some of these classes it can be checked in polynomial time whether all inequalities of the class are satisfied and, if not, a cutting plane can be provided. These are classes of inequalities with *exact separation algorithms*.
- As can be expected for NP-hard problems there will be classes of inequalities the separation of which is NP-complete or where no good separation algorithm is known. In such a case we have to be content with *separation heuristics*. But,

also if exact separation is possible, heuristics may be employed in addition to save computing time.

It might seem that cutting planes offer a chance to escape the inherent complexity of NP-hard combinatorial optimization problems. In this context the following result (see e.g. [61]) is of major importance.

Theorem 5.1. *Let $P = \{x \mid Ax \leq b\}$ be a rational polyhedron such that the encoding length of every inequality in $Ax \leq b$ is at most φ.*
Then for every vector $c \in \mathbf{Q}^n$ the optimization problem $\max\{c^T x \mid Ax \leq b\}$ can be solved in polynomial time (in φ and the encoding length of c) if and only if for every vector $y \in \mathbf{Q}^n$ the separation problem with respect to P can be solved in polynomial time (in φ and the encoding length of y).

This theorem states the equivalence of optimization and separation. Therefore, LP relaxations with polynomial separation algorithms can be solved in polynomial time. The number of inequalities is not relevant as long as cutting planes can be generated in polynomial time. It should be noted that this result is only true if the respective linear programs are solved using the ellipsoid method. Since this method is not practicable, the simplex algorithm is used in practice. No polynomial version of the simplex method could be formulated so far but, in any case, it is a fast algorithm for practical computations.

We now have a theoretical framework for solving combinatorial optimization problems using linear programming and branch-and-bound. The remaining sections of this chapter will be devoted to practical issues of the implementation and possibilities to add further cuts which are not contained in the LP relaxation and not known beforehand.

5.3 Branch-and-Cut with 3-Dicycle Cuts

A branch-and-bound algorithm using the canonical LP relaxation for computing upper bounds can, in principle, be used for solving the LOP to optimality. An interesting theoretical result should be mentioned here. Recall the Lagrangian relaxation with 3-dicycle inequalities of Chap. 4:

$$L(\mu) = \max \sum_{(i,j)\in A_n} c_{ij}x_{ij} + \sum_{(r,s,t)\in T} (2 - x_{rs} - x_{st} - x_{tr})\mu_{rst}$$
$$x_{ij} + x_{ji} = 1, \text{ for } 1 \leq i < j \leq n,$$
$$x_{ij} \in \{0,1\}, \text{ for } 1 \leq i, j \leq n, i \neq j.$$

Because the constraints of the Lagrangian problem define a polyhedron with 0/1 vertices only, it can be shown that the optimum value $\min_{\mu \geq 0} L(\mu)$ is equal to the optimum value of the canonical LP relaxation of the LOP.

We now address the question of how to solve the LP relaxation most effectively and we will see that its exact value can be obtained much faster than the approximate value using the Lagrangian approach.

5.3.1 Solving the 3-Dicycle Relaxation

The 3-dicycle relaxation to be solved at every node of the tree is

$$\max \sum_{(i,j) \in A_n} c_{ij} x_{ij}$$

$$x_{ij} + x_{jk} + x_{ki} \leq 2, \text{ for all distinct nodes } i, j, k \in V_n,$$

$$x_{ij} + x_{ji} = 1, \text{ for } 1 \leq i < j \leq n,$$

$$x_{ij} \geq 0, \text{ for } 1 \leq i, j \leq n, i \neq j.$$

In the branching process some of the variables will be fixed.

When looking for violated 3-dicycle inequalities it makes no sense to look for some sophisticated procedure: it is appropriate to just enumerate all of them and check for violation. But there are some possibilities to speed up the solution process.

An important observation is that it is helpful to try to generate "deeper" cuts. To this end we set z as a convex combination of the current LP optimum x^* and the currently best known feasible solution $\bar{x}$ and first try to cut off z. Only if z cannot be cut off, then 3-dicycle inequalities separating x^* are searched for. Experiments showed that the setting $z = \frac{1}{3}x^* + \frac{2}{3}\bar{x}$ leads to good results and usually reduces the overall CPU time by about 60%.

Furthermore, it is not necessary to add all violated inequalities in every phase, but to limit their number. Significant progress in the first phases is already achieved with few inequalities and thus the LP size only increases slowly. The best limit on the number of added cuts depends on the problem size, but usually limiting the number to about several hundreds is perfect. It is a further improvement to add in the first phases only inequalities which are arc-disjoint, i.e., having the property that no variable occurs in more than one cut.

It is not only important to limit the number of cuts added, but also to select among the available cuts. This is substantiated in Table 5.2. Columns 5 and 6 show the number of LPs solved (nlps_2) and the CPU time (CPU_2) in min:sec if cuts were just selected at random. Alternatively, cuts $f^T x \leq f_0$ with large value $\frac{c^T f}{\|f\|}$ were preferred. The idea is that inqualities with small angle between their normal vector and the normal vector of the objective function should yield progress. The results for this strategy are shown in columns 3 and 4 (nlps_1 and CPU_1). Depending on the problem it can lead to substantial savings in CPU time (see in particular `atp76`).

Table 5.2 Importance of cut selection strategies

Problem	Size	nlps$_1$	CPU$_1$	nlps$_2$	CPU$_2$
EX1	50	15838	7:15	18065	8:38
EX2	50	20993	10:26	24977	13:07
EX4	50	244	0:06	365	0:14
EX5	50	346	0:10	445	0:15
EX6	50	138	0:04	203	0:08
atp66	66	112	0:06	278	0:29
atp76	76	591	0:49	4186	14:28
econ77	77	30	0:02	63	0:04
randD	50	9711	4:14	10097	4:20
Sum		48003	23:16	58679	41:46

5.3.2 An LP Based Heuristic

The following hypothesis proved to be very useful. We assume that the current fractional LP solution x^* should be somehow close to an optimum solution and contain some of its characteristics. Therefore, we exploit x^* for starting a heuristic to find a feasible solution. For every object i we compute

$$s_i = \sum_{j \neq i} x_{ij}^*$$

and then sort the objects according to nondecreasing values s_i. The corresponding linear ordering serves as a starting solution for improvement heuristics.

For proving the usefulness of this idea, we solved the 3-dicycle relaxation for the LOLIB input-output matrices and called the lower bound heuristic for every LP solution. As improvement heuristic we just employed local enumeration as described in Chap. 2.

Since almost all problems could be solved to optimality we list in Table 5.3 only those problems that could not be solved and the artifical problems econ36 – econ77. The table shows in column 3 the deviation of the 3-dicycle upper bound from the optimum solution in percent. For these easy problems the bound is very good and at most 0.052% above the optimum. (The bound for stabu70 is less than 0.0005% off).

In addition, Table 5.3 proves that the LP based heuristic works very well achieving at least 99.26% of the optimum.

5.3.3 Computational Results with 3-Dicycles

In another experiment we solved the same instances as with the branch-and-bound algorithm of Charon and Hudry of Chap. 4. Table 5.4 displays for every problem

Table 5.3 3-Dicycle relaxation

Problem	OPT	3CYC	LPHeu
be75np	790966	0.003%	100.000%
stabu70	422088	0.000%	99.938%
t59b11xx	245750	0.005%	100.000%
econ36	555568	0.003%	99.888%
econ43	675180	0.001%	99.670%
econ47	845374	0.032%	99.889%
econ58	1263005	0.004%	99.358%
econ59	1256708	0.035%	99.726%
econ61	1265164	0.002%	99.881%
econ62	1275989	0.005%	99.263%
econ64	1328547	0.011%	99.627%
econ67	1437471	0.007%	99.969%
econ68	1480971	0.052%	99.770%
econ71	1636218	0.007%	99.624%
econ72	1932752	0.028%	99.795%
econ73	2146505	0.009%	99.648%
econ76	2781838	0.008%	99.735%
econ77	2798507	0.007%	99.994%

the root bound obtained from the 3-dicycle relaxation, the number of nodes in the branch-and-cut tree, the maximum depth of the tree, and the CPU time (in min:sec).

Table 5.4 Branch-and-cut with 3-dicycle inequalities

Problem	Optimum	Root bound	#nodes	#maxlevel	CPU
p40-01	29457	29494.47	3	2	0:00
p40-02	27482	28032.00	5449	28	15:24
p40-03	28061	28354.33	115	12	0:20
p40-04	28740	29298.67	2317	22	7:12
p40-05	27450	28213.33	976	27	9:26
p40-06	29164	29632.33	613	15	1:54
p40-07	28379	29006.00	2777	21	8:30
p40-08	28267	28870.67	1279	16	4:08
p40-09	30578	31183.00	5675	23	17:00
p40-10	31737	32147.67	419	15	1:07
p40-11	30658	31275.00	3985	23	11:43
p40-12	30986	31479.00	571	14	1:48
p40-13	33903	34056.85	10	5	0:00
p40-14	34078	34494.33	241	13	0:39
p40-15	34659	35369.67	5073	28	14:24
p40-16	36044	36199.00	17	5	0:00
p40-17	38201	38217.86	3	2	0:00
p40-18	37562	37694.93	17	5	0:00
p40-19	38956	39117.17	15	5	0:00
p40-20	39658	39812.67	31	7	0:00

The problems are solved in about the same order of magnitude of CPU time. The branch-and-cut algorithm generates much fewer tree nodes, but solving the LPs needs a lot of time. Of course, since the 3-dicycle bound and exact Lagarangian bound are equal, problems which are difficult for one code are difficult for the other one as well. Recent experiments with a bundle method, however, were very promising. Running some iterations of the bundle method on the Lagrangian relaxation seems to offer the chance of identifying important starting inequalities for computing the 3-dicycle bound with a subsequent cutting plane algorithm.

There have also been experiments using interior-point solvers for optimization the LP relaxations. Restarting after adding cuts is difficult here. In [99, 100] restart is realized by backtracking some steps on the path taken when optimizing the previous LP. Results, however, do not suggest employing interior-points methods in this context.

5.4 Generation of Further Cuts

Only 3-dicycle cuts are not sufficient for attempting to solve larger problems. We will now address the question of how to generate further cuts. As further cuts we will only discuss at this point cuts that can basically be used for every combinatorial optimization problem. Special further cuts for the LOP will be described in Chap. 6.

5.4.1 Chvátal-Gomory Cuts

We discuss briefly the principle of general cuts. Their generation has far-reaching consequences at least in theory, but to some extent also in practice.

Chvátal-Gomory cuts can be viewed as being obtained by so-called closure operations which allow for generating new stronger inequalities from an inequality system.

Definition 5.1. Let S be the system $Ax \leq b$ of rational inequalities.

(i) An inequality $d^T x \leq d_0$ with integral d is said to belong to the *elementary closure* of S if there is a rational vector $\lambda \geq 0$ such that $\lambda^T A = d^T$ and $\lfloor \lambda^T b \rfloor \leq d_0$.
(ii) The set of all inequalities belonging to the elementary closure of S is denoted by $e^1(S)$. For $k > 1$ the set $e^k(S)$ is defined as $e^k(S) = e^1(S \cup e^{k-1}(S))$.
(iii) The *closure cl(S) of S* is given as

$$\text{cl}(S) = \cup_{k=1}^{\infty} e^k(S).$$

An inequality is said to have *Chvátal rank* k with respect to S if it is contained in $e^k(S)$, but not in $e^{k-1}(S)$. The number k is a good indicator for the complexity of an inequality.

As an example, assume that for integral variables y and x_i, $i = 1, \ldots, k$ the following constraints have to be satisfied:

$$a_1 x_1 + \ldots + a_k x_k - y = b \tag{5.1}$$

$$-x_i \leq 0, \quad i = 1, \ldots, k. \tag{5.2}$$

By adding suitable multiples of (5.2) to (5.1) and rounding down the right hand side one obtains the valid inequality

$$\lfloor a_1 \rfloor x_1 + \ldots + \lfloor a_k \rfloor x_k - y \leq \lfloor b \rfloor. \tag{5.3}$$

Subtracting (5.1) from (5.3) now yields the inequality

$$(\lfloor a_1 \rfloor - a_1) x_1 + \ldots + (\lfloor a_k \rfloor - a_k) x_k \leq \lfloor b \rfloor - b.$$

Note that this inequality is violated if $x_1 = \ldots = x_k = 0$. This is the standard type of a Chvátal-Gomory cut as it can be obtained directly from the simplex tableau. There the variables $x_1, \ldots, x_n$ are the non-basic variables at the optimum of the current LP relaxation and thus the inequality cuts off this optimum, but is satisfied by all feasible integral points.

A very nice theoretical results is

Theorem 5.2. *A cutting plane algorithm with Chvátal-Gomory cuts either solves an integer programming problem* $\max\{c^T x \mid Ax = b, x \geq 0, x \text{ integer}\}$ *in finitely many steps or verifies that it is unbounded or infeasible.*

For proving this theorem several technicalities have to be observed. The optimal tableau of the LP relaxations has to be lexicographically positive and thus the pivoting rule of the dual simplex has to be modified to preserve lexicographic positivity. A lower bound on a feasible solution (derivable from the data) has to be tested. A cut has to be derived from the row with smallest index. If in the dual simplex the slack variable of a cut becomes basic, then the cut is eliminated from the problem.

It turned out that this approach cannot be utilized in practice straightaway because after the addition of many such cutting planes severe numerical problems occur. On the other hand, however, it was shown that the careful use of Chvátal-Gomory cutting planes can lead to substantial improvements in integer programming.

5.4.2 Maximally Violated Mod-k Cuts

As a further possiblity to generate cuts of general nature in a branch-and-cut algorithm we describe mod-k inequalities which are special kinds of Chvátal-Gomory cuts.

Let $Ax \leq b$ be a system of linear inequalities with integral coefficients and let $k > 1$ be an integer number. Suppose that we scale every inequality r of $Ax \leq b$

by a nonnegative factor μ_r and sum the resulting inequalities. Now, if μ denotes the vector of all factors, assume that all coefficients of $\mu^T A$ are divisible by k and that the remainder on dividing $\mu^T b$ by k is $k-1$. Hence μ satisfies the congruence system

$$\mu^T A \equiv 0^T \quad \mathrm{mod}\, k$$
$$\mu^T b \equiv k-1 \quad \mathrm{mod}\, k.$$

We have $\mu^T b = sk + (k-1)$ for some $s \in \mathbf{Z}$, and therefore $\mu^T b - (k-1)$ is divisible by k. Furthermore $\mu^T A x$ is divisible by k for all $x \in \{x \in \mathbb{Z}^n \mid Ax \leq b\}$, and therefore the inequality

$$\mu^T A x \leq \mu^T b - (k-1)$$

is valid for all feasible integer solutions of $Ax \leq b$. We can express the inequality in an equivalent way as the *mod-k inequality*

$$\frac{1}{k}\mu^T A x \leq \frac{1}{k}\left(\mu^T b - (k-1)\right).$$

Now consider some fractional solution x^* of $Ax \leq b$. In a branch-and-cut algorithm we would like to find an integer k and a vector μ such that the congruence system above is satisfied and the corresponding mod-k inequality is violated by x^* and thus provides a cutting plane (a so-called *mod-k cut*). Since $\mu^T A x^* \leq \mu^T b$ this solution can violate $\mu^T A x \leq \mu^T b - (k-1)$ by at most $k-1$ and the maximal violation can only be achieved if $\mu^T A x^* = \mu^T b$, i.e., if $\mu_r = 0$ for all r with $A_r x^* < b_r$.

We use the separation algorithm for maximally violated mod-k cuts suggested in [25]. Let x^* be a fractional solution of the current LP relaxation $Ax \leq b$. Due to the remarks above, in order to find a maximally violated inequality, we restrict the congruence system to contain only those inequalities that are tight for x^*. We have to choose k and find an integer multiplier vector $\mu \geq 0$ such that the congruence system is satisfied. Note that the coefficients of μ can obviously be restricted to values smaller than k, i.e., $\mu_r \in \{0,1,\ldots,k-1\}$ for all r. It was proved in [25], that only prime numbers have to be considered for k.

Usually, if one cut exists then there are plenty of others. This is due to the fact, that there is a number f of free variables μ_i in the solution whose values can be chosen arbitrarily from $\{0,1,\ldots,k-1\}$. There can be dozens or even a few hundred of these variables and thus there exist k^f different solutions of the congruence system. Even though the map of solutions to cuts is not injective, we will have to address the problem of selecting cuts from the huge set of generated constraints to be added to the linear relaxation.

5.4.3 Mod-2 Cuts

For the special case $k = 2$ an efficient separation procedure is given in [24]. Again, we search for a multiplier vector μ for $Ax \le b$ that satisfies the congruence system. The arguments of previous section show that

$$\frac{1}{2}\mu^T Ax \le \frac{1}{2}(\mu^T b - 1).\qquad(5.4)$$

is valid for $\{x \in \mathbb{Z}^n \mid Ax \le b\}$ and can be violated by at most $\frac{1}{2}$ by a tight fractional solution x^* of $Ax \le b$. For $\lambda = \frac{1}{2}\mu$ and μ satisfying the congruence system, all entries of $\lambda^T A$ are integer and $\lambda^T b - \frac{1}{2} = \lfloor \lambda^T b \rfloor$. Therefore (5.4) can be expressed equivalently as

$$\lambda^T Ax \le \lfloor \lambda^T b \rfloor.$$

This inequality is a so-called $\{0, \frac{1}{2}\}$-*Chvátal-Gomory inequality*. The identification of a suitable vector λ in principle amounts to optimizing over

$$P_{1/2} = \left\{ x \in \mathbb{R}^n \;\middle|\; Ax \le b, \lambda^T Ax \le \lfloor \lambda^T b \rfloor \text{ and } \lambda \in \{0, \tfrac{1}{2}\}^m \text{ with } \lambda^T A \in \mathbb{Z}^n \right\}.$$

As the separation problem of $P_{1/2}$ is NP-hard, it is unlikely that one can optimize over $P_{1/2}$ in polynomial time. But, in [24], possibilities are derived for relaxing $P_{1/2}$ in such a way that separation becomes polynomially solvable.

By a so-called *weaking of inequalities*, a separation algorithm is developed which basically amounts to solving shortest path problems in a specially constructed digraph.

The advantage of this *shortest path mod-2 method* is that there is no restriction on the constraints that are used for the generation of the cut, as long as they can be transformed in such a way that they have exactly two odd coefficients on the left hand side. Another advantage is that the digraph for the shortest path computations is usually sparse. On the other hand, no extension from 2 to bigger prime numbers is possible and the violation of the resulting cut is maximal if and only if all used constraints are tight for x^*.

There is a further problem-independent approach for generating cuts, namely so-called *local cuts* and *target cuts*. Because it is easier to explain these cuts in the context of the linear ordering polytope, their discussion is postponed to Sect. 6.7.

5.5 Implementation of Branch-and-Cut

We will now discuss in more detail implementational issues for the realization of a branch-and-bound algorithm where upper bounds are computed using LP relaxations. Because the LP relaxations are solved with the cutting plane approach the notion "branch-and-cut" was coined for this type of algorithms.

We will keep the discussion on a general level valid for all combinatorial optimization problems. At some point we will add special remarks for the LOP.

In the following it is assumed that the problem is defined on a graph or digraph and that the variables are associated with edges or arcs. Variables and arcs/edges are used as synonyms.

A 0/1-IP formulation of the problem is assumed to be available as

$$\max\{c^T x \mid Ax \leq b, x \in \{0,1\}^n\}.$$

The minimal equation system is also known, but for simplicity not listed explicitly. Further inequalities $Ax \leq b$ in addition to the system $\tilde{A}x \leq \tilde{b}$ are known with

$$\{x \mid Ax \leq b, x \in \{0,1\}^n\} \subset \{x \mid \tilde{A}x \leq \tilde{b}\} \subset \{x \mid Ax \leq b, 0 \leq x \leq 1\}.$$

Note that an explicit list of the constraints of $\tilde{A}x \leq \tilde{b}$ is not needed. We assume that for some classes of inequalities in $\tilde{A}x \leq \tilde{b}$ exact or heuristic separation algorithms are available.

The algorithm will construct a branch-and-cut tree whose nodes represent the subproblems generated. The following states of a tree node will be distinguished:

– *current*: this node is presently worked on,
– *active*: the node is generated, but not considered yet,
– *inactive*: the node has been treated and has active successors,
– *fathomed*: this node and all of its successors have been worked on (such nodes can obviously be deleted from the tree).

A variable can be in one of the following states:

– *active*: the variable is present in the current LP,
– *inactive*: the variable is not present in the current LP,
– *fixed*: the variable is permanently fixed to 1 or 0,
– *set*: the variable is set to 1 or 0, but this is valid only in part of the tree.

Furthermore, let L be the global lower bound on the optimum objective function value (best known feasible solution so far) and let U be the global upper bound (valid for the original problem).

We now describe the main aspects of the branch-and-cut algorithm. The goal is to solve $\max\{c^T x \mid Ax \leq b, x \in \{0,1\}^n\}$ based on the approximation $\tilde{A}x \leq \tilde{b}$.

5.5.1 Initialization

As first LP for the root node $\max\{c^T x \mid 0 \leq x \leq 1\}$ is usually chosen. In principle, the minimal equation system could be used to eliminate variables beforehand. However, this is only reasonable in few cases, such as the LOP. Here variables x_{ij}, $i > j$, can simply be replaced by $1 - x_{ji}$ without effecting the density of the constraint matrix.

One could already introduce some of the inequalities from $Ax \leq b$ if respective information is available. Computational experience has shown that this only has a significant effect if inequalities are selected carefully.

5.5.2 Active Variables

If the number of variables is large, it can be useful to work only with a subset of active variables and assume that the inactive variables are zero. Of course, it has to be checked later (by pricing) if this assumption is correct.

For the LOP, working with only a subset of variables has not been interesting so far. The problem is already difficult for small values of n, and the number $n(n-1)$ of its variables ($\binom{n}{2}$ after elimination, resp.) is still small.

5.5.3 Local Upper Bound

The solution of the current LP gives an upper bound for the respective subproblem and all subproblems generated from it (if all inactive variables price out correctly). LPs are usually solved using the dual simplex method because of the simple restart after addition of constraints or variables.

5.5.4 Branching

If the current subproblem cannot be solved, it has to be partitioned. A common way for partitioning (also frequently used for the LOP) is to choose some variable x_{ij} with fractional value and set it to 1 in one subproblem and to 0 in the other. This way a binary branch-and-cut tree is generated.

Several priority rules for selecting this branching variable are possible. It is common to choose variables close to $\frac{1}{2}$ with high absolute value of their objective function coefficient. The motivation behind this is that setting the branching variable should have some effect on the current problem. Non-binary partitions or generating subproblems using inequalities is also possible.

5.5.5 Fixing and Setting of Variables

Let x^* be the optimum solution of the current LP with value c^* and let r be the vector of reduced costs.
The following holds for a nonbasic variable x_e:

(i) If $x_e = 0$ and $c^* + r_e \leq L$, then one can set $x_e = 0$.

(ii) If $x_e = 1$ and $c^* - r_e \leq L$, then x_e can be set to 1.

This setting of variables is valid for the current node and all of its successors.

If setting is possible at the root node, then it amounts to fixing a variable permanently. If advantageous, the variable could be removed from the problem.

5.5.6 Logical Implications

Setting variables can (depending on the problem) influence other variables as well. E.g. if in the LOP we have the setting $x_{ij} = 1$ and $x_{jk} = 1$, then $x_{ki} = 0$ is implied, because otherwise a 3-dicycle inequality would be violated.

Tests for logical implications can in particular be helpful since it can affect basic variables which are currently fractional. Nonbasic variables are only set or fixed to values they already have (so, locally there is no effect on the problem).

5.5.7 Selection of Nodes

When work on the current node is finished, then the next node for processing has to be chosen. Basic strategies are:

– *Depth-first*
  Choose an active node with highest level in the branch-and-cut tree.
– *Breadth-first*
  Choose an active node with lowest level in the branch-and-cut tree.
– *Best-first*
  Choose a node with smallest difference between its lower and upper bound. (Note that different upper bounds are valid for the nodes.)
– *Strong branching*
  Here work is started on several active nodes and depending on the progress one of them is selected promising the best improvement of the upper bound.

Depth-first is mainly employed if it is at all difficult to find feasible solutions. Otherwise it is inferior compared with the other strategies. Best-first leads to fewer nodes and better running times than breadth-first. Strong branching usually performs best, but at higher effort.

Further heuristics rules for selecting the most promising node can be formulated in addition.

5.5.8 Lower Bounds

Good feasible solutions are very important in order to obtain lower and upper bounds which are close together. Feasible solutions can be computed independent from the branch-and-cut algorithm beforehand or in parallel on a separate processor. As pointed out above, we found LP based heuristics exploiting the current fractional LP solution very useful.

Note that, if better lower bounds are found, fixing further variables can be tried. To this end the reduced costs at the root node have to be stored.

5.5.9 Separation

Separation clearly is at the core of branch-and-cut. It depends on the knowledge about the system $\tilde{A}x \leq \tilde{b}$ and on how many effective separation algorithms are available. Strategies for calling the respective procedures and for the selection of cutting planes found have to be developed.

5.5.10 Elimination of Constraints

With the cutting plane approach only a very small fraction of potential inequalities is actually introduced into the LP, but nevertheless this number also can be huge. It has proved to be reasonable to eliminate inequalities which are not binding at the current LP optimum. However, these inequalities should be stored in a constraint pool, because they could be valuable for other nodes. This is in particular advisable if they have been found with high computational effort.

5.5.11 Constraint Pool

The constraint pool stores inequalities that are needed for the initialization of the next node or were temporarily eliminated. Since the pool can grow very large, suitable data structures for storing inequalities in sparse format are necessary.

5.5.12 Pricing

If the current LP is solved in the active variables then it has to be checked if the setting of the inactive variables is optimal as well. This can be accomplished by evaluating their reduced costs.

Assume that x_e is an inactive variable with value 0. First the column for x_e in the constraint matrix has to be retrieved; let this be a_e. Then the equation system

$$A_B \bar{a}_e = a_e$$

has to be solved. The reduced costs of x_e are

$$r_e = c_e - c_B^T \bar{a}_e = c_e - c_B^T A_B^{-1} a_e.$$

Usually, in LP codes, the dual vector $c_B^T A_B^{-1}$ is available and does not have to be computed. So there is no need for solving an equation system.

If $r_e \geq 0$, for all inactive variables x_e, then the current LP solution is optimal for the complete problem, although possibly only a small subset of the variables has actually been used. If the reduced costs are negative for one or more variables, then some of them have to be activated and the LP has to be augmented correspondingly.

It is possible to perform pricing in several stages by assigning priority classes to inactive variables. Also, not all variables will be examined, if already several have been found that do not price out correctly. Note that pricing incurs significant computational effort.

5.5.13 Infeasible LPs

If there are inactive variables it is possible that the current LP is infeasible. This can be caused by a setting of variables that makes it impossible to achieve feasibility with active variables only. In such a case variables have to be added to regain feasibility. It is not clear how to choose such variables, and there are only heuristic strategies available.

5.5.14 Addition of Variables

Addition of variables is necessary if LPs become infeasible or it can also be caused by pricing. It can also be helpful to activate variables which seem to be important because they are contained in good feasible solutions.

The efficient implementation of the above components is not trivial. In particular, the administration of branch-and-cut nodes, of active variables and of the constraint pool require some effort. But these are tasks which can be solved to a large extent independently of the concrete problem and there are frameworks like ABACUS [74], SCIP [2] or SYMPHONY [109] that facilitate the development of branch-and-cut algorithms. Cutting plane generation is problem dependent in any case and also requires complex algorithms and data structure.

5.6 Some Computational Results

We report about computational results from [107] where also further details on the implementation of mod-2 and mod-k separation can be found.

We have applied the branch-and-cut algorithm with mod-2 and mod-k cuts on ten random problem instances `p40-01 – p40-10`. The problem instances are difficult and no 3-dicycle relaxation had an integral optimum solution. In the following we describe how we deal with the fractional solution x^* of the relaxation. We will also speak about the digraph associated with x^* which consists of the arcs whose associated variables have a positive value.

Since the generation of cuts is time consuming, we incorporated a heuristic element in our separation procedures to possibly generate further cuts at cheap cost. Our idea is based on an interesting property of the linear ordering polytope which will be discussed in the next chapter. Namely, for this polytope a rotation mapping can be defined which converts valid inequalities into different valid inequalites. We enhanced our separation routines by also checking rotated versions of violated inequalities for violation. This lead to the detection of further cuts.

We are using two general strategies for trying to generate mod-k cuts violated by the current fractional solution x^*. The first strategy considers small subdigraphs (of the digraph defined by x^*) and generates all violated cuts that can be found for this digraph. The idea is that by proceeding this way, the relaxation can locally be strengthed considerably and therefore allow for a reasonable bound improvement in the branch-and-cut algorithm. The second strategy applies the mod-k separation routine to the complete digraph and selects cuts afterwards. Cuts for the whole digraph should provide global information which is also important for the algorithm.

The strategies for choosing an appropriate subdigraph differ as follows. The *variable heuristic* limits the number n_V of variables from which the subdigraph is constructed, while the *improvement heuristic* limits the number n_N of the nodes of the subdigraph.

The variable heuristic starts with the variable that occurs in the most constraints of the current LP. (If there are several variables satisfying this condition, then we choose one at random.) Then we continue to successively select all other variables from these constraints and also choose all constraints that contain these new variables. If the limit n_V on the number of variables is reached, the selection is stopped and the subdigraph G' is defined by the variables selected by this procedure.

A practical problem we had to address was that already a slight increase of n_V could lead from subdigraphs for which no cuts were found to subdigraphs where very many violated cuts existed. This phenomenon turned out to be caused by the fact that not all 3-dicycle inequalities for the subdigraph are part of the LP relaxation. After all trivial and binding 3-dicycle inequalities were added, this problem disappeared. We experienced that, for our problem instances, $n_V = \frac{1}{4}n$ and $n_V = \frac{1}{5}n$ are good bounds as the resulting digraph was just big enough to generate a reasonable number of mod-k cuts.

The improvement heuristic uses information from the current LP solution. For every node, the sum of the values of variables corresponding to outgoing arcs is

computed. Then the nodes are linearly ordered with respect to nonincreasing sums. Our heuristic is based on the hypothesis that nodes close together in this ordering will also be neighboring in an optimum solution. For forming the subdigraph we therefore take the first n_N nodes of the ordering, then the second n_N nodes, etc., and generate all mod-k cuts from these n/n_N subgraphs. Values $n_N \in \{\frac{1}{4}n, \frac{1}{5}n, \frac{1}{6}n\}$ lead to suitable results.

The third strategy is to apply the maximally violated mod-k method to the whole digraph and select cuts afterwards. This selection strategy is based on some different criteria. First, we randomly order the columns of the matrix to avoid always generating the first cuts obtained from the system in each iteration. Second, we prefer cuts introducing few non-zero coefficients because dense LPs are usually more difficult than sparse ones. A further criterion is to select only the single basic solution or at most f trivial solutions and not take all k^f possible solutions of the system into account.

For the shortest path calculation of the mod-2 procedure we use Dijkstra's algorithm leading to running time $O(\overline{E} \log \overline{V})$. To avoid inequalities being found more than once we delete all nodes that are part of an already found violated inequality from the list of potential starting nodes.

In the case of rotation we repeat this procedure for all rotation parameters r in a random order until the limit of 250 cuts is reached. For the violation tolerance we use the relatively high value of 0.01, because otherwise the separation procedure finds a lot of cuts without significantly improving the dual bound. In addition we use tailing off which stops the mod-2 separation if the last 10 separations did not improve the bound by more than 0.05 percent.

Table 5.5 Root bounds

Problem	Opt	3-cyc	MV	M2	MV+M2
p40-01	29457	29494.48	29488.09	29457.00	29457.00
p40-02	27482	28032.00	27967.75	27768.25	27768.11
p40-03	28061	28354.33	28297.53	28098.77	28098.59
p40-04	28740	29298.67	29220.46	28976.85	28974.93
p40-05	27450	28213.33	28091.60	27855.72	27849.33
p40-06	29164	29632.30	29553.90	29321.40	29316.30
p40-07	28379	29006.00	28892.08	28651.53	28647.43
p40-08	28267	28870.67	28744.77	28463.46	28460.07
p40-09	30578	31183.00	31078.85	30898.49	30891.87
p40-10	31737	32147.67	32089.26	31897.68	31894.69

Table 5.5 displays, for the instances p40-01 – p40-10, the optimum values, the 3-dicycle bound and the improved bound when additional cuts are added for tightening the LP relaxation. In the first two experiments we added mod-2 and mod-k cuts separately (MV and M2), in the third experiment both separations were employed (MV+M2). The branch-and-cut algorithm was stopped when no more violated inequalities could be found at the root node, i.e., no branching was started.

Table 5.6 verifies that these bounds can be improved considerably with little additional computational effort when rotation is employed to find further cuts.

Improvement with respect to the 3-dicycle relaxation is easier to assess if the *gap closure*

$$100 \times \frac{|c^T x^* - c_{3\text{cycle}}|}{|c_{\text{opt}} - c_{3\text{cycle}}|}$$

is computed, where c_{opt}, $c_{3\text{cycle}}$, and $c^T x^*$ are the optimum objective function value, the 3-dicycle upper bound, and the bound obtained with additional mod-k cuts, respectively. The gap closure measures (in percent) how much of the gap between 3-dicycle bound and optimum value could be closed.

Table 5.6 Root bounds with rotation

Problem	Opt	3-cyc	MV_R	$M2_R$	$MV+M2_R$
p40-01	29457	29494.48	29485.69	29457.00	29457.00
p40-02	27482	28032.00	27956.53	27648.32	27647.84
p40-03	28061	28354.33	28290.45	28061.00	28061.00
p40-04	28740	29298.67	29193.68	28829.76	28829.77
p40-05	27450	28213.33	28085.01	27679.71	27678.77
p40-06	29164	29632.30	29546.30	29186.20	29186.10
p40-07	28379	29006.00	28879.06	28481.02	28480.75
p40-08	28267	28870.67	28735.02	28311.82	28312.41
p40-09	30578	31183.00	31077.59	30697.73	30698.30
p40-10	31737	32147.67	32090.65	31778.55	31778.37

Table 5.7 shows that the gap between 3-dicycle bound and optimum can be closed by 86% on average and that the gap closure is mainly due to mod-2 cuts with rotation.

So with respect to bound improvement the additional separation has proved its advantages. But it is interesting as well to check if this improvement also leads to faster computation times when a provably optimum solution has to be computed.

Table 5.7 Gap Closure

Name	MV	MV_R	M2	$M2_R$	MV+M2	$MV+M2_R$
p40-01	17.05%	23.45%	100.00%	100.00%	100.00%	100.00%
p40-02	11.68%	13.72%	47.96%	69.76%	47.98%	69.85%
p40-03	19.36%	21.18%	87.12%	100.00%	87.19%	100.00%
p40-04	14.00%	18.79%	57.61%	83.95%	57.95%	83.94%
p40-05	15.95%	16.81%	46.85%	69.99%	47.69%	70.03%
p40-06	16.74%	18.36%	66.39%	95.26%	67.48%	95.28%
p40-07	18.17%	20.25%	56.53%	83.73%	57.19%	83.77%
p40-08	20.86%	22.47%	67.46%	92.58%	68.02%	92.48%
p40-09	17.21%	17.42%	47.03%	80.21%	48.23%	80.01%
p40-10	14.22%	13.88%	60.87%	89.88%	61.60%	89.93%

Our experiments revealed that the separation of maximally violated mod-k inequalities did not have a too big effect on the root bound and is, in general, not worth the effort compared to mod-2 separation. Mod-2- separation gives a better gap closure, and because of the better bounds the number of branch-and-nodes needed to solve the problems to optimality is considerably less (even when rotation is not invoked). Therefore, we did not use mod-k separation anymore. Table 5.8 displays the respective results of our computations.

Table 5.8 CPU times (min:sec), number of tree nodes and percentage of separation time

Name	3-cyc			M2			M2$_R$		
	CPU	#nod	sep	CPU	#nod	sep	CPU	#nod	sep
p40-01	0:00	3	30%	0:01	1	70%	0:01	1	75%
p40-02	6:57	5441	4%	26:26	147	87%	34:51	41	91%
p40-03	0:08	115	6%	0:32	5	79%	0:23	1	79%
p40-04	3:18	2317	4%	12:48	69	86%	13:56	19	89%
p40-05	1:41:36	24317	3%	1:29:39	471	85%	1:53:12	73	93%
p40-06	0:48	609	4%	3:02	19	84%	2:03	3	84%
p40-07	4:22	2775	3%	15:47	85	87%	12:05	15	87%
p40-08	1:57	1287	4%	6:21	33	85%	4:45	3	87%
p40-09	7:42	5625	4%	31:57	177	87%	19:41	23	89%
p40-10	0:33	421	4%	3:14	19	86%	3:41	7	87%

Table 5.8 exhibits the following facts. If only 3-dicycle separation is used, the algorithm spends most of its CPU time for solving the linear programs. Mod-2 separation changes this relation. Now separation is responsible for the CPU time. Because of the better bounds, the number of branch-and-cut nodes is drastically decreased. However, since separation is time consuming, the overall solution time for the problems with $n = 40$ is not reduced. But this changes for larger problems. E.g. instance p50-01 can now be solved in half of the time.

Our computational experiments for the linear ordering problem lead us to the conclusion that the incorporation of general cut generation procedures is worthwhile and promising, and should also be tried for other combinatorial optimization problems. The optimal use of mod-2 and mod-k inequalities still has to be explored. To some extent their potential cannot be fully exploited because LPs are becoming more difficult, at least for current LP solvers.

Chapter 6
The Linear Ordering Polytope

Abstract So far we developed a general integer programming approach for solving the LOP. It was based on the canonical IP formulation with equations and 3-dicycle inequalities which was then strengthened by generating mod-k-inequalities as cutting planes. In this chapter we will add further ingredients by looking for problem-specific inequalities. To this end we will study the convex hull of feasible solutions of the LOP: the so-called linear ordering polytope.

6.1 Polyhedral Combinatorics and Basic Results

We recall some necessary definitions and notations. If E is some set we denote by $\mathbf{R}^E$ the set of real vectors having $|E|$ components which are indexed by the members of E. The vector with all components equal to 1 is denoted by $\mathbf{1}$ and the zero vector is denoted by $\mathbf{0}$. The unit matrix is denoted by I and the *k-th unit vector* by e_k. We do not usually state explicitly the row and column dimensions of a matrix or the dimension of a vector if they are clear from the context.

Let $m > 0$ and $x_1, \ldots, x_m \in \mathbf{R}^n, \alpha_1, \ldots, \alpha_m \in \mathbf{R}$. A linear combination $\sum_{i=1}^m \alpha_i x_i$ is called an *affine combination* if $\sum_{i=1}^m \alpha_i = 1$, a *convex combination* if $\sum_{i=1}^m \alpha_i = 1$ and $\alpha_i \geq 0, i = 1, \ldots, m$, and a *conic combination* if $\alpha_i \geq 0, i = 1, \ldots, m$. If $S \subseteq \mathbf{R}^n$, then the *affine hull* aff(S) of S is defined as the set of all (finite) affine combinations of elements of S. Similarly the *convex hull* conv(S) and the *conic hull* cone(S) are defined. By definition we have aff$(\emptyset) = \text{conv}(\emptyset) = \emptyset$. If $S = \text{aff}(S)$ then S is called an *affine space*. The *affine rank* arank$S$ of a set S is the smallest cardinality of a set X such that $S \subseteq \text{aff}(X)$. If $\mathbf{0} \in \text{aff}(S)$ then we have arank$S = \text{rank}S + 1$; otherwise the affine rank is equal to the linear rank.

For $a \in \mathbf{R}^n \setminus \{\mathbf{0}\}$ and $a_0 \in \mathbf{R}$ the set $\{x \in \mathbf{R}^n \mid a^T x = a_0\}$ is called a *hyperplane*. A hyperplane defines the *halfspace* $\{x \in \mathbf{R}^n \mid a^T x \leq a_0\}$. A *polyhedron* (or $\mathscr{H}$-*polyhedron*) P is defined as the intersection of finitely many halfspaces or equivalently as the solution set of a finite system of linear inequalities. More precisely, P is an $\mathscr{H}$-polyhedron if there exists a matrix $A \in \mathbf{R}^{m \times n}$ and a vector $b \in \mathbf{R}^m$ such that

$$P = \{x \mid Ax \leq b\}.$$

We use the abbreviation $P(A,b)$ for the set $\{x \mid Ax \leq b\}$.

A description of a polyhedron by means of linear inequalities is also called the *outer description*. Note that some inequalities may actually be equations. If we want to emphasize that equations are present in the description of P, we explicitly write $P = \{x \mid Ax \leq b, Bx = d\}$.

A set P is a called a $\mathscr{V}$-*polyhedron* if there exist finite sets X and Y such that

$$P = \text{conv}(X) + \text{cone}(Y),$$

i.e., P consists of all vectors $z = x + y$ where $x \in \text{conv}(X)$ and $y \in \text{cone}(Y)$. This type of description is called the *inner description* with *generating sets* X and Y.

According to classical results in polyhedral theory, a set P is a $\mathscr{V}$-polyhedron if and only if it is an $\mathscr{H}$-polyhedron.

A *polytope* is a bounded polyhedron, i.e., $P \subseteq \mathbf{R}^n$ is a polytope if and only if it is a polyhedron and there exist bounds $l, u \in \mathbf{R}^n$ such that $l \leq x \leq u$ for all $x \in P$. In particular, P is a polytope if and only if it is equal to the convex hull of a finite set.

The *dimension* of a set S is defined as $\dim S = \text{arank} S - 1$. (The empty set has dimension -1.) The dimension of a polyhedron P is obtained as follows. If $Bx = d$ is a system of equations such that $\text{aff}(P) = \{x \mid Bx = d\}$, then $\dim P = n - \text{rank} B$, where $\text{rank} B$ is the usual rank of the matrix B. A polyhedron $P \subseteq \mathbf{R}^n$ is *full dimensional* if $\dim P = n$. Therefore, if P is full dimensional, then there exists no equation $a^T x = a_0$, with $a \neq \mathbf{0}$, satisfied by all points in P.

An inequality $a^T x \leq a_0$, $a \neq \mathbf{0}$, is said to be *valid* for a polyhedron P if $P \subseteq \{x \mid a^T x \leq a_0\}$. If $a^T x \leq a_0$ is valid then the set $F = P \cap \{x \mid a^T x = a_0\}$ determines a *face* of P (which may be empty). If $F \neq P$ then F is called a *proper face*. If $|F| = 1$ then the element $v \in F$ is called a *vertex* of P. Vertices cannot be represented as convex combinations of other elements of the polyhedron. If P is a polytope, then $P = \text{conv}(V)$ where V is the set of vertices of P. A maximal nonempty proper face F is called a *facet* of P. A face F is a facet of P if and only if $\dim F = \dim P - 1$. If F is a facet we call $a^T x \leq a_0$ a *facet defining* inequality for P. If $a^T x \leq a_0$ and $b^T x \leq b_0$ are inequalities defining the same facet of P, then one can be obtained from the other by multiplication by a positive constant and by adding multiples of equations valid for P. For a full dimensional polyhedron, facet defining inequalities are unique up to positive multiples.

The following theorem shows two ways for proving that an inequality is facet defining.

Theorem 6.1. *Let P be a polyhedron and $b^T x \leq b_0$ a valid inequality such that $F = \{x \in P \mid b^T x = b_0\}$ is a proper face of P. Let $Dx = d$ be a minimal equation system and $\text{aff}(P) = \{x \mid Dx = d\}$. Then the following statements are equivalent.*

  (i) *F is a facet of P.*
  (ii) *$\dim F = \dim P - 1$.*
  (iii) *If $F \subset \{x \in P \mid c^T x = c_0\}$ where $c^T x \leq c_0$ is a valid inequality for P then there exist $\lambda \in \mathbf{R}^{n-\dim P}$ and $\mu > 0$ such that $c^T = \mu b^T + \lambda^T D$.*

The main idea of polyhedral combinatorics for solving a combinatorial optimization problem $(E, \mathcal{I}, c)$ is to associate a polytope with it as follows. For $I \in \mathcal{I}$ we define the corresponding *characteristic vector* (or *incidence vector*) χ^I by setting

$$\chi_e^I = \begin{cases} 0, & \text{if } e \in I, \\ 1, & \text{otherwise.} \end{cases}$$

The polytope $P_{\mathcal{I}}$ associated with $(E, \mathcal{I}, c)$ is

$$P_{\mathcal{I}} = \text{conv}\left(\{\chi^I \mid I \in \mathcal{I}\}\right).$$

For the LOP we define the *linear ordering polytope* P_{LO}^n as the convex hull of the characteristic vectors of the acyclic tournaments in D_n, i.e.,

$$P_{LO}^n = \text{conv}\left(\{\chi^T \in \{0,1\}^{n(n-1)} \mid T \subset A_n \text{ is an acyclic tournament }\}\right).$$

Hence the vertices of P_{LO}^n correspond exactly to the linear orderings of n objects.

If P_{LO}^n were explicitly known, then the LOP could be solved as the linear programming problem

$$\max\{c^T x \mid x \in P_{LO}^n\}.$$

However, to be able to apply linear programming techniques, the above definition is useless. P_{LO}^n has to be represented as an $\mathcal{H}$-polyhedron. Therefore, it is the main goal of this chapter to study the linear description of P_{LO}^n. The best such description would be a so-called *minimal linear description* $P_{LO}^n = \{x \mid Ax \leq b, Bx = d\}$ where $\text{aff}(P_{LO}^n) = \{x \mid Bx = d\}$ with B of full row rank and where the inequality system $Ax \leq b$ contains exactly one defining inequality for every facet of P_{LO}^n.

Some basic general properties of the linear ordering polytope are easily derived.

Theorem 6.2. *Let $n \geq 2$. Then the system*

$$x_{ij} + x_{ji} = 1, \text{ for all } i, j \in V_n, i < j,$$

is a minimal equation system for P_{LO}^n.

Due to the minimal equation system the dimension of P_{LO}^n is $\binom{n}{2}$. Since P_{LO}^n is not full-dimensional, a facet defining inequality can be represented in different ways. However, the structure of the equation system allows a simple standard representation of inequalities to be defined.

Theorem 6.3. *For every facet of P_{LO}^n there exists an inequality $a^T x \leq \alpha$ defining it such that the vector a has nonnegative integral coefficients and the property that for every pair of nodes $i, j \in V_n$ at least one of the coefficients a_{ij} or a_{ji} is equal to zero.*

We can use this observation to define a *normal form* for facet defining inequalities. Namely, every facet can be represented uniquely by an inequality $a^T x \leq \alpha$ such that all coefficients a_{ij} are nonnegative coprime integers and $\min\{a_{ij}, a_{ji}\} = 0$ for every pair of nodes $i, j \in V_n$. The set of arcs corresponding to positive coefficients is called the *support* of a.

An important general question is to decide whether or not two facet defining inequalities define the same facet. A sufficient condition for the nonequivalence of two inequalities is given in the following theorem.

Theorem 6.4. *Let $a^T x \leq \alpha$ and $b^T x \leq \beta$ be facet defining inequalities for P_{LO}^n, $n \geq 2$, given in normal form. If there exists an arc $(i, j) \in A_n$ with $a_{ij} > 0$ and $b_{ij} = 0$ (or $b_{ij} > 0$ and $a_{ij} = 0$) then the inequalities define different facets.*

Two useful general properties of facet defining inequalities for P_{LO}^n are stated in the following two lemmas.

Theorem 6.5 (Trivial Lifting Lemma). *Let $a^T x \leq \alpha$ be facet defining for P_{LO}^n, $n \geq 2$. Define the vector $\bar{a} \in \mathbf{R}^{(n+1)n}$ by setting $\bar{a}_{ij} = a_{ij}$ for all $(i, j) \in A_n$, and $\bar{a}_{i,n+1} = \bar{a}_{n+1,i} = 0$, for $i = 1, \dots, n$. Then $\bar{a}^T x \leq \alpha$ defines a facet of P_{LO}^{n+1}.*

Theorem 6.6 (Arc Reversal Lemma). *Suppose $a^T x \leq \alpha$ is a facet defining inequality for P_{LO}^n, $n \geq 2$. If $b \in \mathbf{R}^{n(n-1)}$ is defined by $b_{ij} = a_{ji}$ for all $(i, j) \in A_n$, then $b^T x \leq \alpha$ is facet defining for P_{LO}^n.*

Therefore the linear ordering polytope P_{LO}^{n+1} inherits all facets from P_{LO}^n in the sense that the added coefficients can just be set to zero. Furthermore, reversing the arcs in the support of a facet defining inequality yields a new facet defining inequality. However, if the reversed digraph is isomorphic to the original support digraph then the new facet basically has the same structure because only the numbering of the nodes is different.

6.2 Facets of the Linear Ordering Polytope

The facial structure of P_{LO}^n has been investigated in many publications [22, 15, 60, 111, 90, 121, 48, 42, 47]. Interestingly, there has been a lot of independent research because the linear ordering problem occurs with different names in several fields. In this section we review some of the results. We usually give no proofs (except for showing some principles) and do not give definitions precisely that are of minor importance for our exposition.

A simple class of inequalities for P_{LO}^n is given by the so-called **trivial inequalities** $0 \leq x_{ij} \leq 1$, for all $(i, j) \in A_n$, which are always valid for polytopes with 0/1-vertices.

Theorem 6.7. *Trivial inequalities define facets of P_{LO}^n for all $(i, j) \in A_n$. No two of the inequalities $x_{ij} \leq 1$ are equivalent. The normal form of inequality $-x_{ij} \leq 0$ is $x_{ji} \leq 1$.*

The minimal equation system and the trivial inequalities are sufficient for describing P_{LO}^2. Further inequalities are needed for $n \geq 3$. Clearly, inequalities excluding dicycles in tournaments should play a central role.

Theorem 6.8. *For every dicycle C of length three in A_n, the inequality $x(C) \leq 2$ defines a facet of P_{LO}^n.*

Inequalities $x(C) \leq |C| - 1$ for longer dicycles are obviously valid for P_{LO}^n. Let $|C| > 3$ and let i, j be a pair of dicycle nodes with $(i, j) \notin C$ and $(j, i) \notin C$. Then C can be partitioned into C_1 and C_2 such that $C_1 \cup \{(i, j)\}$ and $C_2 \cup \{(j, i)\}$ are dicycles. Now $x(C) = x(C_1) + x(C_2) + x_{ij} + x_{ji} \leq |C_1| - 1 + |C_2| - 1 + 1 = |C| - 1$. All dicycle inequalities for dicycles longer than three are thus implied by 3-dicycle inequalites and therefore do not define facets.

The polytope P_{LO}^3 is contained in $\mathbf{R}^6$, but using the equation system it can be projected to $\mathbf{R}^3$. Figure 6.1 shows the projection of P_{LO}^3. The remaining variables are x_{12}, x_{13} and x_{23}. The vertices are labeled by the corresponding linear orderings.

Fig. 6.1 The (projected) polytope P_{LO}^3

The minimal equation system, trivial inequalities and 3-dicycle inequalities are sufficient to completely describe P_{LO}^3, P_{LO}^4, and P_{LO}^5. Hence their respective number of facets is 8, 20, and 40. In earlier publications [22, 15] it was believed that, in general, the polytope P_C^n defined by equations, trivial and 3-dicycle inequalities has only integral vertices. But this is not the case as already observed in [125].

Figure 6.2 shows a fractional vertex for P_C^6. Dotted edges represent pairs of antiparallel arcs the corresponding variables taking the value $\frac{1}{2}$.

Fig. 6.2 A fractional vertex of P_C^6

The study of this example gives rise to a first class of further inequalities.

Definition 6.1. A digraph $D = (V, F)$ is called a *k-fence* if it has the following properties:

(i) $|V| = 2k, k \geq 3$,
(ii) V can be partitioned into two subsets $U = \{u_1, \ldots u_k\}$ and $L = \{l_1, \ldots, l_k\}$ such that

$$F = \bigcup_{i=1}^{k} \left(\{(u_i, l_i)\} \cup \{(l_i, u_j) \mid j \in \{1, \ldots, k\}, j \neq i\} \right).$$

Every k-fence $D = (V, F)$ defines the *k-fence inequality* $x(F) \leq k^2 - k + 1$ which is valid for $P_{LO}^n, n \geq 2k$. Figure 6.3 shows a 3-fence.

Fig. 6.3 A 3-fence

Note that the fractional vertex shown in Fig. 6.2 violates the inequality associated with 3-fences by $\frac{1}{2}$.

Theorem 6.9. *Let $D = (V, F)$ be a k-fence contained in $D_n, n \geq 2k$. Then the k-fence inequality $x(F) \leq k^2 - k + 1$ defines a facet of P_{LO}^n.*

Proof. To illustrate the technique of how to show, without exhibiting a sufficiently large set of affinely independent vertices, that an inequality is facet defining we give an explicit proof.

Let $D = (V, F)$ be a k-fence, $k \geq 3$. Assume that $U = \{1, 2, \ldots, k\}$ is the set of its upper nodes and $L = \{k+1, k+2, \ldots, 2k\}$ is the set of its lower nodes. We call the arcs $(i, k+i)$, $i = 1, \ldots, k$, *pales* and the other arcs *pickets*. Due to Lemma 6.5 it suffices to show that $x(F) \leq k^2 - k + 1$ is facet defining for P_{LO}^{2k}. Denote this k-fence inequality by $a^T x \leq a_0$.

Suppose that $\{x \in P_{LO}^{2k} \mid b^T x = b_0\} \supseteq \{x \in P_{LO}^{2k} \mid a^T x = a_0\}$ for some inequality $b^T x \leq b_0$ valid for P_{LO}^{2k}. If we can show that there exists a vector $\lambda \in \mathbf{R}^{\binom{2k}{2}}$ and a scalar μ with $b^T = \mu a^T + \lambda^T H$ (where $Hx = \mathbf{1}$ denotes the minimal equation system) then we are done by Theorem 6.1.

Let $\overline{F} = A_n \setminus \{(u, v) \mid (u, v) \in F \text{ or } (v, u) \in F\}$. Because of the structure of the minimal equation system we can make the following assumptions without losing generality

$$b_{ij} = a_{ij} = 1, \text{ for } (i, j) \in F,$$
$$b_{ij} = a_{ij} = 0, \text{for } (i, j) \in \overline{F} \text{ and } i < j.$$

By reversing $k-1$ pales of D or $k-2$ pales and one picket having no endnode with any of these $k-2$ pales in common we obtain an acyclic digraph which induces a partial ordering of the nodes of D. Every extension of such a partial ordering to a linear ordering gives an acyclic tournament T the incidence vector of which satisfies $a^T \chi^T = a_0$. Let u, v be two nodes in U (or V). If we can find a partial ordering (by reversing exactly $k-1$ arcs of D) which implies neither $u \prec v$ nor $v \prec u$ then there exist linear extensions $T_1 = \langle \alpha, u, v, \beta \rangle$ and $T_2 = \langle \alpha, v, u, \beta \rangle$ satisfying $a^T \chi^{T_1} = a^T \chi^{T_2} = a_0$ and therefore $b^T \chi^{T_1} = b^T \chi^{T_2} = b_0$. From this we get

$$0 = b^T \chi^{T_1} - b^T \chi^{T_2} = b_{uv} - b_{vu}.$$

We first show that $b_{21} = b_{12} = 0$. By reversing all pales except for pale $(k, 2k)$ we obtain a partial ordering of the nodes of D which leaves nodes 1 and 2 incomparable (both nodes are sinks in the corresponding digraph) and hence by the above argument $b_{12} = b_{21} = 0$. Repeating this construction for all pairs of nodes in U, resp. in V yields $b_{ij} = 0$ for all $(i, j) \in \overline{F}$.

We now show that there is some scalar $\xi \in \mathbf{R}$ such that $b_{ij} = \xi$ for all arcs (i, j) with $(j, i) \in F$. W.l.o.g. we assume that $b_{k+1,1} = \xi$

The reversal of all pales except for pale $(2, k+2)$ gives a partial ordering which can be extended to the linear ordering

$$T_1 = \langle k+3, k+4, \ldots, 2k, k+1, 2, k+2, 1, 3, 4, \ldots, k \rangle.$$

The reversal of the picket $(k+2, 1)$ and all pales except for the two pales $(1, k+1)$ and $(2, k+2)$ induces a partial ordering which can be extended to the linear ordering

$$T_2 = \langle k+3, k+4, \ldots, 2k, 1, k+1, 2, k+2, 3, 4, \ldots, k \rangle.$$

From this we get

$$\begin{aligned}
0 &= b^T \chi^{T_1} - b^T \chi^{T_2} \\
&= b_{k+1,1} + b_{21} + b_{k+2,1} - b_{1,k+1} - b_{12} - b_{1,k+2} \\
&= \xi + 0 + 1 - 1 - 0 - b_{1,k+2}
\end{aligned}$$

and hence $b_{1,k+2} = \xi$.

Using similar constructions we eventually obtain $b_{ij} = \xi$ for all arcs antiparallel to the arcs of the k-fence.

Defining $\lambda \in \mathbf{R}^{\binom{2k}{2}}$ with components $\lambda_{\{i,j\}}$ for $i < j$ by

$$\lambda_{\{i,j\}} = \begin{cases} 0, & \text{if } (i, j) \in \overline{F}, \\ \xi, & \text{otherwise,} \end{cases}$$

and setting $\mu = 1 - \xi > 0$, we get $b^T = \mu a^T + \lambda^T H$. This finishes the proof. $\qquad \square$

According to Theorem 6.4 different k-fences induce different facets of P_{LO}^n. Hence the number of facets of P_{LO}^n, $n \geq 6$, which are induced by k-fences is

$$\sum_{k=3}^{\lfloor \frac{n}{2} \rfloor} \left[\binom{n}{2k} \binom{2k}{k} k! \right] = \sum_{k=3}^{\lfloor \frac{n}{2} \rfloor} \frac{n!}{(n-2k)!k!}.$$

For $k > 3$ the k-fences can be considered as a generalization of the 3-fence. Looking at the 3-fence from a different point of view, namely by focusing on the structure of its dicycles, leads to another generalization and the rich class of Möbius ladders. Let $C_1, \ldots, C_k$ be different dicycles in $D_n = (V_n, A_n)$ such that

(i) C_i and C_{i+1}, for $i \in \{1, \ldots, k-1\}$, have exactly one arc in common. This arc is called e_i. C_k and C_1 have exactly the arc e_k in common.
(ii) C_i and C_j, for $j \notin \{i-1, i+1\}$ (resp. $j \notin \{k, 2\}$, if $i = 1$), have no common arc.

We define the digraph $D = (V, A)$ by setting $V = \cup_{i=1}^k V(C_i)$ and $A = \cup_{i=1}^k C_i$ and say that D is *generated* by the dicycles $C_1, \ldots, C_k$. Conversely, whenever a digraph $D = (V, A)$ is said to be generated by k dicycles, we implicitly assume that $V = \cup V(C_i)$, $A = \cup C_i$ and that the common arcs are denoted by e_i as in (i).

A dicycle C_j is called *right-adjacent* to a dicycle C_i if C_j and C_i have some node v in common and if all dicycles C_l contain this node, $l \in \{j, j+1, \ldots, i\}$, for $j < i$, and $l \in \{j, j+1, \ldots, k, 1, 2, \ldots, i\}$, for $j > i$. A dicycle C_j is called *left-adjacent* to a dicycle C_i if C_j and C_i have some node v in common and if all dicycles C_l contain this node, $l \in \{i, i+1, \ldots, j\}$, for $j > i$, resp. $l \in \{i, i+1, \ldots, k, 1, 2, \ldots, j-1, j\}$, for $j < i$. If C_j is both left- and right-adjacent to C_i then all dicycles have a common node.

Figure 6.4 gives a different drawing of the 3-fence showing that it can also be considered a Möbius ladder consisting of three 4-dicycles. The figure also illustrates why the name Möbius ladder was chosen.

Fig. 6.4 A 3-fence drawn as Möbius ladder

Not all digraphs of Möbius ladder structure yield facets. They have to be generated by an odd number of dicycles and these dicycles have to be short.

Definition 6.2. Let $D = (V, M)$ be a subdigraph of D_n which is generated by the k dicycles $C_1, \ldots, C_k$, i.e., $V = \cup V(C_i), M = \cup C_i$, and which satisfies the following properties:

 (i) $k \geq 3$ and k is odd.
 (ii) The length of C_i is three or four, $i = 1, \ldots, k$.
 (iii) The degree of each node $u \in V(M)$ is at least three.
 (iv) If two dicycles C_i and $C_j, 2 \leq i + 1 < j \leq k$, have a node, say v, in common then C_j is either left-adjacent or right-adjacent to C_i but not both.
 (v) Given any dicycle $C_j, j \in \{1, \ldots, k\}$, set $J = \{1, \ldots, k\} \cap (\{j - 2, j - 4, \ldots\} \cup \{j + 1, j + 3, \ldots\})$. Then the set $M \setminus \{e_i \mid i \in J\}$ contains exactly one dicycle, namely C_j.

Then D is called a *Möbius ladder*.

Theorem 6.10. *Let $D = (V, M)$ be a Möbius ladder in D_n generated by the k dicycles $C_1, C_2, \ldots, C_k$. Then the Möbius ladder inequality*

$$x(M) \leq |M| - \frac{k+1}{2}$$

defines a facet of P_{LO}^n for $n \geq |V|$.

Möbius ladder inequalities are mod-2 inequalities w.r.t. to trivial and dicycle inequalities. Namely, let $D = (V, M)$ be generated by $C_1, C_2, \ldots, C_k$ and let $F = \{e_1, \ldots, e_k\}$ be the set of common arcs of adjacent dicycles. Then the addition of the inequalities

$$x(C_1) \leq |C_1| - 1$$

$$\vdots$$

$$x(C_k) \leq |C_k| - 1$$

$$\sum_{e \in F} x_e \leq |M| - |F|$$

gives $2x(M) \leq 2|M| - k$ and hence, after dividing by 2 and rounding down the right hand side, the Möbius ladder inequality $x(M) \leq |M| - \frac{k+1}{2}$.

There are further classes of facet-defining and valid inequalities which we will not describe here:

 – *diagonal inequalites* [51],
 – *Z_m-inequalities* [111],
 – *t-reinforced k-fences* [90],
 – *augmented k-fences* [91],
 – *Paley inequalities* [58],
 – *facets from stability-critical graphs* [81, 40],
 – *generalizations of Möbius ladders* [46],
 – *new facets by rotations* [14].

Concerning separation, the situation is not too fortunate. The separation problem for k-fences is already NP-hard. The only inequalities so far which can be separated in a systematic way (and in polynomial time) are certain Möbius ladder inequalities as they are mod-2 inequalities [24].

We will discuss how to obtain further cutting planes in the subsequent sections.

6.3 Computation of Complete Descriptions

In addition to identifying classes of facet defining inequalities, there has always been interest in deriving complete linear descriptions for polytopes associated with small instances of combinatorial optimization problems. Though real problems are usually large scale, it is worthwhile to put investigations into small polytopes. Facets derived for small polytopes can give hints for generalizations to facets for larger polytopes and studying small polytopes can also yield information on the relative importance of the different classes of facets. Furthermore, in particular in the case of the LOP (because of the trivial lifting property), facet-defining inequalities for small instances are also facet-defining for large instances and can thus be used in computations.

In general, a $\mathcal{V}$-polyhedron $\mathrm{conv}(V) = \mathrm{conv}(\{v_1 \ldots, v_m\})$ can be transformed to an $\mathcal{H}$-polyhedron according to the following reformulation.

$$\mathrm{conv}(V) = \{x \mid \text{there exists } y \text{ such that } Vy = x, \mathbf{1}^T y = 1, y \geq 0\}$$

$$= \left\{ x \;\middle|\; \text{there exist } y \text{ such that } \begin{pmatrix} V & -I \\ -V & I \\ \mathbf{1}^T & 0 \\ -\mathbf{1}^T & 0 \\ -I & 0 \end{pmatrix} \begin{pmatrix} y \\ x \end{pmatrix} \leq \begin{pmatrix} 0 \\ 0 \\ 1 \\ -1 \\ 0 \end{pmatrix} \right\}$$

$$= \left\{ x \;\middle|\; \text{there exists } y \text{ such that } \begin{pmatrix} y \\ x \end{pmatrix} \in P(D,d) \right\},$$

where D and d are defined by the second-to-last row.

By projecting $P(D,d)$ onto the subspace $y = 0$ (by elimination of the y-variables) we obtain a polyhedron $P(A,b)$ with the property

$$x \in P(A,b) \iff \text{there exists } y \text{ such that } \begin{pmatrix} y \\ x \end{pmatrix} \in P(D,d)$$

$$\iff x \in \mathrm{conv}(V).$$

As a first step one should determine the minimal equation system for $\mathrm{conv}(V)$ by performing Gaussian elimination in the system

$$\begin{pmatrix} V & -I \\ \mathbf{1}^T & 0 \end{pmatrix} \begin{pmatrix} y \\ x \end{pmatrix} = \begin{pmatrix} 0 \\ 1 \end{pmatrix}.$$

The system is brought into an equivalent form

$$\begin{pmatrix} I & * & * \\ 0 & 0 & D' \end{pmatrix} \begin{pmatrix} y \\ x \end{pmatrix} = \begin{pmatrix} * \\ d' \end{pmatrix}$$

with D' as large as possible. The system $D'x = d'$ is an equation system for $\text{conv}(V)$. Possible redundant equations still contained in it can be removed to obtain a minimal system.

$P(A,b)$ can be obtained algorithmically by using Fourier-Motzkin elimination. The following algorithm eliminates the j-th variable from the system $Dx \leq d$ and yields a system $Ax \leq b$ which is defined in the remaining variables only and has feasible solutions if and only if $Dx \leq d$ is feasible.

FourierMotzkin(D, d, j)

(1) Let $k = 0$ and partition the rows $M = \{1, 2, \ldots, m\}$ of D into
$$N = \{i \in M \mid d_{ij} < 0\},$$
$$P = \{i \in M \mid d_{ij} > 0\},$$
$$Z = \{i \in M \mid d_{ij} = 0\}.$$
(2) For all $i \in Z$, set $k = k + 1$ and
$$A_{k\cdot} = D_{i\cdot} \text{ and } b_k = d_i.$$
(3) For all $(s,t) \in (N \times P)$ set $k = k + 1$ and
$$A_{k\cdot} = d_{tj}D_{s\cdot} - d_{sj}D_{t\cdot} \text{ and } b_k = d_{tj}d_s - d_{sj}d_t.$$
(4) The resulting system is $Ax \leq b$.

To turn this principle approach into an effective algorithm further ingredients have to be added, e.g. for avoiding redundancy and finding good elimination orders. All details can be found in [34]. Here we report on the most important findings for the linear ordering polytope.

The complete linear description of P_{LO}^6 consists of the 15 equations forming the minimal equation system, 30 trivial inequalities, 40 3-dicycle inequalities, 120 3-fence inequalities and 360 inequalites from each of the two types of Möbius ladders M_1 (defined on four 3-dicycles and one 4-dicycle) and M_2 (the reversal of M_1).

Figure 6.6 displays two facet defining inequalities of P_{LO}^7. All solid arcs shown have coefficient 1 in the inequality, the dotted arc has coefficient 2. The right hand side of the inequality on the left side is 13 and the other inequality has right hand side 14.

The complete linear description of P_{LO}^7 given in [112] consists of 87,472 facets. To visualize how many facets are structurally different we define equivalence classes. Two facet defining inequalities are said to belong to the same σ-*class* if one can be obtained from the other just by renaming the nodes. The vertices satisfying a facet defining inequality with equation are called *roots* of the facet.

$x_e \leq 1$

$x(C) \leq 2$

$x(F) \leq 7$ $x(M_1) \leq 8$ $x(M_2) \leq 8$

Fig. 6.5 Facet defining inequalities of P_{LO}^6

Fig. 6.6 Facet defining inequalities for P_{LO}^7

The facets of P_{LO}^7 can be partitioned into 27 σ-classes which are displayed in Table 6.1. Eight of these classes are obtained from other classes just by arc reversal. They are marked with a "*". The table also gives the number of vertices on each facet and the number of nonequivalent facets in each class.

Table 6.1 substantiates the importance of trivial (F1) and 3-dicycle (F2) inequalities. Every such inequality defines a facet with 2520 roots, i.e., a facet containing half of the vertices of P_{LO}^7. A 3-fence facet (F3) contains only 126 vertices and there are facets containing only a few roots more than required for the facet dimension.

A further partition of facets of P_{LO}^n into equivalence classes is possible. Let $P = \text{conv}(V)$ be a polytope in $\mathbf{R}^d$. An affine mapping ψ from $\mathbf{R}^d$ to itself is called a *rotation mapping for P* if $\psi(V) = V$. A rotation mapping transforms a facet F of P to a facet $\psi(F)$.

Table 6.1 The polyhedral structure of P_{LO}^7

Class	#vertices on facet	#different facets
F1	2,520	42
F2	2,520	70
F3	126	840
F4, F4*	126	5,040
F5, F5*	67	10,080
F6	44	5,040
F7, F7*	104	5,040
F8, F8*	67	10,080
F9, F9*	67	10,080
F10, F10*	44	5,040
F11	67	5,040
F12	104	5,040
F13	67	5,040
F14	126	840
F15, F15*	104	5,040
F16	104	5,040
F17, F17*	28	2,520
F18	28	2,520
F19	28	5,040

Arc reversal is one such rotation mapping. A second one was introduced in [14]. For fixed $r \in \{1,\dots,n\}$ the rotation mapping ψ is defined by

$$\psi(x)_{rj} = x_{jr}, \text{ for all } 1 \le j \le n, j \ne r,$$
$$\psi(x)_{jr} = x_{rj}, \text{ for all } 1 \le j \le n, j \ne r,$$
$$\psi(x)_{ij} = x_{ij} + x_{jr} + x_{ri} - 1, \text{ for all } 1 \le i, j \le n, i \ne r, j \ne r.$$

It was shown in [14] that if $a^T x \le \alpha$ defines the facet F of P_{LO}^n then

$$\sum_{i=1, i \ne r}^{n} \sum_{j=1, j \ne r}^{n} a_{ij} \psi(x)_{ij}$$
$$= \sum_{i=1, i \ne r}^{n} \left(\sum_{j=1, j \ne r}^{n} a_{ij}(x_{ij} + x_{jr} - x_{ir}) + a_{ir}x_{ri} + a_{ri}x_{ir} \right)$$
$$\le \alpha$$

defines the facet $\psi(F)$ of P_{LO}^n. E.g. this mapping transforms a 3-dicycle inequality to a trivial inequality and vice versa. We say that two facets belong to the same P_{LO}^n-*class* if they can be converted to each other either by renaming the nodes or by applying a rotation mapping.

Using a parallel computer and the adjacency approach described in [34] we could compute at least a lower bound on the number of facets of P_{LO}^8. We could not prove that this lower bound gives the exact number of facets, because in 11 out of 12,231 cases we could not compute the adjacent facets. These cases are the facets with the maximum number of roots (among them the facets defined by trivial, 3-dicycle, 3-fence and 4-fence inequalities). At least 67.5% of the facet σ-classes were first discovered by our computations. This is a very conservative estimation, based on the observation that no facets of P_{LO}^8 with coefficients larger than 2 were known before, while for that percentage of facet σ-classes the minimal coefficient is 3.

Table 6.2 summarizes the current state of knowledge about linear ordering polytopes.

Table 6.2 Facet structure of P_{LO}^n

n	#vertices	#different facets	#σ-classes	P_{LO}^n-classes
3	6	8	2	1
4	24	20	2	1
5	120	40	2	1
6	720	910	5	2
7	5,040	87,472	27	6
8	40,320	$\geq 488{,}602{,}996$	$\geq 12{,}231$	$\geq 1{,}390$

It is impressive how huge the number of facets of combinatorial polytopes is even for small instances. And, moreover, none of them is superfluous and they are also present in larger instances. We will therefore below turn to the question whether we can make use of this wealth of inequalities in practical computations.

6.4 Differences between Facets

It is a natural question which classes of facets are most or least useful in cutting plane algorithms. For reasons of efficiency, the choice of the inequalities to use in a cutting plane algorithm is typically dictated by whether efficient exact or heuristic algorithms are known for the corresponding separation problem. Up to now no ultimate measure of the quality of a valid inequality with respect to its application in a branch-and-cut algorithm could be established. From our computations of descriptions of small polytopes we now have 12231 different classes of (necessary!) facet defining inequalities available for all linear ordering problems on at least 8 nodes. Even on parallel hardware it does not seem to make sense to call a separation procedure for every class. In the following we therefore want to exhibit differences between the facets in order to possibly obtain insight into their usefulness for branch-and-cut algorithms.

Our goal is to find out if there is a measure such that a priority rule can be given for the application of the facets of a small problem instance relaxation in a branch-and-cut algorithm.

In [57, 58] the notion of *strength of a relaxation* was introduced. The strength of a relaxation is meant to be a measure of how well a relaxation approximates a polyhedron in comparison to another weaker relaxation. The strength is only defined for certain types of combinatorial polyhedra, namely polyhedra of blocking type [57] and of anti-blocking type [58]. In the case of the LOP the anti-blocking type is of interest.

Definition 6.3. Given polytopes P and Q of anti-blocking type, P is said to be an α-relaxation of Q for some $\alpha \geq 1$ if $Q \supseteq \frac{1}{\alpha}P = \{\frac{x}{\alpha} \mid x \in P\}$. The strength $\text{str}(P,Q)$ of Q with respect to P is the minimum value of α such that P is an α-relaxation of Q.

Notice that $\text{str}(P,Q) \geq 1$ and $\text{str}(P,Q) = 1$ if and only if $P = Q$. In general, $\text{str}(P,Q)$ could be infinite.

While P_{LO}^n is not of blocking or anti-blocking type we can apply this concept because of its close relation to the acyclic subdigraph polytope P_{AC}^n. The *acyclic subdigraph polytope*

$$P_{AC}^n = \text{conv}\{\chi^B \in \{0,1\}^{n(n-1)} \mid B \text{ is an acyclic arc set in } A_n\}$$

is of anti-blocking type.

Because $P_{LO}^n = \{x \in P_{AC}^n \mid x_{ij} + x_{ji} = 1, 1 \leq i < j \leq n\}$, the linear ordering polytope P_{LO}^n is a face of P_{AC}^n. For any nonnegative objective function, an optimal solution of the linear ordering problem is an optimal solution of the acyclic subdigraph problem.

Furthermore, a facet defining inequality for P_{LO}^n in normal form is valid, resp. facet defining for P_{AC}^n. Hence, if we restrict ourselves to nonnegative objective functions (which we can do w.l.o.g. for the LOP), we can interpret a cutting plane algorithm for solving the *linear ordering problem* as a cutting plane algorithm for the *acyclic subdigraph problem* using exclusively facet defining inequalities for P_{LO}^n. This observation justifies the discussion of strength of relaxations given by inequalities for P_{LO}^n.

The *strength of an inequality* with respect to a polytope P is the strength of the relaxation obtained by adding the inequality. We will compute the strength of classes of inequalities with respect to the following polytopes.

Trivial relaxation:

$$P_T^n = \{x \mid x_{ij} + x_{ji} = 1, \text{ for all } 1 \leq i, j \leq n,$$
$$x_{ij} \geq 0, \text{ for all } 1 \leq i, j \leq n\}.$$

Dicycle relaxation:

$$P_C^n = \{x \mid x_{ij} + x_{ji} = 1, \text{ for all } 1 \le i, j \le n,$$
$$x_{ij} \ge 0, \text{ for all } 1 \le i, j \le n$$
$$x(C) \le 2, \text{ for all dicycles } C \text{ of length } 3\}.$$

Obviously, $P_C^n = P_T^n \cap \{x \mid x(C) \le 2, \text{ for all dicycles } C \text{ length } 3\}$.

Let $f^T x \le f_0$ be a facet defining inequality for P_{LO}^n in normal form. Following [58] we define the *trivial strength* of this inequality as

$$s_T(f) = \frac{\max\{f^T x \mid x \in P_T^n\}}{\max\{f^T x \mid x \in P_{LO}^n\}} = \frac{\max\{f^T x \mid x \in P_T^n\}}{f_0}$$

and the *dicycle strength* of f as

$$s_C(f) = \frac{\max\{f^T x \mid x \in P_C^n\}}{\max\{f^T x \mid x \in P_{LO}^n\}} = \frac{\max\{f^T x \mid x \in P_C^n\}}{f_0}.$$

Note that $\max\{f^T x \mid x \in P_T^n\} = \mathbf{1}^T f$ since $f^T x \le f_0$ is in normal form.

Figure 6.7 shows the trivial strengths of the σ-classes of facets of P_{LO}^8 (except for the trivial inequalities).

Fig. 6.7 Trivial strength of σ-classes of facets of P_{LO}^8

Figure 6.8 displays the only two facets of P_{LO}^8 with $s_T > 1.5$. The right hand sides of these facets are 17 and 13, the trivial strengths being $s_T = 1.52941$ and $s_T = 1.53846$, respectively. The 3-dicycle inequality has trivial strength 1.5.

Figure 6.9 displays the dicycle strength of the σ-classes of facets (except for the trivial and 3-dicycle inequalities). The values of $\max\{f^T x \mid x \in P_C^n\}$ were computed using the branch-and-cut algorithm for the linear ordering problem described earlier.

Fig. 6.8 Facet defining inequalities with $s_T = 1.52941$ and $s_T = 1.53846$

The 4-fence inequality is the facet with maximum dicycle strength $s_C = 1.07692$ among all facets of P_{LO}^8.

Fig. 6.9 Dicycle strength of σ-classes of facets of P_{LO}^8

As shown by Goemans and Hall [58], the trivial strength of P_{AC}^n is $2 - o(1)$, where $o(1)$ is nonnegative and tends to 0 as the number of nodes n tends to infinity. They prove that the highest trivial strength of the known inequalities of P_{LO}^n is attained asymptotically for the augmented k-fence inequality with a value of only 1.52777 and they present new valid inequalities, called Paley inequalities, which they prove to be facet defining for $n = 11$ and $n = 19$ with trivial strengths 1.57143 and 1.59813, respectively. They conclude that the strongest facets of the acyclic subdigraph polytope are unknown. Concerning the dicycle strength they

show that the value for P_{AC}^n must be at least $\frac{4}{3}$ since the Paley inequalities achieve this bound asymptotically. Figure 6.10 displays the Paley graph for 11 nodes.

Fig. 6.10 Paley graph on 11 nodes

A further measure for the usefulness of a facet (proposed by Naddef and Rinaldi in [101]) could be the distance of the facet to a given relaxation. For P_{LO}^n we consider the *distance with respect to the 3-dicycle relaxation* and define

$$d_c(f) = \frac{f^T x^* - f_0}{|f|},$$

where $f^T x^* = \max\{f^T x \mid x \in P_C^n\}$.

Figure 6.11 shows this distance for the facets of P_{LO}^8 (except for the trivial and 3-dicycle inequalities). The strongest facet in this sense is the 4-fence inequality with dicycle distance $d_c = 0.25$.

A further indication of the usefulness of a facet could be its volume. Volume computation is very difficult and we refer to [34] where volume computation methods are described and where it is shown that there is a correlation between the volume of facets and their number of roots.

If one examines correlations between these measures then it turns out that dicycle strength and dicycle distance are correlated whereas there seems to be a negative correlation between trivial and dicycle strength, i.e., a facet with large dicycle strength has small trivial strength and vice versa. There seems to be a weak

Fig. 6.11 Distance of σ-classes of facets of P_{LO}^8 relative to P_C^8

correlation between number of roots and trivial strength but no correlation between number of roots and dicycle strength or distance.

6.5 Separation of Small Facets

We now describe a possibility for employing the "small" facets in practical computations. We do not attempt to generalize promising classes because the separation problem for well-structured inequalities like fences and Möbius ladders is already hard. Of course, the study of generalizations is an interesting topic for theoretical research. We rather develop a general procedure which has the chance to find violated inequalities of any facet class.

Let $f^T x \leq f_0$ be a small facet defining inequality for the LOP (in normal form) with support graph $D_f = (V_f, A_f)$. If x^* is the current fractional solution, we have to identify a subgraph $D' = (V', A')$ of D_n with $|V'| = |V_f|$ and an isomorphism between V' and V_f such that $f^T x^* > f_0$ on this subgraph.

We use the following method for identifying the subset V' and the isomorphism.

Small facet separation

(1) Choose a subset of nodes $W \subseteq V$, $|W| = |V_f|$ and a 1-1 correspondence
between V_f and W. This gives an initial configuration and a configura-
tion value obtained by evaluating $f^T x^*$ according to the isomorphism
(left hand side of the small facet defining inequality).
(2) As long as possible, generate configurations with higher value. This
can be done by varying the correspondence between W and V_f or by
replacing nodes of W by nodes of $V \setminus W$.
(3) If the value is greater than f_0, a cutting plane is found.

In principle, it is possible to find the best configuration by solving a quadratic
assignment problem with $|V| \times |V_f|$ variables. But since the quadratic assignment
problem is NP-hard in general, we have to use heuristics. In the following, let the
current configuration be represented by a mapping $\sigma : V \rightarrow V_f \cup \{0\}$ with $\sigma(w) = u$,
if $w \in W$ and u is the node of V_f associated with w in the 1-1 correspondence, and
$\sigma(w) = 0$, if $w \in V \setminus W$.

Figure 6.12 illustrates the principle problem and how to proceed with local mod-
ifications. Here a violated facet defining inequality from a facet class with a support
graph on 7 nodes is searched for.

Fig. 6.12 Separation of small facets

We can use virtually any of the heuristic or meta-heuristic approaches described
in this book. For our computational experiments we have chosen several methods.
Here we briefly describe a deterministic improvement and a simulated annealing
heuristic. A more elaborated GRASP algorithm is given in [34].

Deterministic Improvement

(1) For every pair of nodes $u, v \in V$ perform the following:

   (1.1) Define a new configuration σ' by setting $\sigma'(u) = \sigma(v)$, $\sigma'(v) = \sigma(u)$, and $\sigma'(w) = \sigma(w)$ for $w \in V \setminus \{u,v\}$.
   (1.2) If the new configuration has higher value, set $\sigma = \sigma'$.

(2) If the current configuration could be improved for at least one pair, then repeat step (1).
(3) For every triple of nodes $u, v, w \in W$ perform the following:

   (3.1) Consider all possible new configurations that can be obtained by exchanging the assignments of u, v and w.
   (3.2) If one of these new configurations, say σ', has higher value than the current configuration, set $\sigma = \sigma'$.

(4) If the current configuration could be improved for at least one triple, then repeat step (1).

This deterministic heuristic follows the principle of local search and can be characterized as a combination of 2-node and 3-node exchanges. Obviously, the modification capabilities of this heuristic are limited, but based on experience from many local search algorithms one can expect that reasonable configurations are determined.

In our separation approach we do not exploit special structural information about small facets to find good configurations. For such unstructured searches, usually also meta-heuristic algorithms give good results. The second approach is a simulated annealing scheme.

Simulated Annealing

(1) Choose an initial parameter ϑ, a repetition factor r, and a stopping parameter ε.
(2) Perform the following as long as $\vartheta > \varepsilon$:

   (2.1) Repeat the following steps r times
      (2.1.1) Let σ be the current configuration. Choose two nodes $u, v \in V$ at random, and define a new configuration σ' by setting $\sigma'(u) = \sigma(v)$, $\sigma'(v) = \sigma(u)$, and $\sigma'(w) = \sigma(w)$ for $w \in V \setminus \{u,v\}$. Let Δ be the difference of the new configuration value and the old one.
      (2.1.2) If $\Delta > 0$ set $\sigma = \sigma'$.
      (2.1.3) If $\Delta \leq 0$, compute a random number p, $0 \leq p \leq 1$ and set $\sigma = \sigma'$ if $p \leq \frac{e^{\Delta}}{\vartheta}$.

(3) Update ϑ and r.

To give an impression of how many violated inequalites one can expect from small facet separation we list in Tables 6.3 and 6.4 the number of inequalities found for the facet classes of P_{LO}^7.

Table 6.3 Number of small facets of classes F3–F11 found

Problem	F3	F4	F5	F6	F7	F8	F9	F10	F11
econ59	26	163	11	1	1	–	105	3	–
econ64	1	47	1	–	1	–	9	3	–
econ71	1	8	–	–	–	–	35	–	–
econ72	19	34	2	2	5	3	66	8	–
econ76	4	147	1	–	1	–	17	3	–
econ77	–	42	9	–	25	17	49	10	–
randB	195	1435	41	11	103	82	134	28	50
randC	130	1035	26	2	55	37	54	11	30

In these experiments problem instances were selected which could not be solved at the root node with 3-dicycle inequalities only. Small facet separation was invoked when no more 3-dicycle inequalities were violated. Now, all problems could be solved at the root node.

Table 6.4 Number of small facets of classes F12–F19 found

Problem	F12	F13	F14	F15	F16	F17	F18	F19
econ59	1	6	108	67	11	102	15	–
econ64	1	–	18	6	3	5	1	–
econ71	1	3	26	24	1	13	–	–
econ72	1	1	79	34	19	55	6	–
econ76	–	–	28	18	–	9	2	–
econ77	–	4	141	24	–	74	17	–
randB	75	23	107	134	70	1	4	–
randC	36	17	36	63	25	2	3	–

For the linear ordering problem we make use of the fact that the facets of small polytopes are globally facet defining even when lifted trivially. Fractional LP solutions for the linear ordering problem are dense, since we are working on the complete directed graph and since a solution vector contains $\binom{n}{2}$ arcs with value 1. Therefore, we could not develop a reasonable shrinking procedure and perform small facet separation in the complete digraph with arc weights given by the fractional solution.

6.6 Computational Experiments with Small Facets

Extensive computational experiments have been carried out to find the best strategy for employing facets from small relaxation. They are documented in depth in [34, 35] and we only cite the major insights here.

The linear ordering problems were always solved as follows. In a first phase, only 3-dicycle inequalities were generated. If not enough 3-dicycle cuts are found, then the heuristics for small facet generation are invoked. Because of the huge number of facet classes various approaches have been tested. All computations were carried out on a PC cluster where one processor was the master processor handling the core of the branch-and-cut and the other processors solved the separation problem for small facets. The facet classes taken into account were partitioned on the available processors.

6.6.1 Comparison of Heuristics

It turned out that simulated annealing and GRASP were more effective that the deterministic approach. Annealing and GRASP were of about the same quality with respect to finding inequalties. Because of the easier tuning of parameters, eventually GRASP was preferred.

6.6.2 Cutting Plane Selection

Usually, at each LP phase in the branch-and-cut algorithm many violated inequalities were generated. This was in particular the case because many parallel processors were active. It is well-known that it is not advisable to add all violated inequalities because the linear programs become more difficult. The following strategies were tested.

S1 Select some cuts at random.

S2 Select cuts with priority depending on the amount of violation $\delta = f^T x^* - f_0$ (higher violation preferred). This strategy is called *distance*.

S3 Select cuts with priority depending on the distance d between x^* and the hyperplane $f^T x = f_0$, i.e., $d = \frac{f^T x^* - f_0}{\|f\|} = \frac{\delta}{\|f\|}$.

S4 If $f = kc$ with $k > 0$, then $f^T x \leq f_0$ is obviously the "best" inequality that can be added to the linear program, since f_0 is an upper bound of the objective function which is attained for the roots of $f^T x \leq f_0$. Therefore, an idea is to prefer cutting planes $f^T x \leq f_0$ with f as parallel as possible to the objective function c, i.e., those which maximize $\frac{c^T f}{\|f\|}$. Moreover, it might be assumed that a hyperplane $f^T x = f_0$ which is nearly parallel to the direction of the objective

function can bound more symmetrically in all directions when reoptimizing. Since the angle ϕ between c and f satisfies $\cos\phi = \frac{c^T f}{\|c\|\|f\|}$, we call this strategy *angle*.

S5 Add cutting planes such that the "expected" following LP solution y^* minimizes $c^T x$. Since, of course, y^* is not known before reoptimizing, we take $-f$ as approximating the direction of reoptimization. Let $\delta = f^T x^* - f_0$. Then with $y^* = x^* - \lambda f$ and $f^T y^* = f_0$ one obtains $\lambda = \frac{\delta}{f^T f}$ and hence

$$c^T y^* = c^T x^* - \frac{\delta c^T f}{f^T f}.$$

Our experiments suggest using strategy S4. It usually leads to the least number of subproblems and also to the least CPU time. Note that this is also the preferrable cut selection strategy if only the 3-dicycle relaxation is solved.

6.6.3 Number of Classes Taken into Account

Of course, the gap between the LP bound and integer optimum decreases if more cutting planes are available. Figure 6.13 displays the gap closure achieved when using more and more facet classes where in the final case all classes are employed. It is seen that there is almost no effect any more if more than 1000 classes are used.

Fig. 6.13 Varying the number of P_{LO}^8 σ-classes for problem (EX2)

6.6.4 Facet Selection

A next point to be addressed is the question which of the 12229 σ-classes of facets of P_{LO}^8 should be used for separation in addition to trivial and 3-dicycle inequalities. We have seen above that they have different properties and it should be found out which of them are important.

In the experiments the 100 best facets relative to the following four criteria were chosen:

  C1 number of roots,
  C2 trivial strength,
  C3 dicycle strength,
  C4 dicycle distance.

Interestingly, it turned out that the number of roots (supposedly correlated to the volume) does not generally outperform the other choices. It seems that the trivial strength is a better indicator for the usefulness of a facet. But it was also observed that this is highly problem dependent.

As the final outcome of all experiments it turned out that, if CPU time in the parallel setup should be minimum, only the 14 facet classes with most roots should be taken into account and the 300 cuts best with respect to criterion "angle" or "distance" should be added in each phase. Despite parallelism, separation of small facets is still very time consuming and only one or two such phases where introduced at every node of the branch-and-cut tree. But, since all cuts found are globally valid, a pool of cuts was kept which could be checked by every node.

In this way the total elapsed time could be reduced by 50% compared with a sequential implementation using 3-dicycles only. Note that, when using 3-dicycles only, there is no reasonable use of parallelization (except for parallel processing of the tree), whereas it makes no sense to use small facet separation in a sequential code. The number of nodes in the branch-and-cut tree was significantly reduced. At the expense of more CPU time the number of nodes could be reduced further, but the above strategy seems to give the best trade-off between less CPU time or less number of nodes in the tree. The separation of small facets proved to be useful, and more research on their optimal employment is needed. However, it seems that, due to their small number of nonzero coefficients, they are not the key to achieving a breakthrough for solving really large problems.

6.7 Local Cuts and Target Cuts

There is a further general possibility for finding violated inequalities which is in particular easily applicable for the LOP. This approach is somehow related to the small facet cut generation as discussed above in that the original problem and the associated fractional solution are projected to a smaller problem, and cuts for the smaller

problem are lifted back to the orginal one. The difference is that not "template cuts" (i.e., inequalities which are basically known) are sought, but that general cuts can be derived. The approach for generating *local cuts* was invented and extensively used for solving large traveling salesman problems [5]. We only describe the principal idea.

Assume that the given problem is large and that for the current fractional solution x^* either no cut can be found, or searching for cuts would be much too time consuming because of the size of the problem.

The idea now is to go to a smaller problem instance by some way of "shrinking" the original feasible set P to a smaller set Q, at the same time transforming x^* to a point y^*. The goal is to do the shrinking such that, if x^* is infeasible for P, then also y^* is infeasible for Q. Now it is tried to find an inequality $b^T x \leq b_0$ separating y^* from Q which can basically be accomplished by solving the problem

$$\max b^T y^* - b_0$$
$$b^T y - b_0 \leq 0, \text{ for all } y \in Q.$$

In the case of violation (solutions with objective function value greater than 0 exist), this problem would be unbounded, so some normalization conditions are needed for b. E.g., this could be realized by bounding the components of b, but we do not elaborate on this here.

Because of the trivial lifting property of the linear ordering polytope we can assume that Q is defined by a subset of the nodes of the original problem (i.e., consists of the characteristic vectors of all linear orderings of the nodes of this subset) and that y^* is just x^* restricted to this subset.

The new problem can now be solved iteratively in the following way. We start with some linear orderings $\chi^1, \ldots, \chi^k \in Q$ giving a subset $\overline{Q} \subseteq Q$. Then the linear programming problem (augmented by normalization conditions)

$$\max b^T y^* - b_0$$
$$b^T \chi^i - b_0 \leq 0, \text{ for all } i = 1, \ldots, k$$

is solved. Let z^* be the optimum objective function value with corresponding inequality $b^{*T} x \leq b_0^*$. If $z^* \leq 0$ then no violated inequality can be found for Q. If $z > 0$ then it has to be checked whether the computed inequality is valid for Q or not. In the first steps one can try to find points in Q violating this inequality by running heuristics for finding linear orderings with high value w.r.t. the objective function b^*. If violating linear orderings are found then they can be added to $\overline{Q}$ and the above problem is solved again. If the heuristic does not find a violating ordering then the problem has to be solved exactly by optimizing over Q. This amounts to solving a small LOP. If the optimum objective function value is still positive, then a cutting plane for the small problem has been found, otherwise the shrinking procedure has failed to generate a cut.

If a cut has been found, it has to be lifted to a valid inequality $a^T x \leq a_0$ for the original problem. In the case of the LOP it can be taken as is and a cutting plane for the original problem can be added.

There is a slight addition to this procedure which is helpful in practical computations. Usually the cut is tight just at a single vertex of Q. In order to get a better inequality this cut has to be "tilted" to include further vertices until eventually a facet cut is constructed.

With the so-called *target cuts* the same intention as with local cuts is followed. We want to find cuts for large problems by finding cuts for a problem originating from the true one by some shrinking procedure.

Let P, x^*, Q and y^* be defined as above.

In a first step, all linear orderings $\chi^1, \ldots, \chi^m \in Q$ are enumerated. Let x^0 be an interior point of $\mathrm{conv}(\{\chi^1, \ldots, \chi^m\})$ (e.g., $x^0 = \frac{1}{m} \sum \chi^i$). Note that in our case $\mathrm{conv}(\{\chi^1, \ldots, \chi^m\}) = P_{LO}^k$, if Q is defined by k nodes.

The separation problem can be solved by determining the optimum solutions of the following pair of dual linear programs.

$$\min \sum_{i=1}^m \lambda_i$$

$$\sum_{i=1}^m \lambda_i (\chi_i - x^0) = x^* - x^0$$

$$\lambda_i \geq 0$$

$$\max (x^* - x^0)^T u$$

$$(\chi_1 - x^0, \ldots, \chi_m - x^0)^T u \leq 1$$

If $\sum \lambda_i^* > 1$ for the pair λ^* and u^* of optimum solutions then the violated inequality $u^{*T}(x - x^0) \leq 1$ is found. This inequality is facet defining for P_{LO}^k.

The main difference between the two approaches is that target cuts generate facets for the subproblem while local cuts generate facets only after tilting. But on the other hand, local cut generation can work with subsets of Q (*delayed column generation*), and a cut can be found without enumerating the set Q. Delayed column generation can also be adopted for target cuts, but this is more complicated (details can be found in [16]).

Chapter 7
Further Aspects

Abstract In this chapter we want to address some issues of interest for the LOP which we have not included in the previous chapters and point to some areas for possible further research.

7.1 Approximative Algorithms

It is surprising that, although the LOP is a classical difficult combinatorial optimization problem, not much is known about heuristics with approximation guarantees.

For a nonnegative objective function, a trivial heuristics guarantees an approximation of 50%. Namely, take a random ordering, then either the arc weights of the acyclic tournament induced by this ordering, or the arc weights of the reverse tournament sum to at least half of the total weight, and so $\frac{1}{2}$-approximation is trivial. In [104] it is proved that it is NP-hard to approximate the LOP with a factor better than $\frac{65}{66}$.

As a somewhat contrasting result, it was shown in [18] that the LOP is "asymptotically easy". To be more precise: under certain mild probability assumptions, the ratio between the objective function values of the best and of the worst solution is arbitrarily close to 1 with probability tending to 1 if the problem size goes to infinity. On the other hand, intuitively this result is not surprising. If all entries of the matrix are drawn from a uniform distribution then there is no specific structure and all permuted matrices look similar. And, since one half of the entries are added, the expected difference between values of feasible solutions should be small.

But, as we have seen in the preceding chapters, it is not very difficult to come up with solutions close to optimal for arbitrary instances using elaborate heuristics. This suggests that better approximation results should be possible to narrow the gap $\left[\frac{1}{2}, \frac{65}{66}\right]$ which is so far an open problem for possible polynomial time approximation.

There are some approximation results on special variants of the problem like the minimum feedback arc set or the acyclic subdigraph problem. For a survey of

approximation results and a discussion of further interesting aspects and open questions concerning tournaments see [31].

7.2 Integrality Gaps of LP Relaxations

The key to solving the LOP with branch-and-bound or branch-and-cut algorithms is the strength of the relaxations employed. There has been interesting research on the so-called *integrality gap* of LP relaxations. For some given nonnegative weight function c, let $c_{\text{opt}}(c)$ and $c_{\text{LP}}(c)$ be the corresponding values of the optimum linear ordering and of the LP relaxation, respectively. Then the integrality gap of the LOP w.r.t. this relaxation is defined as

$$\sup_{c>0} \frac{c_{\text{LP}}(c)}{c_{\text{opt}}(c)}.$$

In [104] it is shown that the gap of the standard LP relaxation with 3-dicycle inequalities is arbitrarily close to 2. The proof is based on the existence of a class of digraphs $D_\varepsilon = (V,A)$ with the property that any acyclic subset of the arcs does not contain more than approximately $(1+\varepsilon)\frac{|A|}{2}$ arcs.

But, also when further facet-defining inequalities are added, this gap is not improved too much. If all k-fence inequalities were added, the integrality gap would still be $2-\varepsilon$. The situation could be a little better when small Möbius ladder inequalities (on up to 7 nodes) are added. In this case, the integrality gap can be shown to be at least $\frac{33}{17} - \varepsilon$, but it could be strictly bounded away from 2.

Based on the approximation results of [104] cited above we can conclude that, unless P = NP, no LP relaxation can have an integrality gap less than $\frac{66}{65}$ if it is solvable in polynomial time.

The integrality gap of the 3-dicycle relaxation is closely related to the dicycle strength discussed in the previous chapter. Namely, if the objective function is given as the left hand side of a facet-defining inequality in normal form, then the integrality gap is exactly the dicycle strength of this facet.

7.3 Degree of Linearity

As a measure for the "triangularity" of an (n,n)-matrix $C = (c_{ij})$ the *degree of linearity* is defined as

$$\lambda(C) = \frac{\sum\limits_{\sigma(i)<\sigma(j)} c_{ij}}{\sum\limits_{i \neq j} c_{ij}}$$

for an optimum permutation σ. Consider the trivial relaxation of the LOP:

$$\max \sum_{(i,j)\in A_n} c_{ij}x_{ij}$$

$$x_{ij} + x_{ji} = 1, \text{ for } 1 \le i < j \le n,$$

$$x_{ij} \in \{0,1\}, \text{ for } 1 \le i, j \le n, i \ne j.$$

This relaxation is called the *tournament relaxation* because its feasible solutions are exactly the tournaments in A_n. It can be solved simply by setting, for $i < j$, $x_{ij} = 1$ and $x_{ji} = 0$, if $c_{ij} \ge c_{ji}$, and $x_{ij} = 0$ and $x_{ji} = 1$, otherwise. Therefore, the constraint "$x_{ij} \in \{0,1\}$" can be replaced by "$x_{ij} \ge 0$" and the relaxation can also be viewed as an LP relaxation.

For a problem in normal form, the optimum value is just the sum of all weights. If $c_T(c)$ denotes this optimum value and $c_{\text{opt}}(c)$ the maximum weight of an acyclic tournament, then

$$\frac{c_T(c)}{c_{\text{opt}}(c)} = \frac{1}{\lambda(C)},$$

i.e., the integrality gap of this relaxation is the inverse of the degree of linearity, and thus lies in the interval $[1,2]$. The results on the integrality gap of the 3-dicycle relaxation show that the worst case bound 2 is indeed tight.

In analogy with the 3-dicycles, we now have a connection with the trivial strength of facets. Namely, if the objective function is given as the left hand side of a facet-defining inequality in normal form, then the integrality gap of the tournament relaxation is exactly the trivial strength of this facet.

The integrality gaps of the tournament and of the 3-dicycle relaxation should be correlated. Figure 7.1 displays the corresponding integrality gaps for a set of random problems with $n = 44$.

Fig. 7.1 Correlation between 3-dicycle and tournament gap

There is a significant correlation but, e.g., for instances where the 3-dicycle gap is 1, the tournament gap varies between 1.232 and 1.314. Note that the gaps are far from their worst case bound.

We have performed some experiments for learning more about the distribution of the degree of linearity. To this end we have generated random problems in normal form with different densities and sizes of coefficients.

Tables 7.1–7.3, respectively, show minimum, maximum and average degrees of linearity for random problems with densities from 10% to 100% and integral coefficients drawn uniformly from the interval $[0, U]$, where $U = 1, 2, 20, 200, 2000$. For every combination 1000000 instances were solved.

Table 7.1 Miminum degree of linearity ($n = 14$, 1000000 problems)

	10%	20%	30%	40%	50%	60%	70%	80%	90%	100%
1	0.667	0.750	0.744	0.721	0.720	0.702	0.700	0.691	0.683	0.681
2	0.667	0.759	0.750	0.740	0.730	0.720	0.707	0.703	0.689	0.686
20	0.694	0.769	0.772	0.764	0.750	0.731	0.725	0.719	0.710	0.701
200	0.690	0.765	0.777	0.767	0.740	0.734	0.729	0.712	0.705	0.701
2000	0.679	0.765	0.776	0.758	0.750	0.730	0.723	0.720	0.713	0.701

Table 7.2 Maximum degree of linearity ($n = 14$, 1000000 problems)

	10%	20%	30%	40%	50%	60%	70%	80%	90%	100%
1	1.000	1.000	1.000	1.000	1.000	1.000	0.967	0.944	0.923	0.923
2	1.000	1.000	1.000	1.000	1.000	1.000	0.974	0.956	0.929	0.910
20	1.000	1.000	1.000	1.000	1.000	1.000	0.989	0.970	0.957	0.944
200	1.000	1.000	1.000	1.000	1.000	1.000	1.000	0.975	0.966	0.943
2000	1.000	1.000	1.000	1.000	1.000	0.999	0.987	0.974	0.947	0.931

Table 7.3 Average degree of linearity ($n = 14$, 1000000 problems)

	10%	20%	30%	40%	50%	60%	70%	80%	90%	100%
1	0.988	0.950	0.907	0.872	0.844	0.822	0.803	0.787	0.773	0.760
2	0.991	0.957	0.918	0.885	0.857	0.835	0.815	0.799	0.785	0.773
20	0.994	0.971	0.940	0.909	0.882	0.859	0.840	0.823	0.808	0.795
200	0.994	0.971	0.939	0.908	0.881	0.858	0.838	0.822	0.807	0.794
2000	0.994	0.971	0.938	0.908	0.881	0.858	0.838	0.821	0.807	0.794

We conducted further experiments with smaller and larger problems. On average, the following tendencies can be observed and are confirmed by the results shown here:

– Problems with smaller entries have lower degree of linearity.
– The minimum degrees of linearity can be found in very sparse problems with small entries.
– Problems with higher density have a lower degree of linearity.
– Problem examples giving values close to the worst case bound are extremely rare.

Since the degrees of linearity observed in the above experiments are far from the worst case 0.5, we tried to generate more instances in order to have a chance to get lower degrees. Table 7.4 shows some results suggesting again that worst case examples must be extremely rare. In addition, the table also displays the maximum integrality gap encountered for 3-dicycle relaxations. It should be difficult to generate small instances where the 3-dicycle bound exceeds the optimum by more than 5%.

The discussion of facet strengths in the preceding chapter showed that there are instances with $n = 8$ and trivial gap 1.538 (corresponding to the degree of linearity 0.650) and 3-dicycle gap 1.077. In our experiments we did not even come close to these values which are themselves still far away from the worst case. Probably such bad cases really only occur for large n and sparse digraphs (as supported by the theoretical results). Worst cases should be extremely rare and therefore it is no surprise that they did not come up in our experiments because we could generate only a very small fraction of possible problem instances.

Table 7.4 Computational results for random problems

n	Density	#problems	BB	min DoL	max DoL	3-cyc gap
10	1.0	10008189	0.1%	0.6889	1.0000	1.0323
20	0.1	8500349	0.6%	0.7647	1.0000	1.0357
20	0.5	327350	30.1%	0.7238	0.9333	1.0221
20	1.0	545636	63.1%	0.6737	0.8263	1.0365
30	0.1	1312510	26.4%	0.8269	1.0000	1.0217
30	0.2	88478	55.1%	0.8000	1.0000	1.0159
30	1.0	1974	99.8%	0.6713	0.7379	1.0471
35	1.0	118	81.4%	0.6756	0.7109	1.0446

Table 7.4 also gives the percentage of problems (BB) which required a branch-and-bound, i.e., for which the 3-dicycle relaxation did not give an integral solution. For larger problems with comparatively low degree of linearity a solution with the 3-dicycle LP only cannot be expected. The percentage of 81.4% for $n = 35$ is obviously too low, but caused by the fact that only very few instances could be solved.

7.4 Semidefinite Relaxations

Let $T_1, \ldots, T_m$, $m = n!$, denote all spanning acyclic tournaments of $D_n = (V_n, A_n)$ with their characteristic vectors χ^i. Throughout this monograph algorithms for solving the LOP to optimality were based on linear relaxations of

$$P_{LO}^n = \mathrm{conv}\big(\{\chi^i \mid i = 1, \ldots, m\}\big)$$

for computing upper bounds on the optimum objective function value, i.e., the bounds were obtained by solving linear programs of the form $\max\{c^T x \mid Ax \leq b\}$ where $P_{LO}^n \subseteq \{x \mid Ax \leq b\}$. We have experienced that the computation of these bounds is very time consuming and for difficult problems the upper bounds are fairly weak.

A more powerful alternative could be the employment of so-called *semidefinite relaxations* which require the solution of semidefinite programming problems. With symmetric (n,n)-matrices $C, A_1, \ldots, A_m$ and scalars $b_1, \ldots, b_m$ a semidefinite program (SDP) can be defined as

$$\min\big\{\langle C, X \rangle \mid \langle A_i, X \rangle = b_i, i = 1, \ldots, m, X \succeq 0\big\}.$$

Here $\langle C, X \rangle$ denotes the product $\sum_{i=1}^n \sum_{j=1}^n c_{ij} x_{ij}$ and $X \succeq 0$ denotes the requirement that X is a positive semidefinite matrix. Semidefinite programs can e.g. be solved by interior point methods. We cannot go into details on semidefinite programming here, but refer to the state-of-the-art survey [113].

There a several ways of coming up with a semidefinite relaxation for the LOP. A first one is the following where the central idea is that instead of P_{LO}^n we now consider

$$\mathscr{M}_{LO}^n = \mathrm{conv}\big(\{\chi^i (\chi^i)^T \mid i = 1, \ldots, m\}\big).$$

This set lies in the space of symmetric (n,n)-matrices and in addition every matrix in $\mathscr{M}_{LO}^n$ is positive semidefinite. The additional property that we are interested only in matrices of rank 1 cannot be handled by SDP solvers and hence it is relaxed. Of course, every constraint for the original problem can be transformed into a constraint in the new space. The big advantage is that also quadratic constraints now become linear and can thus be employed in models. A disadvantage of this approach is that the dimension of the problem space is squared, i.e., in the case of the LOP increases from $O(n^2)$ to $O(n^4)$.

When using SDP relaxations it is useful to go from the 0/1-model to a ± 1-model. To this end we first replace the usual characteristic vectors χ^i for tournaments by ± 1-vectors ξ^i by setting

$$\xi^i = 2\chi^i - 1, i = 1, \ldots, m.$$

Inequalities for the original characteristic vectors can easily be transformed into the new vectors. E.g., the 3-dicycle inequalities now read

$$-1 \leq x_{ij} + x_{jk} - x_{ik} \leq 1,$$

for all $i < j < k$. So essentially the model is not changed, but a major benefit is that now, in the semidefinite program, we can request in addition that the diagonal elements have value 1 as they correspond to squares of 1 or -1.

Instead of starting with a linear formulation of the LOP we could as well start with a quadratic formulation and then construct a semidefinite relaxation by going from vectors to matrices as described above. Such an approach was taken by Newman in [103]. It is based on a quadratic programming formulation of the LOP with ±1-variables y_{ik} defined as

$$y_{ik} = \begin{cases} -1, & \text{the position of } i \text{ in the ordering is less than } k, \\ +1, & \text{the position of } i \text{ in the ordering is at least } k. \end{cases}$$

The formulation of Newman is

$$\max \sum_{(i,j)\in A_n} \sum_{1\leq h \leq l \leq n} \frac{1}{4} c_{ij} (y_{ih} - y_{i,h-1})(y_{jl} - y_{j,l-1})$$

$$(y_{ih} - y_{i,h-1})(y_{jl} - y_{j,l-1}) \geq 0, \quad \text{for all } i,j \in V_n, h,l = 1,\ldots n,$$

$$y_{ih} y_{ih} = 1, \quad \text{for all } i \in V_n, h = 1,\ldots n,$$

$$y_{i0} = -1, \text{ for all } i \in V_n,$$

$$y_{in} = 1, \quad \text{for all } i \in V_n,$$

$$\sum_{i,j \in V_n} y_{i,\frac{n}{2}} y_{j,\frac{n}{2}} = 0,$$

$$y_{ih} \in \{-1,+1\}, \text{ for all } i,h = 1,\ldots,n.$$

The quality of the semidefinite relaxation based on this model is analyzed in [103]. Surprisingly, for the worst case examples of the 3-dicycle relaxation which give an integrality gap arbitrarily close to 2, the gap of this relaxation is only at most 1.64.

7.5 Context Independent Solvers

General purpose heuristics are based on models that treat the objective function evaluation as a black box, making the search algorithm context independent. The evaluation of the objective function can be seen as the call of an oracle which returns the objective value of a feasible solution. The solution algorithm itself does not know anything about the structure of this objective function and therefore cannot exploit structural properties. In particular, it is unknown to the solver whether the objective function is linear or nonlinear and therefore it possibly cannot apply the most effective search strategies for a given situation.

Meta-heuristics can be used to create solution procedures that are context independent. The original genetic algorithmic designs were based on this model. The

advantage of this design is that the same solver can be applied to a wide variety of problems. The obvious disadvantage is that the solutions found by context independent solvers might be inferior to those of specialized procedures when both are allotted the same amount of computer effort (e.g. total search time).

Context independent solvers (also called general purpose or black box optimizers) based on meta-heuristics have found their home in commercial implementations. A standard evolutionary solver that is a context independent GA implementation is included in the Premium Solver Platform of Frontline Systems, Inc (`www.frontsys.com`). Opttek Systems, Inc (`www.opttek.com`) commercializes OptQuest, a context independent solver based on scatter search. Other GA-based commercial implementations of general purpose optimizers are Evolver by Palisade Corporation (`www.palisade.com`) and Pointer by Synaps, Inc. (`www.synaps-inc.de`).

In [21] a hybrid meta-heuristic for a class of problems based on a context independent paradigm is proposed. The method is restricted to those problems whose solutions are represented by permutations. This class includes a wide range of problems such as the traveling salesman problem, the quadratic assignment problem, various single machine sequencing problems, and the linear ordering problem, to mention only a few. The procedure is a combination of scatter search and of tabu search. The scatter search framework provides a means for diversifying the search throughout the exploration of the permutation solution space. Two improvement methods are used to intensify the search in promising regions of the solution space. Improved solutions are then used for combination purposes within the scatter search design. We briefly describe this design in the next paragraphs.

As described in Chapter 3, there are three elements that we need to define in any evolutionary method: a way to generate solutions, a way to combine solutions and a way to maintain a set (population) of solutions. The procedure in [21] follows the standard scatter search design to maintain the set of solutions (reference set). A generator of solutions (permutations), which focuses on diversification and not on the quality of the resulting solutions, is used at the beginning of the search to build the initial set P of solutions. The generator, proposed by Glover [55], uses a systematic approach to creating a diverse set of permutations. This contrasts with the typical GA approach of randomly generating an initial set of solutions from which to start the evolutionary search. In order to obtain a set of solutions of reasonable quality and diversity, an improvement method is applied to the solutions in P. The improvement method consists of two phases, a simple local search based on exchange moves and a tabu search. The TS is based on a short-term memory function and is applied only to the most promising solutions.

In order to design a context independent combination methodology that performs well across a wide collection of different problems, a set of ten combination methods is proposed (cm_1 to cm_{10}), from which one is probabilistically selected according to its performance in previous iterations. Solutions in the reference set are typically ordered according to their objective function value. So, the best solution is the first one in this set and the worst is the last one. A score is updated for every method as follows. If a solution obtained with combination method cm_i qualifies to be the

jth member of the current reference set, then $b - j + 1$ is added to the score of cm_i. Therefore, combination methods that generate good solutions accumulate higher scores and are used more often.

In the experiments of [21] the authors considered four well known problems in which solutions are represented by permutations. We report in Table 7.5 the comparison between the proposed solver, called SS-TS, and two commercial packages, Frontline and Opttek, when solving the linear ordering and the traveling salesman problem. Table 7.5 reports the average percentaged deviation between the best solution obtained with each method and the optimal solution of 49 input-output instances of LOLIB (IO set) and 31 instances of TSPLIB (a library for the traveling salesman problem [123]). The table also shows the average CPU time in seconds.

Table 7.5 Comparison with commercial solvers

	OptQuest	Frontline	SS-TS
Linear Ordering Problem			
Deviation	8.5%	16.1%	0.0%
CPU time	301	300	25
Traveling Salesman Problem			
Deviation	311.4%	8.4%	5.7%
CPU time	5772	5628	23

Table 7.5 shows that the proposed scatter search with a tabu search improvement method yields higher quality solutions on average when compared to two commercially available software packages. To make a fair comparison, a fixed number of objective function evaluations has been set as a termination criterion for all procedures. We include the execution time to show the advantage of using a specialized code that does not include additional costly routines, such as those associated with graphical output or databases to store all visited solutions. It must be noted that, although we mentioned in previous chapters that the input-output instances of LOLIB are relatively easy to solve with the meta-heuristic methodologies, the solvers in Table 7.5 are context independent and they do not employ the knowledge, properties and structure of the problem as meta-heuristics can do.

7.6 Difficulty of LOP Instances

A problem instance is not in itself difficult, but it is difficult with respect to the solution algorithm. Clearly, for branch-and-bound or branch-and-cut, a problem should be difficult if the relaxations used do not provide good bounds.

Figures 7.2 and 7.3 display the correlation between CPU times for solving the problems of the previous section to optimality and the respective integrality gaps of the tournament and the 3-dicycle LP. Obviously, problems with higher gaps are more difficult for branch-and-cut algorithms.

Fig. 7.2 Correlation between tournament gap and CPU time

Fig. 7.3 Correlation between 3-dicycle gap and CPU time

From the perspective of heuristic and meta-heuristic methodologies we can differentiate three categories of methods:

a) constructive methods,
b) local search based methods,
c) population based methods.

An instance is easy to solve with a constructive method if it contains useful information associated with partial neighborhoods. In other words, if we have useful evaluations both for selecting a good candidate and for placing it in a good position to extend the partial solution under construction. In this case we say that the constructive process is guided by context information.

An instance is easy to solve with a local search based method if it contains useful information associated with complete neighborhoods. If the evaluation associated with the moves in a neighborhood permits us to discriminate among them, the local search can progress in the solution space. Min-max and max-min problems typically present flat landscapes and do not contain information at all associated with the neighborhoods. Move values are usually 0 in these problems, and therefore decisions to select moves must be taken at random.

An instance is easy to solve with a population based procedure if its combination method is able to obtain good solutions when applied to good solutions. The conjecture that information about the relative desirability of alternative solutions is captured in different forms by different solutions motivated these approaches. Their success on a particular instance is based on its ability to capture the structures in good solutions responsible for their quality, and to transfer these structures to the combined solutions.

7.7 Sparse Problems

When most of the arcs of a LOP have zero weight, then it is preferrable to treat it as an acyclic subdigraph problem. As pointed out in the introduction ASP and LOP are trivially equivalent. If the set of arcs with positive weights is sparse, then a maximum weighted acyclic subdigraph arc set will not be a tournament, but can be extended to a tournament.

In the sparse case, we do not optimize the objective function over P_{LO}^n but over the polytope $P_{AC}(D)$ instead, where $D = (V,A)$ is the subdigraph of D_n obtained by eliminating all arcs with weight zero. $P_{AC}(D)$ is defined as

$$P_{AC}(D) = \text{conv}\{\chi^B \in \{0,1\}^A \mid B \text{ is an acyclic arc set in } A\}.$$

Further details on the ASP, in particular on the facet structure of $P_{AC}(D)$ can be found in [72]. Note that P_{LO}^n is a face of $P_{AC}(D_n)$ obtained by requiring equality in all 2-dicycle inequalities $x_{ij} + x_{ji} \leq 1$. Therefore, the facial structures of the two polytopes are closely related.

The canonical IP formulation of the ASP is

$$\max \sum_{(i,j)\in A} c_{ij} x_{ij}$$
$$x(C) \leq |C| - 1, \text{ for all dicycles } C \text{ in } A,$$
$$x_{ij} \in \{0,1\}, \text{ for } (i,j) \in A.$$

Note that in the case of the ASP we have to exclude all dicycles of length k, $k \geq 2$, because the digraph is not complete. Furthermore the tournament equations do not apply here. (For LOP instances in normal form, there are no 2-dicycles.)

If the integrality constraints are replaced by "$0 \leq x_{ij} \leq 1$", then we obtain the *dicycle relaxation* of the ASP. The separation problem for dicycle inequalities can easily be solved in polynomial time, as the following observation shows. Let x^* be the current fractional LP solution and define $y^* = 1 - x^*$. Then $x^*(C) \leq |C| - 1$ if and only if $y^*(C) \geq 1$. So, for separating dicycle inequalites, the shortest dicycle in A with arc weights given by y^* is computed (using shortest path techniques). If the shortest such diycle has length greater than or equal 1, then no dicycle inequality is violated, otherwise this dicycle yields a cut.

Therefore, the dicycle relaxation can be solved in polynomial time and can serve as a basis for branch-and-bound and branch-and-cut algorithms which are designed analogously as for the LOP.

Note that if we add the inequality system

$$x_e - x_f + x_g \leq 1, \text{ for all arcs } e, f, g \in A \text{ and } V(f) \subseteq V(\{e,g\}),$$

then we obtain an IP formulation of the node induced acyclic subdigraph problem which can serve as a basis for branch-and-cut algorithms.

7.8 A Simple Dual Heuristic

Let a LOP in normal form be given. Denote its optimum solution by c_{opt} and let Q denote the sum of all arc weights.

Let B be some dicycle in A_n and define

$$B_{\min} = \min\{c_{ij} \mid (i,j) \in B\}.$$

Then $B_{\min} \geq 0$ and obviously $c_{\mathrm{opt}} \leq Q - B_{\min}$. We modify the objective function by setting

$$c'_{ij} = \begin{cases} c_{ij} - B_{\min}, & (i,j) \in B, \\ c_{ij}, & \text{otherwise.} \end{cases}$$

With another dicycle B' and $B'_{\min} = \min\{c'_{ij} \mid (i,j) \in B'\}$ we can improve the upper bound to $c_{\mathrm{opt}} \leq Q - B_{\min} - B'_{\min}$.

Of course, only dicycles with strictly positive minimum arc weight are of interest. We can iterate this procedure as long as we find such dicycles and get an upper bound for the optimum solution of the LOP.

DualHeuristic(C)

(1) Set $UB = Q$ and $c'_{ij} = c_{ij}$, for all $(i,j) \in A_n$.
(2) While $A_n \setminus \{(i,j) \mid c'_{ij} = 0\}$ is not acyclic:

    (2.1) Let B be a dicycle in $A_n \setminus \{(i,j) \mid c'_{ij} = 0\}$.
    (2.2) Set $B_{\min} = \min\{c'_{ij} \mid (i,j) \in B\}$.
    (2.3) Set $UB = UB - B_{\min}$.
    (2.4) For every $(i,j) \in B$ set $c'_{ij} = c'_{ij} - B_{\min}$.

Surprisingly, this heuristic yields fairly reasonable upper bounds. Table 7.6 gives the average percentage of this bound w.r.t the best known lower and upper bounds. We also list the percentage of this bound in terms of the sum of all objective function coefficients.

Table 7.6 Comparison of simple upper bounds

Problem class	% Sum	%best UB	%best LB
IO	96.01	100.46	100.46
Random A1	70.38	101.62	114.77
Random A2	94.46	101.29	101.55
Random B	76.48	105.07	105.73
MB	91.11	100.81	100.81
XLOLIB	85.16	106.78	109.66
SGB	73.16	100.89	100.89
Spec	85.59	104.57	106.90

Thus, in the case that LP bounds cannot be computed, these bounds provide useful information for assessing heuristics.

Now consider the following linear program.

$$\max \sum_{(i,j) \in A_n} c_{ij} x_{ij}$$
$$x(C) \le |C| - 1, \text{ for all dicycles } C \text{ in } A_n, |C| \ge 2,$$
$$x_{ij} \le 1, \text{ for } (i,j) \in A_n,$$
$$x_{ij} \ge 0, \text{ for } (i,j) \in A_n.$$

For a LOP in normal form, its optimum value is equal to the value of the 3-dicycle relaxation. The dual of this program is

$$\min \sum_{B \text{ dicycle in } A_n} (|B| - 1)y_B + \sum_{(i,j) \in A_n} z_{ij}$$

$$\sum_{\substack{B \text{ dicycle in } A_n \\ (i,j) \in B}} y_B + z_{ij} \geq c_{ij}, \text{ for all } (i,j) \in A_n,$$

$$y_B \geq 0, \text{ for all dicycles } B \text{ in } A_n,$$

$$z_{ij} \geq 0, \text{ for all } (i,j) \in A_n.$$

Let $\mathscr{B}$ denote the set of dicycles used in the dual heuristic and set the dicycle variables

$$y_B = \begin{cases} B_{\min}, & B \in \mathscr{B}, \\ 0, & \text{otherwise}, \end{cases}$$

and

$$z_{ij} = c_{ij} - \sum_{B \in \mathscr{B},\, (i,j) \in B} y_B.$$

Then this setting of variables is feasible for the dual program and

$$Q - \sum_{B \in \mathscr{B}} B_{\min} = \sum_{(i,j) \in A_n} \left(z_{ij} + \sum_{B \in \mathscr{B},\, (i,j) \in B} y_B \right) - \sum_{B \in \mathscr{B}} B_{\min}$$

$$= \sum_{(i,j) \in A_n} z_{ij} + \sum_{B \in \mathscr{B}} (|B| - 1)y_B.$$

So the heuristic can be interpreted as finding a feasible solution for the dual of the 3-dicycle relaxation. Therefore it gives an upper bound on the value of the 3-dicycle relaxation and thus an upper bound on the LOP.

7.9 Future Research

The current state-of-the-art of solution algorithms for the LOP seems not really satisfactory. Whereas very good feasible solutions can be found for large problems using the comprehensive toolbox of heuristics, fairly small problems can still be difficult for exact algorithms and cannot be solved to optimality in reasonable time. So there is definite need for further research on algorithms, mainly for exact algorithms but also for heuristics. We expect progress on the following topics.

Black Box Solvers

As we have seen in the computational experiments shown in Section 7.5, the quality of the solutions obtained with context independent or black box solvers is still moderate. Although it is expected that general purpose methods obtain medium quality

solutions, as compared with the high-quality obtained with the specialized methods, the truth is that these generic methods need further development. Considering that most commercial solvers are based on this technology, we should study their associated models and heuristics algorithms to achieve better results.

Solution of the 3-Dicycle Relaxation

The determination of this bound is at the core of exact algorithms. The solution time needed is still not satisfactory. Possibly a combination of primal and dual methods, i.e., cutting plane and bundle approaches bears the potential for improvement.

Separation Algorithms

Besides the trivial 3-dicycle enumeration and the separation of certain Möbius ladders as mod-2 inequalities, there are no other exact separation algorithms. It is known that the separation of k-fences is NP-hard. As the computation for small polytopes shows there is a wealth of further facet defining inequalities the structure of which is unexplored. So there is room for theoretical research on the complexity of separation and the development of further exact algorithms.

Local Cuts and Target Cuts

Local cuts and target cuts offer the possibility of generating cutting planes for large problems and generating cuts which are different from the ones used in separation procedures so far. It could be promising to study their effect.

Branch-and-Bound Revisited

Most of the present day branching schemes are binary, i.e., some variable is fixed to 0 or to 1 and the corresponding two subproblems are created. The effect of just changing one variable is small and it could be advantageous to employ more complicated schemes.

An obvious possibility is to branch by fixing nodes at certain positions of the ordering. Fixing a node at a position has significantly more consequences than just fixing one variable. Furthermore, branching on inequalities could be an option and parallelization is always applicable to branch-and-bound.

Different Models

All computational methods for solving the LOP are based on the binary variables x_{ij} where

$$x_{ij} = \begin{cases} 1, & \text{if } i \text{ is before } j \text{ in the ordering,} \\ 0, & \text{otherwise.} \end{cases}$$

But other modeling approaches could be explored as well. E.g., one could use binary variables x_{ijk} for three nodes i, j, k defined as

$$x_{ijk} = \begin{cases} 1, & \text{if } j \text{ is after } i \text{ and before } k \text{ in the ordering,} \\ 0, & \text{otherwise.} \end{cases}$$

An IP formulation of the LOP with these variables is

$$\max \sum w_{ijk} x_{ijk}$$
$$x_{ijk} + x_{ikj} + x_{jik} + x_{jki} + x_{kij} + x_{kji} = 1$$
$$x_{ijk} + x_{ikj} + x_{kij} - x_{ijl} - x_{ilj} - x_{lij} = 0$$
$$x_{ijk} \in \{0, 1\}$$

where

$$w_{ijk} = \frac{c_{ij} + c_{ik} + c_{kj}}{n - 2},$$

if c denotes the original objective function. The canonical LP relaxation gives the 3-dicycle bound as the standard formulation with 2-index variables.

For $n = 4$ the convex hull of feasible 0/1 vectors has further facets defined by the inequalities

$$x_{ijk} - x_{ijl} - x_{ljk} \leq 0$$
$$x_{ijk} - x_{ijl} - x_{lik} - x_{ilk} \leq 0$$
$$x_{ijk} - x_{ikl} - x_{ljk} - x_{ilk} \leq 0$$
$$x_{ijk} + x_{lkj} - x_{ikl} - x_{ilk} - x_{lji} - x_{lij} \leq 0$$

for pairwise distinct i, j, k, l.

If all these inequalities are added, then the LP relaxation for $n = 6$ is integral, i.e., all 3-fence and Möbius ladder inequalities are implied.

The exploration of alternative models could possibly lead to more effective algorithms.

General Integer Programming

Since we have difficult problem instances already for small values of n, the study of general integer programming approaches like lift-and-project could be worthwhile.

Determination of Optimum Solutions Using SDP

In theory, semidefinite programming bounds are stronger than linear programming bounds. First experiments with quadratic models and semidefinite relaxations support this. On the other hand, the problem dimension is drastically increased. But, in the range of a few hundred nodes, this could be manageable and there should be a thorough investigation of the power of semidefinite relaxations in the context of the LOP. First experiments on small problems reveal that the 3-dicycle bound seems to be easily improvable by employing SDP.

References

1. ACHATZ, H., KLEINSCHMIDT P. AND LAMBSDORFF, J. *Der Corruption Perceptions Index und das Linear Ordering Problem*, ORNews **26** (2006), 10–12.
2. ACHTERBERG, T. *SCIP: Solving constraint integer programs*, Math. Programming Computation **1** (2009), 1–41.
3. ADOLPHSON, D. AND HU, T.C.: *Optimal linear ordering*, SIAM J. on Applied Mathematics **25** (1973), 403–423.
4. AARTS, E. AND LENSTRA, J.K. *Local search in combinatorial optimization*, Princenton University Press, 2003.
5. APPLEGATE, D.L., BIXBY, E.B., CHVÁTAL, V. AND COOK, W.J.: *The traveling salesman problem: A computational study*, Princeton University Press, 2006.
6. AUJAC, H.: *La hiérarchie des industries dans un tableau des echanges industriels*, Rev. Economique **2** (1960).
7. BALUJA, S.: *An empirical comparison of 7 iterative evolutionary function optimization heuristics*, School of computer science, Carnegie Mellon University, 1995.
8. BECKER, O.: *Das Helmstädtersche Reihenfolgeproblem – die Effizienz verschiedener Näherungsverfahren*, in: *Computer uses in the social sciences*, Bericht einer Working Conference des Inst. f. höh. Studien u. wiss. Forsch., Wien, 1967.
9. BÉKÉSI, J., GALAMBOS, G., OSWALD, M. AND REINELT, G.: *Comparison of approaches for solving coupled task problems*, Technical Report, U Heidelberg, 2008.
10. BERTACCO, L., BRUNETTA, L. AND FISCHETTI, M.: *The linear ordering problem with cumulative costs*, European J. of Operational Research **189** (2008), 1345–1357.
11. BHANDARKAR, S.M. AND CHIRRAVURI, S.: *A study of massively parallel simulated annealing algorithms for chromosome reconstruction via clone ordering*, Parallel Algorithms and Applications **9** (1996), 67–89.
12. BOENCHENDORF, K.: *Reihenfolgenprobleme / Mean-flow-time sequencing*, Mathematical Systems in Economics 74, Verlagsgruppe Athenäum, Hain, Scriptor, 1982.
13. BOESE, K.D., KAHNG, A.B. AND MUDDU S.: *A new adaptive multi-start technique for combinatorial global optimisation*, Operations Research Letters **16** (1994), 103–113.
14. BOLOTASHVILI, G., KOVALEV, M. AND GIRLICH, E.: *New facets of the linear ordering polytope*, Siam J. Discrete Math. **12** (1999), 326–336.
15. BOWMAN, V.J.: *Permutation polyhedra*, SIAM J. Appl. Math. **22** (1972), 580–589.
16. BUCHHEIM, CH., LIERS, F. AND OSWALD, M.: *Local cuts revisited*, Operations Research Letters **36** (2008), 430–433.
17. BUCHHEIM, CH., WIEGELE, A. AND ZHENG, K.: *Exact algorithms for the quadratic linear ordering problem*, INFORMS J. on Computing **22** (2010), 168–177.
18. BURKARD, R.E. AND FINCKE, U.: *Probabilistic asymptotic properties of some combinatorial optimization problems*, Discrete Appl. Math. **12** (1985), 21–29.

19. CAMPOS, V., GLOVER, F. LAGUNA, M. AND MARTÍ, R.: *An experimental evaluation of a scatter search for the linear ordering problem*, J. of Global Optimization **21** (2001), 397–414.

20. CAMPOS, V., LAGUNA, M. AND MARTÍ, R.: *Scatter search for the linear ordering problem*, in: Corne, D., Dorigo, M. and Glover, F. (eds.), *New Ideas in Optimisation*, McGraw-Hill, 1999, 331–341.

21. CAMPOS, V., LAGUNA, M. AND MARTÍ, R.: *Context-independent scatter and tabu search for permutation problems*, INFORMS J. on Computing **17** (2005), 111–122.

22. DE CANI, J.S.: *Maximum likelihood paired comparison ranking by linear programming*, Biometrika **56** (1969), 537–545.

23. DE CANI, J.S.: *A branch & bound algorithm for maximum likelihood paired comparison ranking*, Biometrika **59** (1972), 131–135.

24. CAPRARA, A. AND FISCHETTI, M.: $\{0, \frac{1}{2}\}$-*Chvátal-Gomory cuts*, Math. Programming **74** (1996), 221–235.

25. CAPRARA, A., FISCHETTI, M. AND LETCHFORD, A.N.: *On the separation of maximally violated mod-k cuts*, Math. Programming **87** (2000), 37–56.

26. CHANAS, S. AND KOBYLANSKI, P.: *A new heuristic algorithm solving the linear ordering problem*, Computational Optimization and Applications **6** (1996), 191–205.

27. CHARON, I. AND HUDRY, O.: *Lamarckian genetic algorithms applied to the aggregation of preferences*, Annals of Operations Research **80** (1998), 281–297.

28. CHARON, I. AND HUDRY, O.: *The noising methods: A generalization of some metaheuristics*, European J. of Operational Research **135** (2001), 86 101.

29. CHARON, I. AND HUDRY, O.: *A branch-and-bound algorithm to solve the linear ordering problem for weighted tournaments*, Discrete Applied Mathematics **156** (2006), 2097–2116. `http://www.enst.fr/˜charon/tournament/median.html`

30. CHARON, I. AND HUDRY, O.: *A survey on the linear ordering problem for weighted or unweighted tournaments*, 4OR **5** (2007), 5–60.

31. CHARON, I. AND HUDRY, O.: *An updated survey on the linear ordering problem for weighted or unweighted tournaments*, Annals of Operations Research **175** (2009), 107–158.

32. CHENERY, H.B. AND WATANABE, T.: *International comparisons of the structure of production*, Econometrica **26** (1958), 487–521.

33. CHIARINI, B., CHAOVALITWONGSE, W. AND PARDALOS, P.M.: *A new algorithm for the triangulation of input-output tables*, in: Pardalos, P.M., Migdala, A. and Baourakis, G. (eds.), *Supply chain and finance*, World Scientific, 2004, 254–273.

34. CHRISTOF, T.: *Low-dimensional 0/1-polytopes and branch-and-cut in combinatorial optimization*, Shaker, 1997.

35. CHRISTOF, T. AND REINELT, G.: *Combinatorial optimization and small polytopes*, Top **4** (1996), 1–64.

36. COELLO, C. A., VAN VELDHUIZEN, D. A. AND LAMONT, G. B.: *Evolutionary algorithms for solving multi-objective problems*, Kluwer Academic / Plenum Publishers, 2002.

37. CONDORCET, M.J.A.N.: *Essai sur l'application de l'analyse à la probabilité des décisions rendues à la pluralité des voix*, Paris, 1785.

38. DAKIN, R.J. *A tree search algorithm for mixed integer programming problems*, Computer Journal **8** (1965), 250–255.

39. DAVIS, L. *Handbook of genetic algorithms*, International Thomson Computer Press, 1996.

40. DOIGNON, J.-P., FIORINI, S. AND JORET, G.: *Facets of the linear ordering polytope: A unification for the fence family through weighted graphs*, J. of Math. Psychology **50** (2006), 251–262.

41. DREO, J., PETROWSKI, A., SIARRY, P. AND TAILLARD, E.: *Metaheuristics for hard optimization*, Springer, 2006.

42. DRIDI, T.: *Sur les distribution binaires associées à des distributions ordinales*, Math. Sci. hum. **69** (1980), 15–31.

43. DUARTE, A., LAGUNA, M. AND MARTÍ, R.: *Tabu search for the linear ordering problem with cumulative costs*, Computational Optimization and Applications (2009), doi: 10.1007/s10589-009-9270-5.

44. FEO, T. AND RESENDE, M.G.C.: *A probabilistic heuristic for a computationally difficult set covering problem*, Operations Research Letters **8** (1989), 67–71.

45. FEO, T. AND RESENDE, M.G.C.: *Greedy randomized adaptive search procedures*, J. of Global Optimization **2** (1995), 1–27.

46. FIORINI, S.: $\{0, 1/2\}$-*cuts and the linear ordering problem: Surfaces that define facets*, SIAM J. on Discrete Mathematics **20** (2006), 893–912.

47. FISHBURN, P.C.: *Binary probabilities induced by rankings*, SIAM J. Disc. Math. **3** (1990), 478–488.

48. FISHBURN, P.C.: *Induced binary probabilities and the linear ordering polytope: a status report*, Math. Soc. Sciences **23** (1992), 67–80.

49. GARCIA, C.G., PÉREZ-BRITO, D., CAMPOS, V. AND MARTÍ, R.: *Variable neighborhood search for the linear ordering problem*, Computers and Operations Research **33** (2006), 3549–3565.

50. GAREY, M.R. AND JOHNSON, D.S.: *Computers and intractability: A guide to the theory of NP-completeness*, Freeman, 1979.

51. GILBOA, I.: *A necessary but insufficient condition for the stochastic binary choice problem*, J. of Math. Psychology **34** (1990), 371–392.

52. GLOVER, F. AND LAGUNA, M.: *Tabu search*, Kluwer Academic Publishers, 1997.

53. GLOVER, F.: *Heuristics for integer programming using surrogate constraints*, Decision Sciences **8** (1977), 371–392.

54. GLOVER, F.: *Future paths for integer programming and links to artificial intelligence*, Computers and Operations Research **13** (1986), 533–549.

55. GLOVER, F.: *A template for scatter search and path relinking*, in: Hao, J.-K., Lutton, E., Ronald, E., Schoenauer, M. and Snyers, D. (eds.), *Artificial Evolution*, LNCS 1363, Springer, 1998, 13–54.

56. GLOVER, F., KLASTORIN, T., AND KLINGMAN, D.: *Optimal weighted ancestry relationships*, Management Science **20** (1974), 1190–1193.

57. GOEMANS, M.X.: *Worst-case comparision of valid inequalities for the TSP*, Mathematical Programming **69** (1995), 335–349.

58. GOEMANS, M.X. AND HALL, L.A.: *The strongest facets of the acyclic subgraph polytope are unknown*, in: *Proc. of the 5th Int. IPCO Conference*, LNCS 1084, Springer, 1996, 415–429.

59. GRÖTSCHEL, M., JÜNGER, M. AND REINELT, G.: *A cutting plane algorithm for the linear ordering problem*, Operations Research **32** (1984), 1195–1220.

60. GRÖTSCHEL, M., JÜNGER, M. AND REINELT, G.: *Facets of the linear ordering polytope*, Math. Programming **33** (1985), 43–60.

61. GRÖTSCHEL, M., LOVASZ, L. AND SCHRIJVER, A.: *Geometric algorithms and combinatorial optimization*, Springer, 1988.

62. GRUNDEL, D.A. AND JEFFCOAT, D.E.: *Formulation and solution of the target visitation problem*, in: *Proc. of the AIAA First Intelligent System Technical Conference*, Chicago, 2004.

63. HANSEN, P. AND MLADENOVIC, N.: *Variable neighborhood search*, in: Glover, F. and Kochenberger, G. (eds.) *Handbook of Metaheuristics* (2003), 145–184.

64. HELLMICH, K.: *Ökonomische Triangulierung*, Rechenzentrum Graz, Heft 54, 1970.

65. HELMSTÄDTER, E.: *Die Dreiecksform der Input-Output-Matrix und ihre möglichen Wandlungen im Wachstumsprozeß*, in: Neumark, F. (ed.): *Strukturwandlungen einer wachsenden Wirtschaft* Schriften des Vereins für Socialpolitik, NF Band 30/II (1964), 1059–1063.

66. HERTZ,A. AND WIDMER, M.: *Guidelines for the use of meta-heuristics in combinatorial optimization*, European J. of Operational Research **151** (2003), 247–252.

67. HICKERNELL, F.J. AND YUAN, Y.: *A simple multistart algorithm for global optimization*, OR Transactions **1(2)** (1997).

68. HOLLAND, J.H.: *Adaptation in natural and artificial systems*, University of Michigan Press, 1975.

69. HOLUB, H.-W. AND SCHNABL, H.: *Input-Output-Rechnung: Input-Output-Tabellen*, Oldenbourg, 1982.

70. Huang, G. and Lim, A.: *Designing a hybrid genetic algorithm for the linear ordering problem*, in: Cantu-Paz, E. et al. (eds.): *Proc. of Genetic and Evolutionary Computation – GECCO 2003*, LNCS 2723, Springer, 2003, 1053–1064.

71. Johnson, D., Aragon, C.R., McGeoch, L.A. and Schevon, C.: *Optimization by simulated annealing: An experimental evaluation; Part I, graph partitioning*, Operations Research **37** (1989), 865–892.

72. Jünger, M.: *Polyhedral combinatorics and the acyclic subdigraph problem*, Heldermann, Berlin, 1985.

73. Jünger, M. and Mutzel, P.: *2-layer straightline crossing minimization: Performance of exact and heuristic algorithms*, J. of Graph Algorithms and Applications **1** (1997), 1–25.

74. Jünger, M. and Thienel, S.: *The ABACUS System for branch-and-cut-and-price algorithms in integer programming and combinatorial optimization*, Software: Practice and Experience **30** (2000), 1325–1352.
   http://www.informatik.uni-koeln.de/abacus

75. Kaas, R.: *A branch and bound algorithm for the acyclic subgraph problem*, European J. of Operational Research **8** (1981), 355–362.

76. Kang, S.: *Linear ordering and application to placement*, in: *Proc. of the 20th Annual ACM IEEE Design Automation Conference*, Miami Beach, Floria, USA (1983), 457–464.

77. Kemeny, J.G.: *Mathematics without numbers*, Daedalus **88** (1959), 577–591.

78. Kernighan, B.W. and Lin, S.: *An efficient heuristic procedure for partitioning graphs*, Bell Systems Technical Journal **49** (1979) 291–308.

79. Kirkpatrick, S., Gelatt C.D. and Vecchi, M.P.: *Optimization by simulated annealing*, Science **222** (1983), 671–680.

80. Knuth, D.E.: *The Stanford GraphBase: A platform for combinatorial computing*, Addison-Wesley, 1993.

81. Koppen, M.: *Random utility representation of binary choice probabilities: critical graphs yielding critical necessary conditions*, J. of Math. Psychology **39** (1995), 21–39.

82. Korte, B. and Oberhofer, W.: *Zwei Algorithmen zur Lösung eines komplexen Reihenfolgeproblems*, Unternehmensforschung **12** (1968), 217–231.

83. Korte, B. and Oberhofer, W.: *Zur Triangulation von Input-Output-Matrizen*, Jahrbücher für Nationalökonomie und Statistik **182** (1969), 398–433.

84. Korte, B. and Vygen, J. : *Combinatorial optimization. Theory and algorithms*, Springer-Verlag, Berlin 2002.

85. Laguna, M. and Martí R.: *Scatter search: Methodology and implementations in C*, Kluwer Academic Publishers, 2003.

86. Laguna, M., Martí R. and Campos, V.: *Intensification and diversification with elite tabu search solutions for the linear ordering problem*, Computers and Operations Research **26** (1999), 1217–1230.

87. Lenstra, H.W. Jr.: *The acyclic subgraph problem*, Report BW26, Mathematisch Centrum, Amsterdam, 1973.

88. Leontief, W.: *Quantitative input-output relations in the economic system of the United States*, The Review of Economics and Statistics **18** (1936).

89. Leontief, W.: *Input-output economics*, Oxford University Press, 1966.

90. Leung, J. and Lee, J.: *Reinforcing old fences gives new facets*, Tech. Rep 90-22, Dptm. of Operations Research, Yale University, 1990.

91. Leung, J. and Lee, J.: *More facets from fences for linear ordering and acyclic subgraph polytopes*, Discrete Applied Mathematics **50** (1994), 185–200.

92. LOLIB: *A library of benchmark instances for linear ordering problem*, www.ifi.uni-heidelberg.de/groups/comopt/software/LOLIB.

93. Martí, R., Reinelt, G. and Duarte, A.: *LOP instances and best known solutions*, http://heur.uv.es/optisom/LOLIB, 2010.

94. Martí, R., Laguna, M., Glover, F. and Campos, V.: *Reducing the bandwidth of a sparse matrix with tabu search*, European J. of Operational Research **135** (2001), 450–459.

95. Martí, R.: *Multi-start methods*, in: Glover, F. and Kochenberger, G. (eds.) *Handbook of Metaheuristics*, Kluwer Academic Publishers, 2003, 355–368.

96. MAYNE, D.Q. AND MEEWELLA, C.C.: *A non-clustering multistart algorithm for global optimization*; in: Bensoussan, A. and Lions, J.L. (eds.) *Analysis and Optimization of Systems*, Lecture Notes in control and information sciences 111, Springer, 1988, 334–345.

97. METROPOLIS, N., ROSENBLUTH, A. ROSENBLUTH, M., TELLER, A. AND TELLER, E.: *Equation of state calculations by fast computing machines*, J. of Chem. Physics **21** (1953), 1087–1092.

98. MICHALEWICZ, Z.: *Genetic algorithms + data structures = evolution programs*, Springer, 1994.

99. MITCHELL, J.E. AND BORCHERS, B: *Solving real-world linear ordering problems using a primal-dual interioir point cutting plane method*, Annals of Operations Research **62** (1996), 253–276.

100. MITCHELL, J.E. AND BORCHERS, B: *Solving linear ordering problems*, in: Frenk, H., Roos, K., Terlaky, T. and Zhang, s. (eds.), *High Performance Optimization*, Applied Optimization Vol. 33) Kluwer, 1999, 349–366
`www.rpi.edu/~mitchj/generators/linord.`

101. NADDEF, D. AND G. RINALDI, G.: *The crown inequalities for the symmetric traveling salesman polytope*, Mathematics of Operations Research **17** (1992), 308–326.

102. NEMHAUSER, G.L. AND WOLSEY, L.A. *Integer and Combinatorial Optimization*, John Wiley and Sons, New York 1999.

103. NEWMAN, A.: *Cuts and orderings: On semidefinite relaxations for the linear ordering problem*, in: *Proc. of APPROX 2004*, LNCS 3122, Springer, 195–206.

104. NEWMAN, A. AND VEMPALA, S.: *Fences are futile: On relaxations for the linear ordering problem*, in: *Proc. of the 8th Int. IPCO Conference*, LNCS 2081, Springer, 2001, 333–347.

105. ONWUBOLU, G.C. AND BABU, B.V.: *New optimization techniques in engineering*, Studies in Fuzziness and Soft Computing, Vol. 141, Springer, 1994.

106. OSMAN, I.H. AND KELLY, J.P.: *Meta-heuristics: Theory and applications*, Kluwer Academic Publishers, 1996.

107. OSWALD, M., REINELT, G. AND SEITZ, H.: *Applying mod-k-cuts for solving linear ordering problems*, Top **17** (2009), 158–170.

108. POETSCH, G.: *Lösungsverfahren zur Triangulation von Input-Output-Tabellen*, Rechenzentrum Graz, Heft 79, 1973.

109. RALPHS, T.K. AND LADANYI, L.: *SYMPHONY: A parallel framework for branch and cut*, Rice University, 2000.
`http://branchandcut.org`

110. REEVES, C.: *Modern heuristics techniques for combinatorial optimization problems*, Advanced topics in computer science, Mc Graw-Hill, 1995.

111. REINELT, G.: *The linear ordering problem: Algorithms and applications*, research and exposition in mathematics 8, Heldermann, 1985.

112. REINELT, G.: *A note on small linear ordering polytopes*, Discrete & Computational Geometry **10** (1993), 67–78.

113. RENDL, F.: *Semidefinite relaxations for integer programming*, in: Jünger, M., Liebling, Th.M., Naddef, D., Nemhauser, G.L., Pulleyblank, W.R., Reinelt, G., Rinaldi, G., Wolsey, L.A. (eds.) *50 Years of Integer Programming 1958-2008: From the Early Years to the State-of-the-Art*, Springer, 2009, 687–726.

114. RIBEIRO, C.C., UCHOA, E. AND WERNECK, R.F., *A hybrid GRASP with perturbations for the Steiner problem in graphs*, INFORMS Journal on Computing **14** (2002), 228 - 246.

115. RINNOOY KAN, A.H.G. AND TIMMER, G.T.: *Global optimization*, in: Nemhauser, G.L., Rinnooy Kan, A.H.G. and Todd, M.J. (eds.) *Handbooks in operations research and management science, Vol. 1*, North Holland, 1989, 631–662.

116. SCHIAVINOTTO, T. AND STÜTZLE, T.: *The linear ordering problem: Instances, search space analysis and algorithms*, J. of Mathematical Modelling and algorithms **3** (2004), 367–402.

117. SCHRIJVER, A.: *Theory of linear and integer programming*, John Wiley and Sons, New York 1999.

118. SCHRIJVER, A.: *Combinatorial optimization. Polyhedra and efficiency*, Springer-Verlag, Berlin 2003.
119. SLATER, P.: *Inconsistencies in a schedule of paired comparisons*, Biometrika **48** (1961), 303–312.
120. SOLIS, F. AND WETS, R.: *Minimization by random search techniques*, Math. of Operations Research **6** 1981, 19–30.
121. SUCK, R.: *Geometric and combinatorial properties of the polytope of binary choice probabilities*, Math. Social Sciences **23** (1992), 81–102.
122. TRANSPARENCY INTERNATIONAL:
      `www.transparency.org/policy_research/surveys_indices/cpi`.
123. TSPLIB: *A library of benchmark instances for the traveling salesman problem*,
      `www.ifi.uni-heidelberg.de/groups/comopt/software/TSPLIB95`.
124. WOLSEY, L.A. *Integer Programming*, John Wiley and Sons, New York 1998.
125. YOUNG, H. P.: *On permutations and permutation polytopes*, Mathematical Programming Study **8** 1978, 128–140.

Index